

# IP (v4) ADDRESSING: INTRODUCTION

- **IP (v4) address:** 32-bit identifier for host, router **interface**

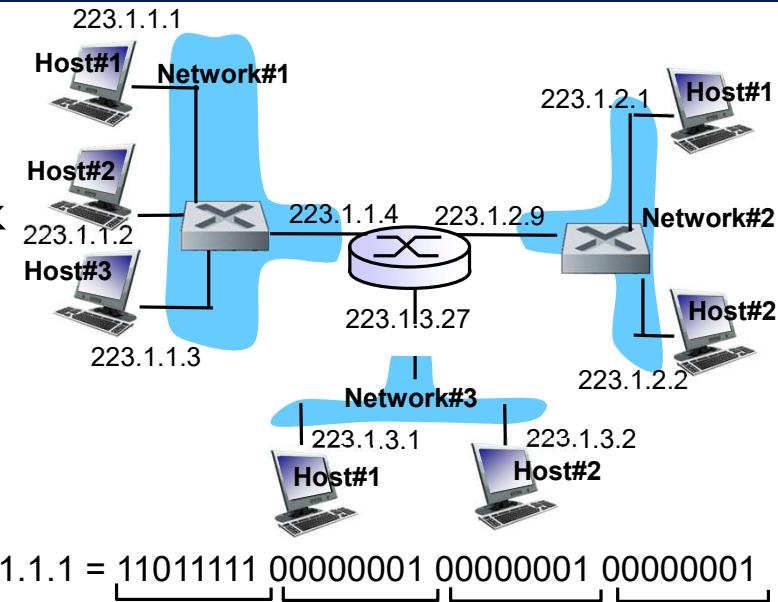
- **interface:** connection between host/router and physical link

- router's typically have multiple interfaces
- host typically has few (one/two?) interfaces (e.g., Ethernet, WiFi)

- **IP addresses are associated with each interface of L3 device**

- **32-bits as 4-octets define {network ID, Host ID}**

- **Historical Class-based Addressing:**



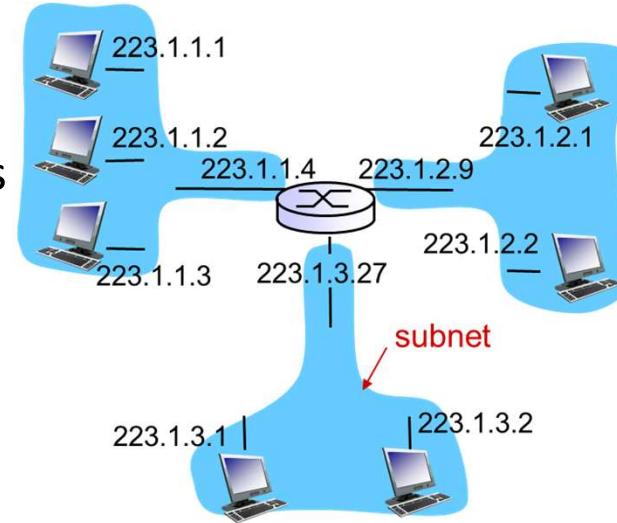
$223.1.1.1 = \underline{11011111} \underline{00000001} \underline{00000001} \underline{00000001}$

Address Class	Address Range	First Octet	Subnet Mask	# IP Addresses in n/w	# Networks
Class A (0)	0. 0 – 127.255.255.255	0-127	255.0.0.0 (8)	$2^{24}$	$128 (2^7)$
Class B (10)	128.0 – 191.255.255.255	128-191	255.255.0.0 (16)	$2^{16}$	$16384 (2^{14})$
Class C (110)	192.0 – 223.255.255.255	192-223	255.255.255.0 (24)	$2^8$	$2097152 (2^{21})$
Class D (1110)	224.0 – 239.255.255.255	224-239	-	$2^{28}$ (multicast groups)	-
Class E (1111)	240.0 – 255.255.255.255	240-255	-	$2^{28}$ (reserved space)	-

# IP ADDRESSING: CIDR

## CIDR: Classless InterDomain Routing

- subnet portion of address of arbitrary length
- Format: **a.b.c.d/x**; where x is # bits in subnet portion of address



network consisting of 3 subnets

- **IP address:**
  - subnet part - high order bits
  - host part - low order bits
- **What's a subnet ?**
  - Each isolated network
  - device interfaces with same subnet part of IP address
  - can physically reach each other *without intervening router*

## IP ADDRESSES: HOW TO GET ONE?

**Q:** how does *network* get subnet part of IP addr?

**A:** gets allocated portion of its provider ISP's address space

ISP's block	<u>11001000</u>	<u>00010111</u>	<u>00010000</u>	<u>00000000</u>	200.23.16.0/20
Organization 0	<u>11001000</u>	<u>00010111</u>	<u>00010000</u>	00000000	200.23.16.0/23
Organization 1	<u>11001000</u>	<u>00010111</u>	<u>00010010</u>	00000000	200.23.18.0/23
Organization 2	<u>11001000</u>	<u>00010111</u>	<u>00010100</u>	00000000	200.23.20.0/23
...	.....	.....	....	....	....
Organization 7	<u>11001000</u>	<u>00010111</u>	<u>00011110</u>	00000000	200.23.30.0/23

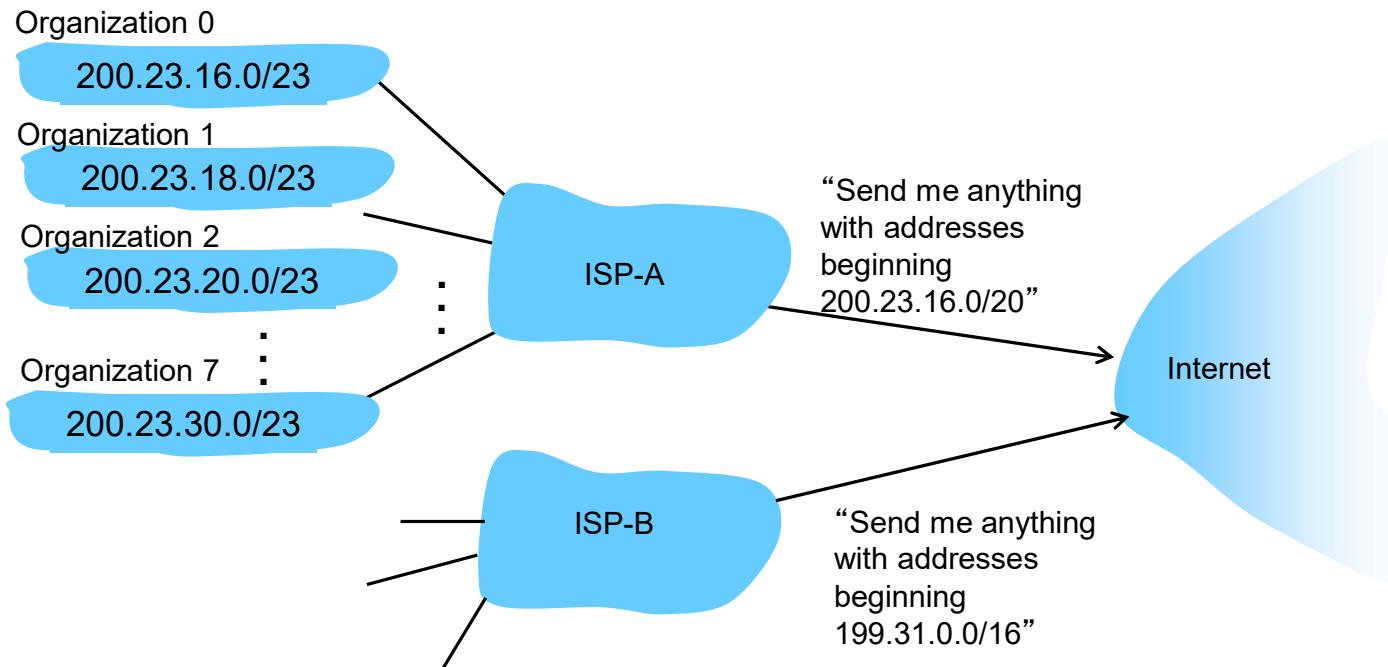
**Q:** how does an ISP get block of addresses?

**A:** ICANN: Internet Corporation for Assigned Names and Numbers

- allocates addresses (Public IPs)
- manages DNS
- assigns domain names, resolves disputes

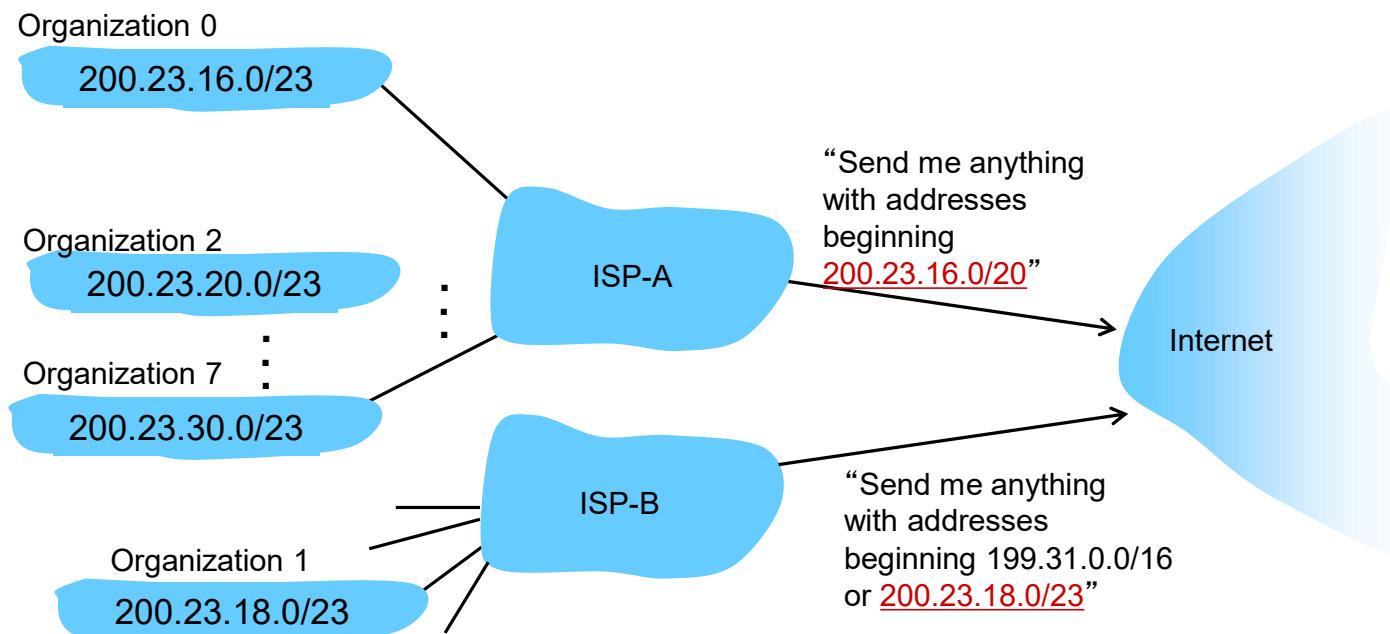
# HIERARCHICAL ADDRESSING: ROUTE AGGREGATION

hierarchical addressing allows efficient advertisement of routing information:



## HIERARCHICAL ADDRESSING: MORE SPECIFIC ROUTES

ISP-B has a more specific route to Organization 1

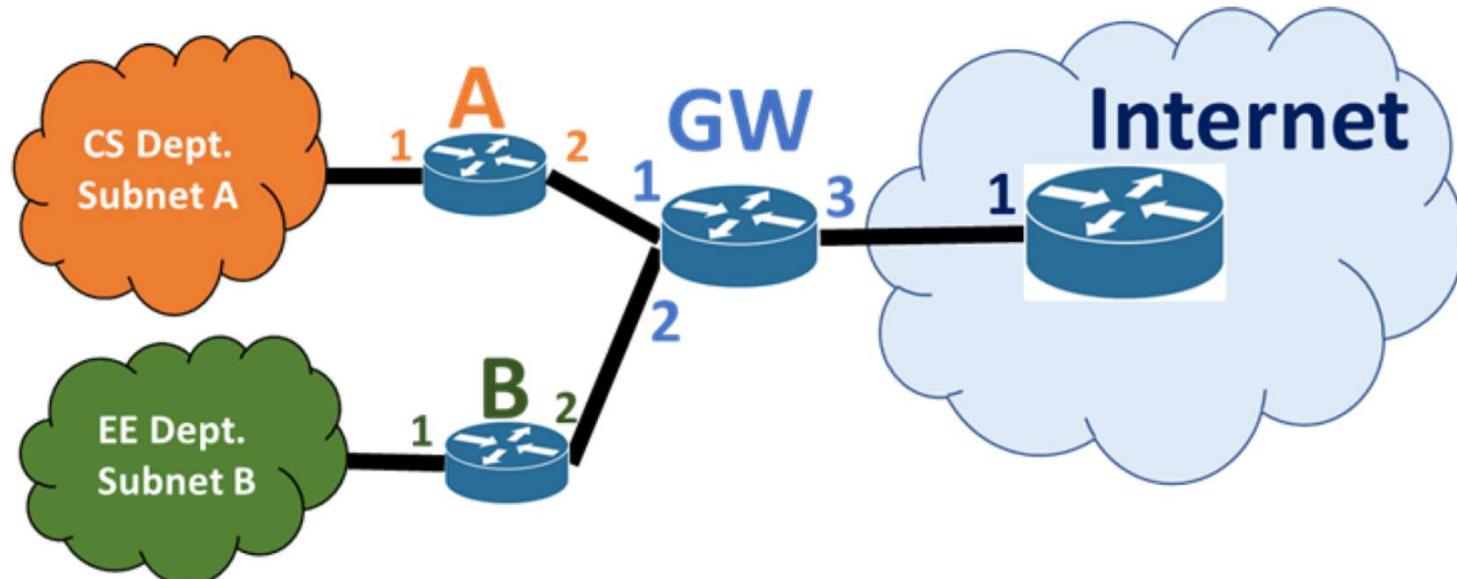


## SAMPLE QUESTIONS

2. An Institute has the following chunk of CIDR-based IP addresses available with it: **128.160.128.0/20**. From this chunk, it needs to allocate half of the addresses to the CS department, say Subnet A, and a quarter to the EE department, say Subnet B. Suggest how it can allocate the addresses (CIDR block) accordingly. Precisely answer the following: Note: range of IP addresses means Start and End IPs. **(10 pts)**
- a. Current IP address range, total number of hosts the Institute can support **(2 pts)**
  - b. CIDR block to be allocated for CS department (subet A) **(1 pts)**
  - c. Start and End IP addresses for Subnet A **(2 pts)**
  - d. CIDR block to be allocated for EE Department (subnet B) **(1 pts)**
  - e. Start and End IP addresses for Subnet B **(2 pts)**
  - f. Start and End IP addresses for the addresses retained by Institute **(2 pts)**

## SAMPLE QUESTIONS

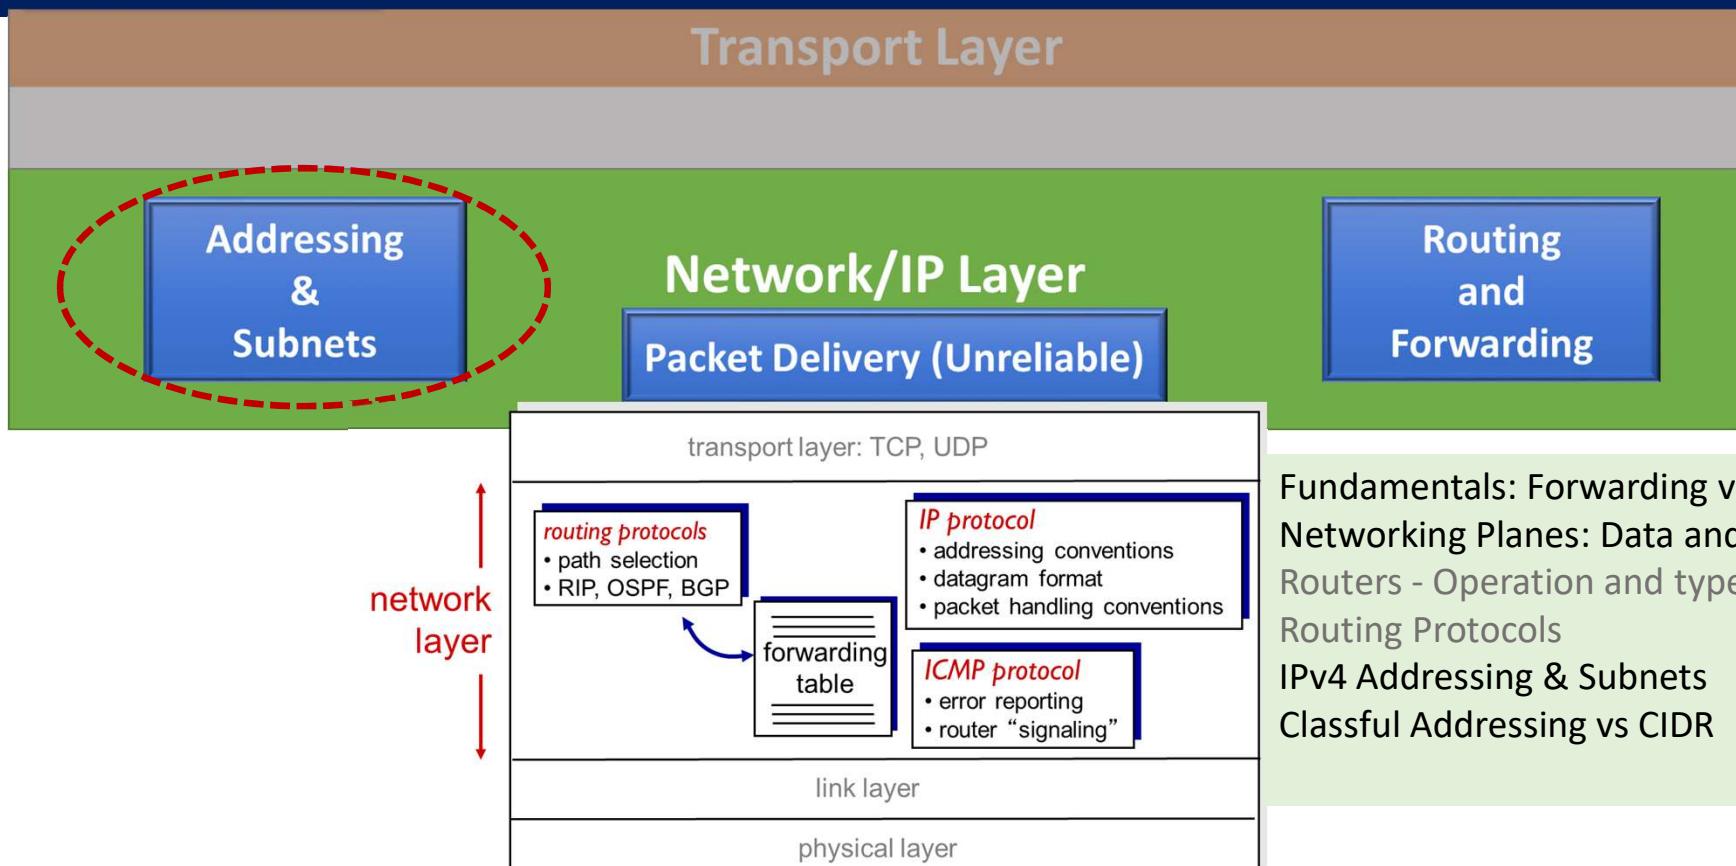
- a. As a map to the above case, describe the routing information (routers forwarding tables) that needs to be set at each of the routers in the below diagram: (5 pts)  
(Refer to the example table to fill the information):



Destination Address Range	Outgoing Interface

# NETWORK LAYER SERVICES

## Transport Layer



**Reading Material:** [RFC1122](#): by Robert Braden  
“Requirements for Internet Hosts — Communication Layers”

**Today's Focus:**  
IP Addressing – DHCP and NAT  
IPv6 Protocol and Addressing

## IP ADDRESSES: HOW TO GET ONE?

**Q:** How does a *host* get IP address?

- hard-coded by system admin in a file
  - Windows: control-panel->network->configuration->tcp/ip->properties
  - UNIX: /etc/rc.config
  - Linux: /etc/sysconfig/network-scripts/ifcfg-ethXX
- **DHCP: Dynamic Host Configuration Protocol: [[RFC 2131](#)]**
  - dynamically get address from a server
  - “plug-and-play”
- DHCP can return more than just allocated IP address on subnet:
  - address of first-hop router for client
  - name and IP address of DNS sever
  - network mask (indicating network versus host portion of address)

# DHCP: DYNAMIC HOST CONFIGURATION PROTOCOL

**goal:** allow host to *dynamically* obtain its IP address from network server when it joins network

- can renew its lease on address in use
- allows reuse of addresses (only hold address while connected/“on”)
- support for mobile users who want to join network (more shortly)

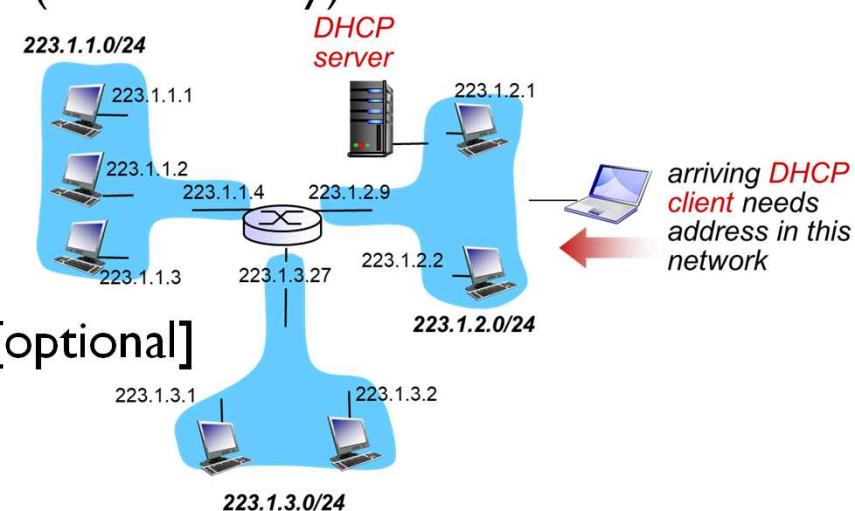
## DHCP overview:

1. host broadcasts “**DHCP discover**” msg [optional]

2. DHCP server responds with “**DHCP offer**” msg [optional]

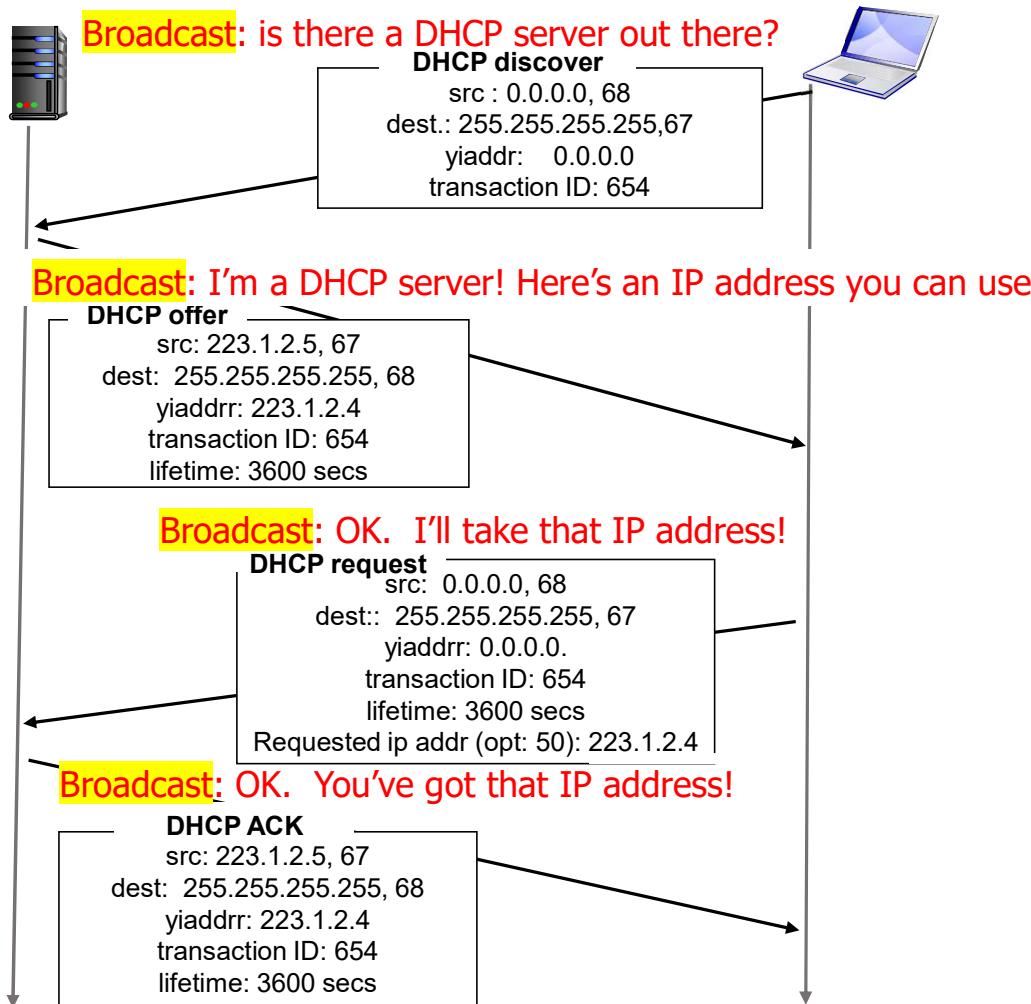
3. host requests IP address: “**DHCP request**” msg

4. DHCP server sends address: “**DHCP ack**” msg

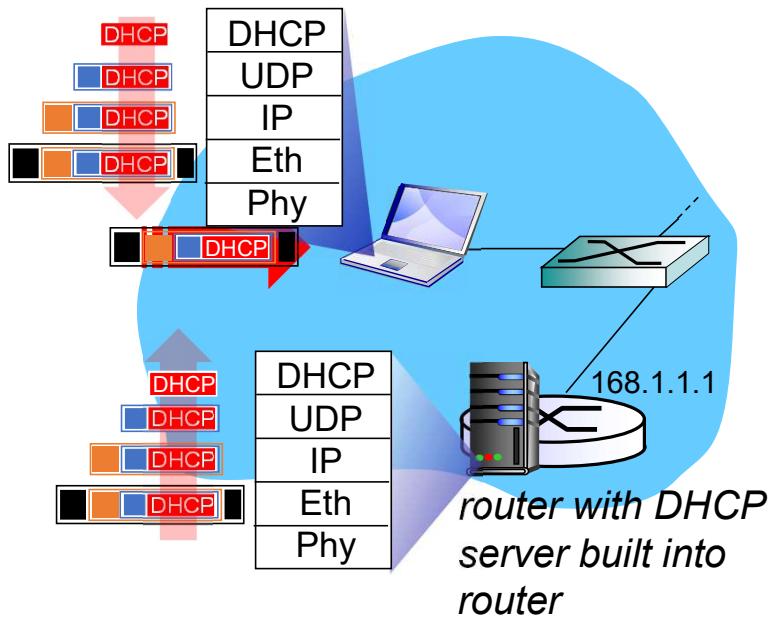


# DHCP CLIENT-SERVER SCENARIO

DHCP server: 223.1.2.5      New arriving client

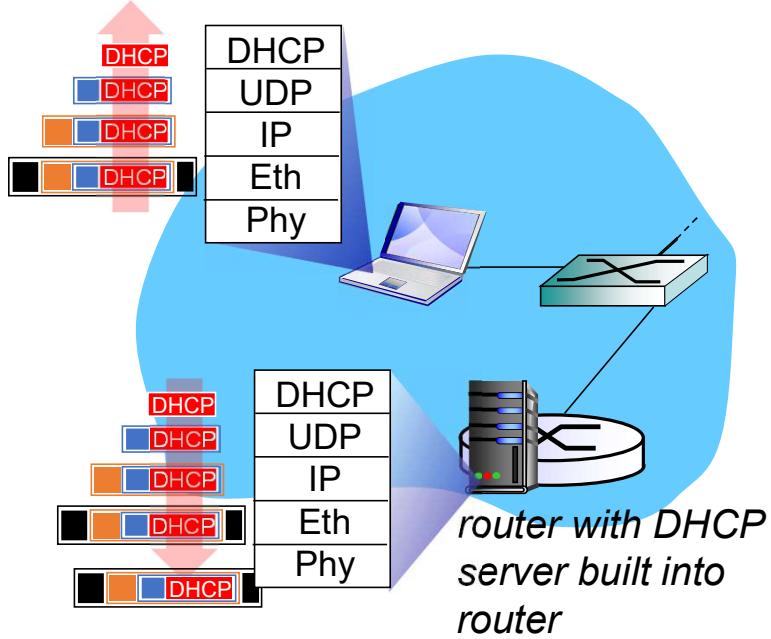


## DHCP: EXAMPLE



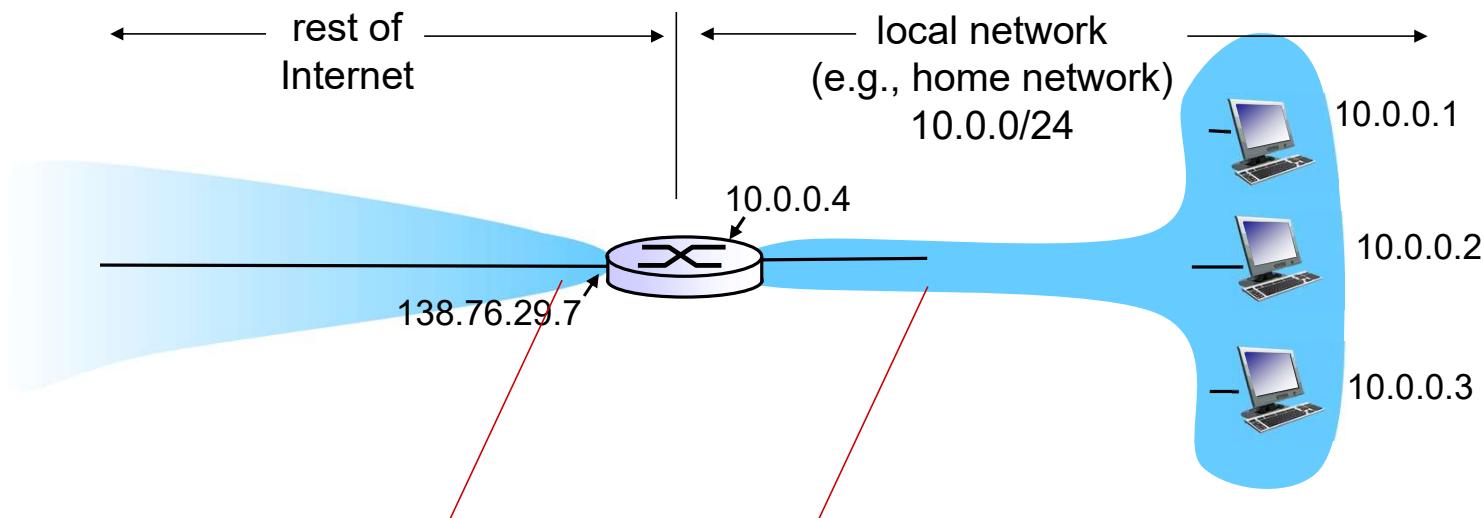
- connecting laptop needs its IP address, addr of first-hop router, addr of DNS server: use DHCP
- ❖ DHCP request encapsulated in UDP, encapsulated in IP, encapsulated in 802.1 Ethernet
- ❖ Ethernet frame broadcast (dest: FFFFFFFFFFFF) on LAN, received at router running DHCP server
- ❖ Ethernet demuxed to IP demuxed, UDP demuxed to DHCP

## DHCP: EXAMPLE



- DCP server formulates DHCP ACK (broadcast (dest: FFFFFFFFFFFF)) containing client's IP address, IP address of first-hop router for client, name & IP address of DNS server
  - ❖ encapsulation of DHCP server frame forwarded to client, demuxing up to DHCP at client
  - ❖ client now knows its IP address; name and IP address of DNS server; IP address of its first-hop router

## NAT(v4): NETWORK ADDRESS TRANSLATION



*all* datagrams *leaving* local network have *same* single source IP address i.e. NAT IP address: 138.76.29.7, mapped to different source port numbers.

datagrams with source or destination in this network have 10.0.0/24 address for source, destination (as usual)

## NAT: NETWORK ADDRESS TRANSLATION

### *motivation:*

local network uses just one IP as far as outside world is concerned

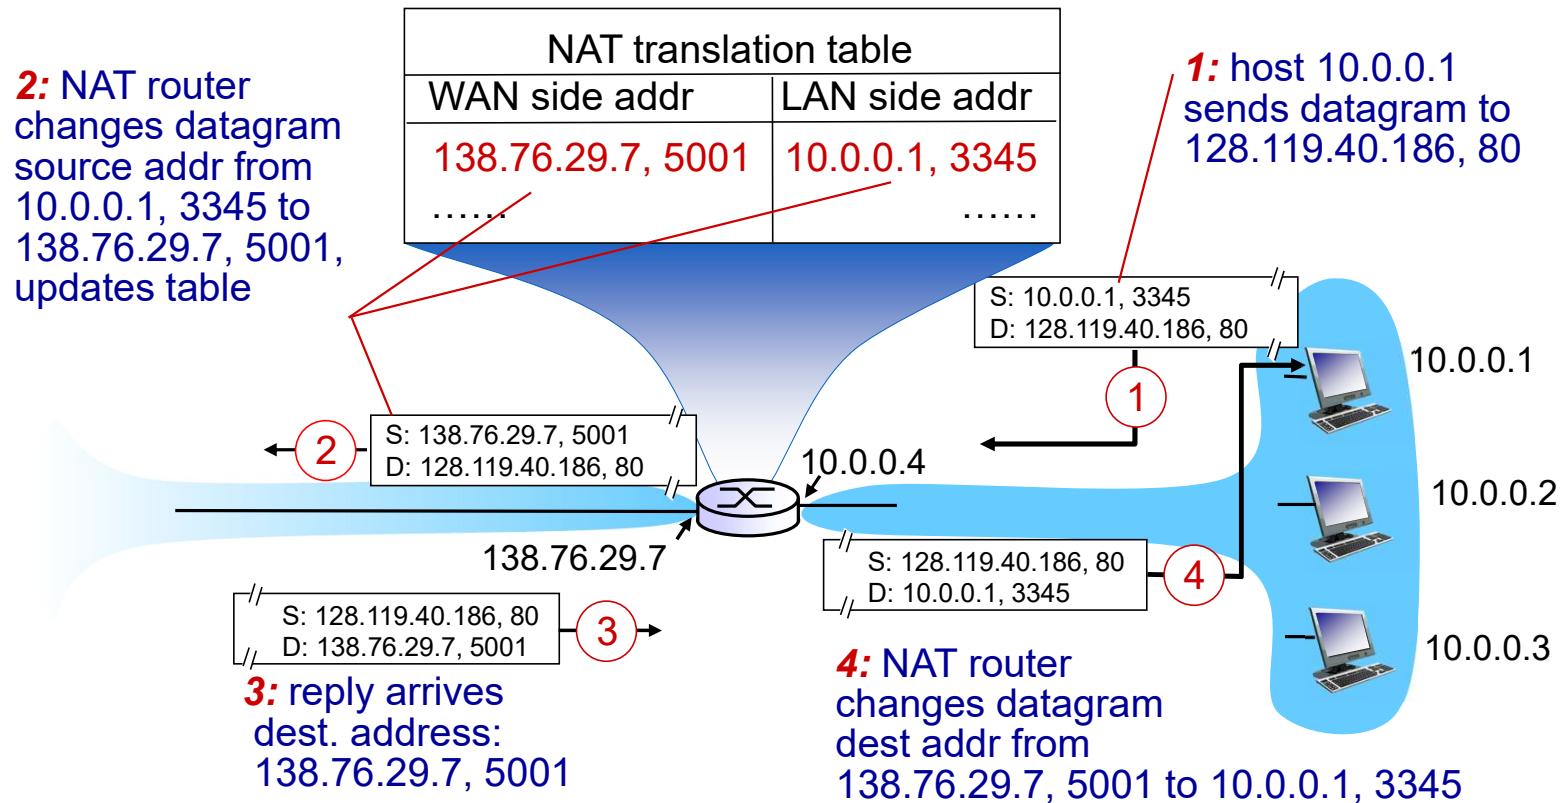
- range of addresses not needed from ISP: just one IP address for all devices
- can change addresses of devices in local network without notifying outside world
- can change ISP without changing addresses of devices in local network
- devices inside local net not explicitly addressable, visible by outside world (a security plus)

## NAT: NETWORK ADDRESS TRANSLATION

*implementation:* NAT router must:

- *outgoing datagrams: replace* (source IP address, port #) of every outgoing datagram to (NAT IP address, new port #)  
... remote clients/servers will respond using (NAT IP address, new port #) as destination addr
- *remember (in NAT translation table)* every (source IP address, port #) to (NAT IP address, new port #) translation pair
- *incoming datagrams: replace* (NAT IP address, new port #) in dest. fields of every incoming datagram with corresponding (source IP address, port #) stored in NAT table

# NAT: NETWORK ADDRESS TRANSLATION



## NAT: NETWORK ADDRESS TRANSLATION

- 16-bit port-number field:
  - 64,000 simultaneous connections with a single LAN-side address!
- NAT is controversial:
  - routers should only process up to layer 3
  - violates end-to-end argument
    - NAT possibility must be taken into account by app designers, e.g., P2P applications
  - address shortage should instead be solved by IPv6

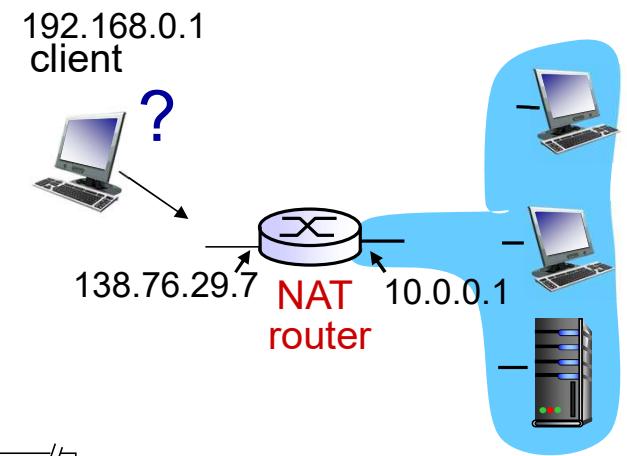
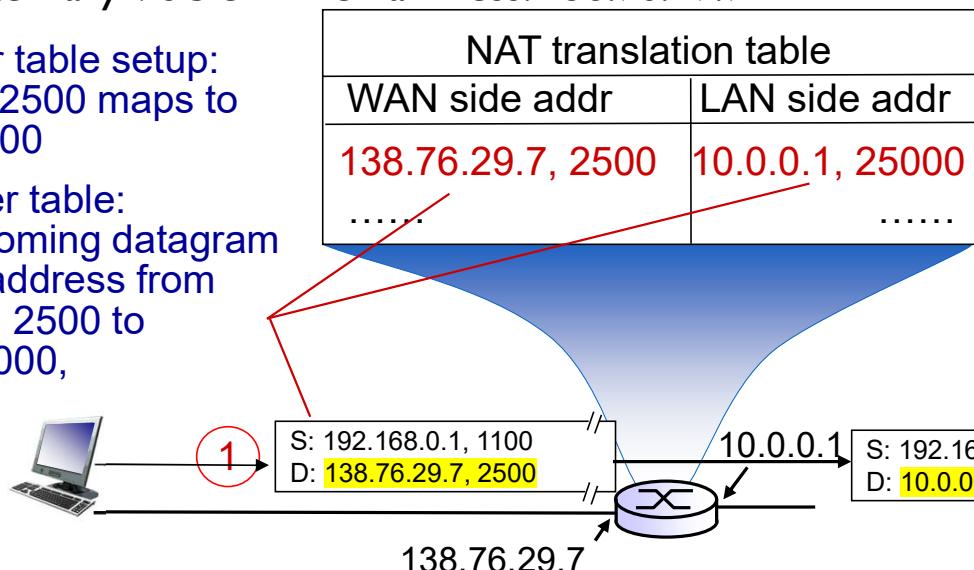
## NAT TRAVERSAL PROBLEM

- client wants to connect to server with address **10.0.0.1:25000**

- server address 10.0.0.1 local to LAN (client can't use it as destination address)
- only one externally visible NATed address: 138.76.29.7

**1:** NAT router table setup:  
138.76.29.7, 2500 maps to  
10.0.0.1, 25000

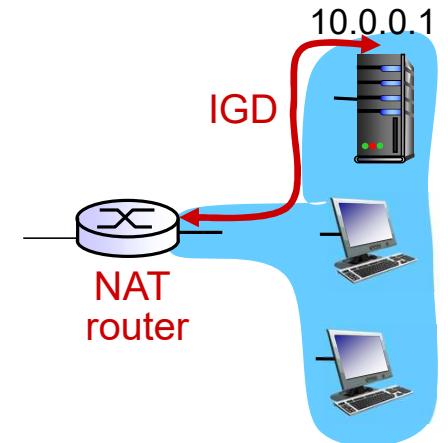
**2:** NAT router table:  
changes incoming datagram  
destination address from  
138.76.29.7, 2500 to  
10.0.0.1, 25000,



- solution 1:** Explicit Port Forwarding Rule Setup at the NAT routers
- statically configure NAT to forward incoming connection requests at a given port to a server.
  - e.g., (138.76.29.7, port 2500) always forwarded to (10.0.0.1 port 25000)

## NAT TRAVERSAL PROBLEM

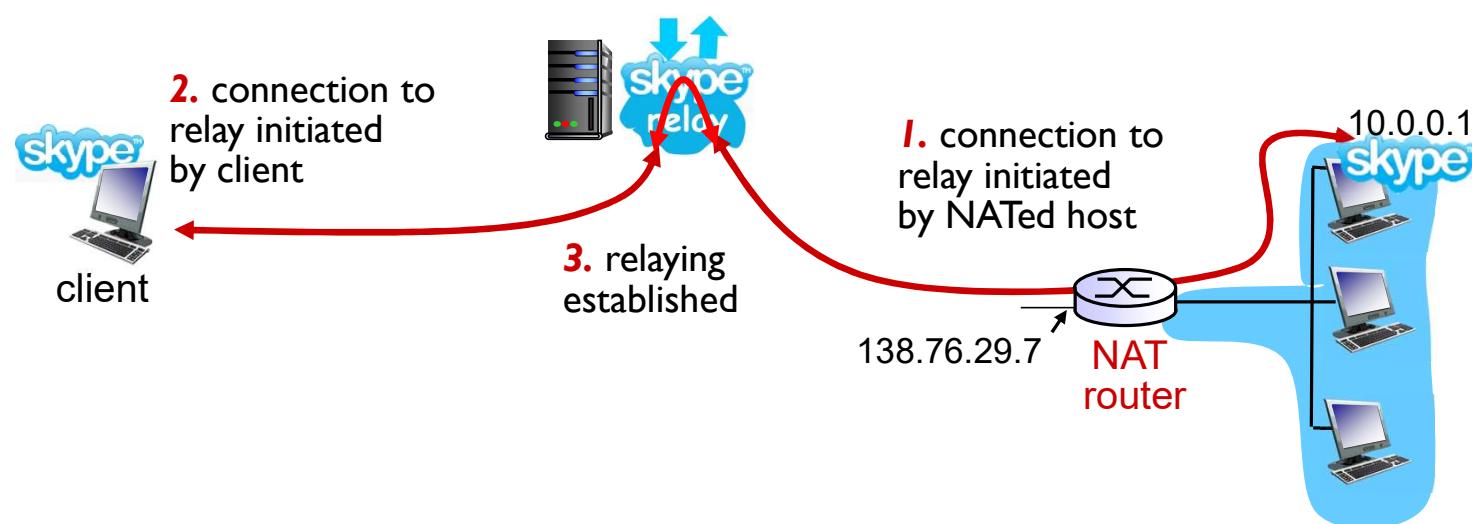
- **solution 2:** Universal Plug and Play (UPnP) Internet Gateway Device (IGD) Protocol. Allows NATed host to:
  - ❖ learn public IP address (138.76.29.7)
  - ❖ add/remove port mappings (with lease times)
- App. in host requests mapping for specific public port #:
- $(pvt. \text{ } IP \text{ } addrs, pvt. \text{ } port \text{ } \#) \Leftrightarrow (public \text{ } IP \text{ } addrs, public \text{ } port \text{ } \#)$ .
- NAT accepts request and creates necessary mapping;  
i.e., *automate the static NAT port map configuration*



Session Traversal Utilities for NAT (STUN)  
<https://www.rfc-editor.org/rfc/rfc5389>

## NAT TRAVERSAL PROBLEM

- **solution 3:** relaying (used in Skype)
  - NATed client establishes connection to relay
  - external client connects to relay
  - relay bridges packets between two connections



Traversal Using Relays around NAT (TURN)

# IPv6: MOTIVATION

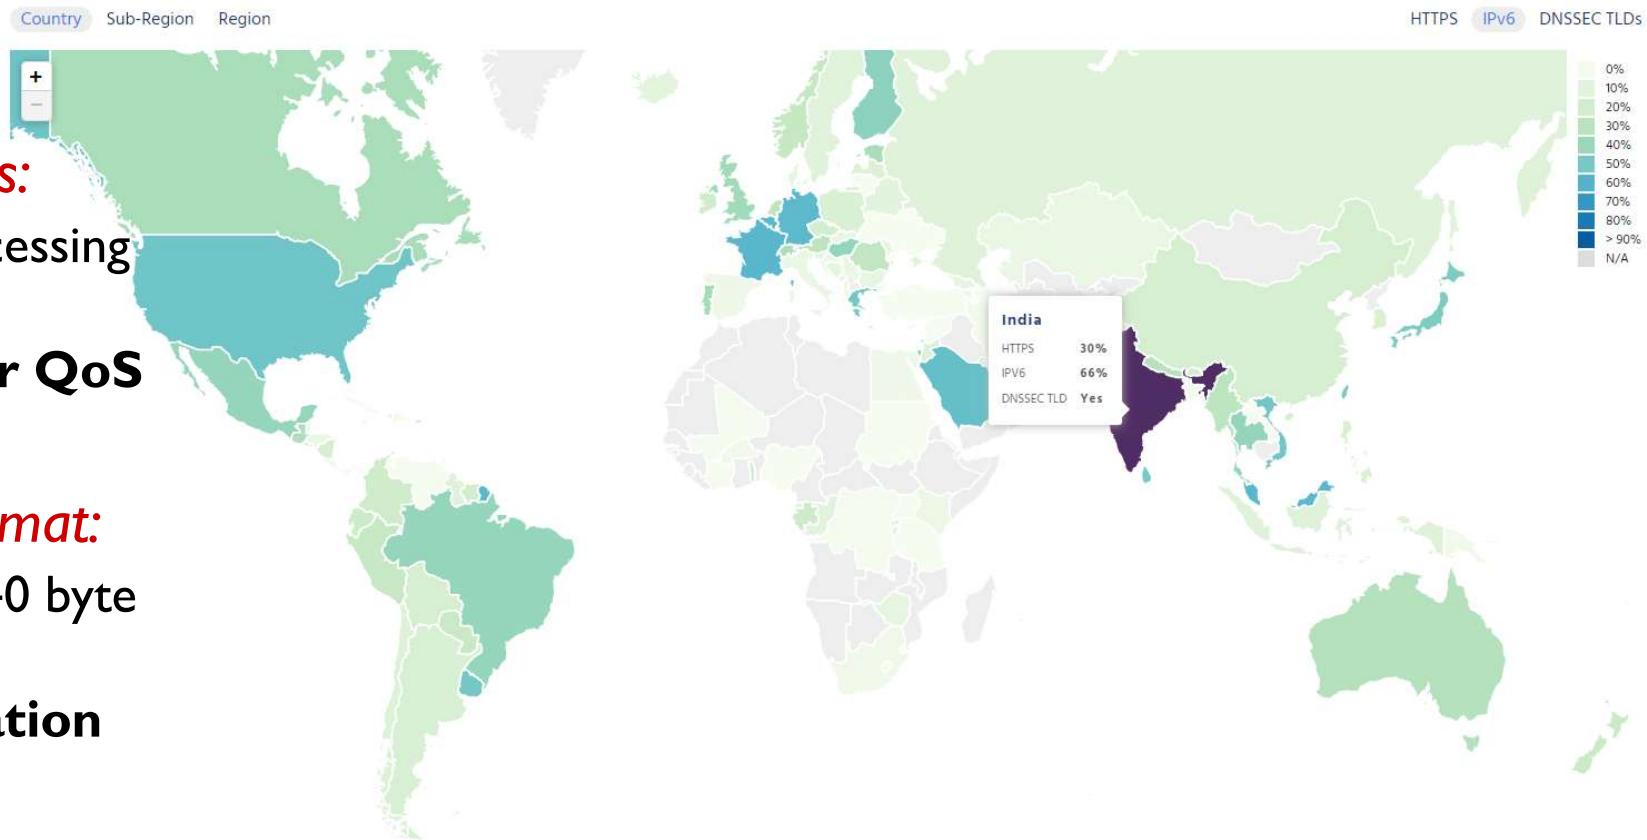
- *initial motivation:* 32-bit address space soon to be completely allocated.

- *other motivations:*

- **speed up processing /forwarding**
- **facilitate better QoS**

## IPv6 datagram format:

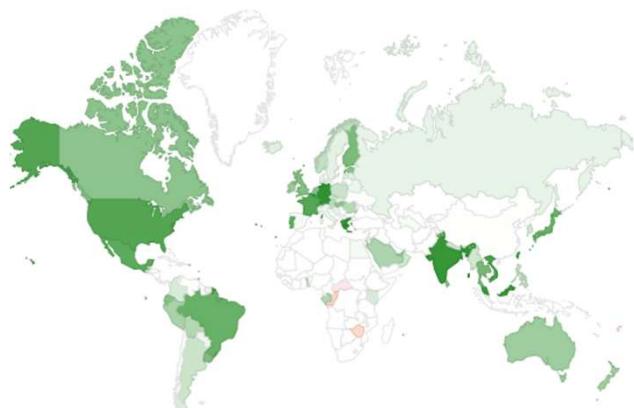
- **fixed-length 40 byte header**
- **no fragmentation allowed.**
- **No checksum**



<https://pulse.internetsociety.org/technologies>

## IPv6: ADOPTION

- 30~52% of traffic is IPv6.
- *Long! time for deployment, use*
  - 24+ years and counting!
  - *Why?*
  - think of application-level changes in last 2 decades:
    - Facebook, streaming media, Skype, Zoom, Hangout, Meet, Online Classrooms



India has reached an IPv6 adoption rate of around 60%; that is nearly double of US (September 2020 report).

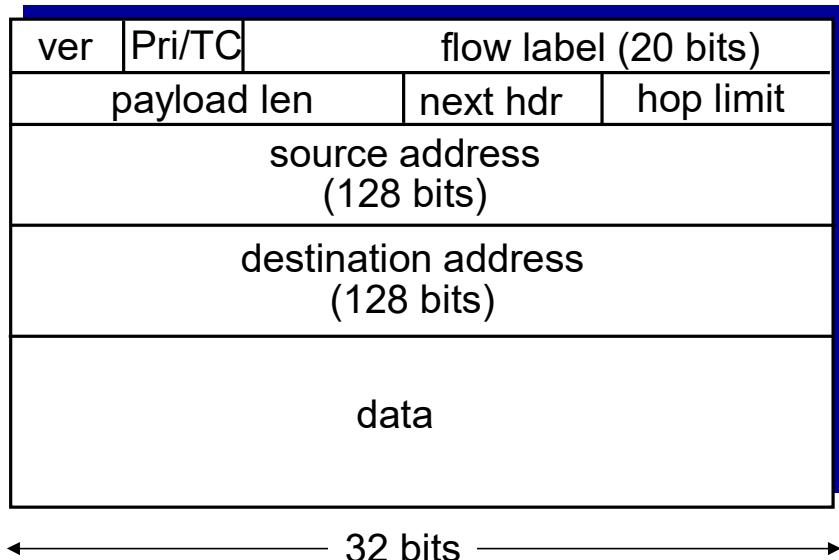
APNIC places India at more than 70% preferring IPv6. [\[154\]](#)

## IPv6 DATAGRAM FORMAT

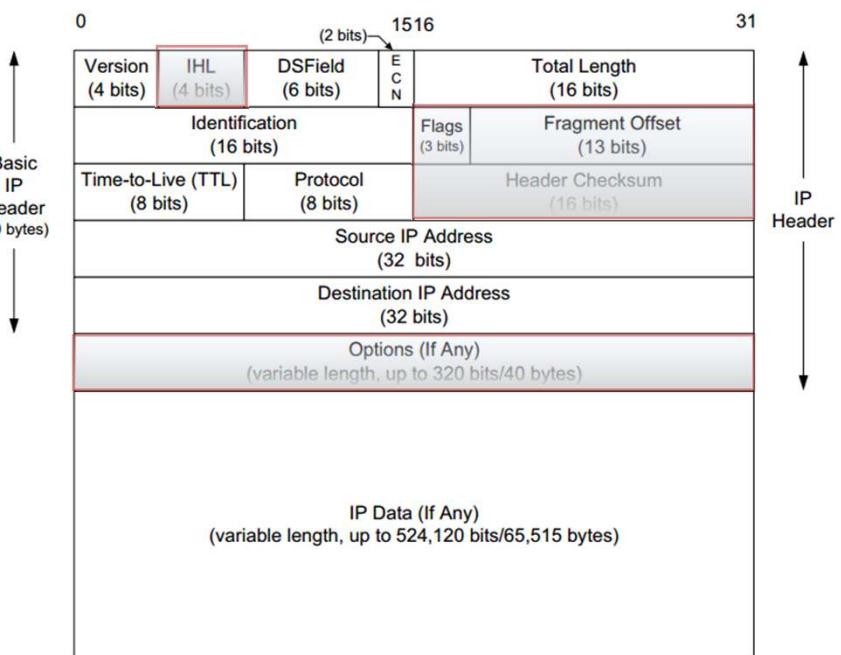
**priority:** identify priority among datagrams in flow

**flow Label:** identify datagrams in same “flow.” (concept of “flow” not well defined).

**next header:** identify the following (upper layer) protocol for data



VS



**checksum:** removed entirely to reduce processing time at each hop

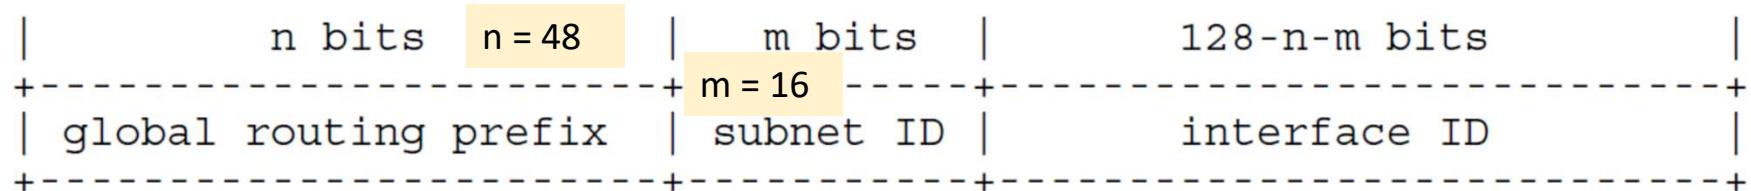
**options:** allowed, but outside of header, indicated by “Next Header” field

## IPV6 ADDRESSING (RFC 4921)

- **Modes of Addressing**

- Unicast: Link-local vs Global
- Multicast (FF00::/8)
- **Anycast.**

Address type	Binary prefix	IPv6 notation
-----	-----	-----
Unspecified	00...0 (128 bits)	::/128
Loopback	00...1 (128 bits)	::1/128
Multicast	11111111	FF00::/8
Link-Local unicast	1111111010	FE80::/10
Global Unicast	(everything else)	



- **Interoperability/Compatibility with IPv4**

- IPv4-Compatible IPv6 Address (First 96 bits =0, Last 32 bits = IPv4) – outdated.
- IPv4-Mapped IPv6 Address. (Last 32 bits=IPv4, next higher 16 bits=FFFF)
- Well Known Prefix: 64.ff9b::/96 or NSP specific Prex64::/n [[RFC 6052](#)]
- **Programming Interface:** new version of socket interface
  - [RFC 3493](#) : Basic Socket Interface Extension for Ipv6

## CHALLENGES (CHANGES) IN ADOPTING IPV6

- *Adoption*

- How to ensure every router in the world supports IPv6 from same date-time.
- Backward-compatibility is key.

- *Interoperability*

- Some devices/networks may still be IPv4 based.

- *Routing – Addressing*

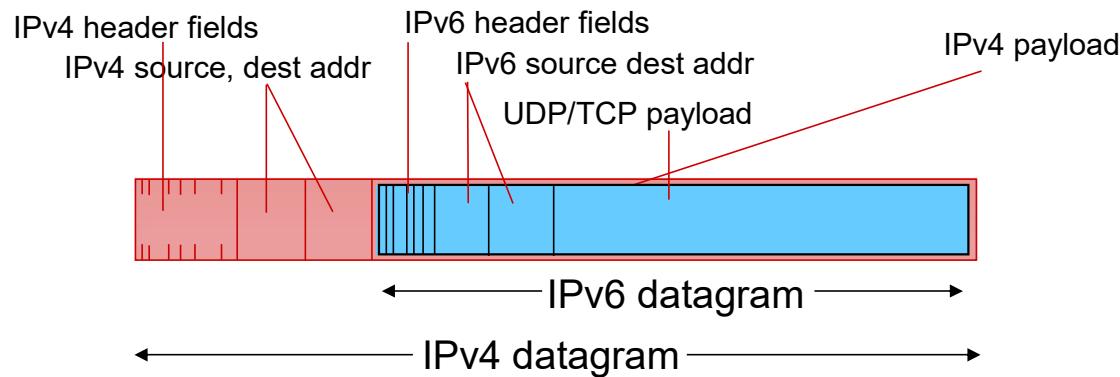
- Routing algorithms need to support both v4 and v6
- Routes may encompass multiple v4 and v6 hops.

- *Supporting Infrastructure:* new version of supporting protocols: ICMP/DHCP/NDP/DNS

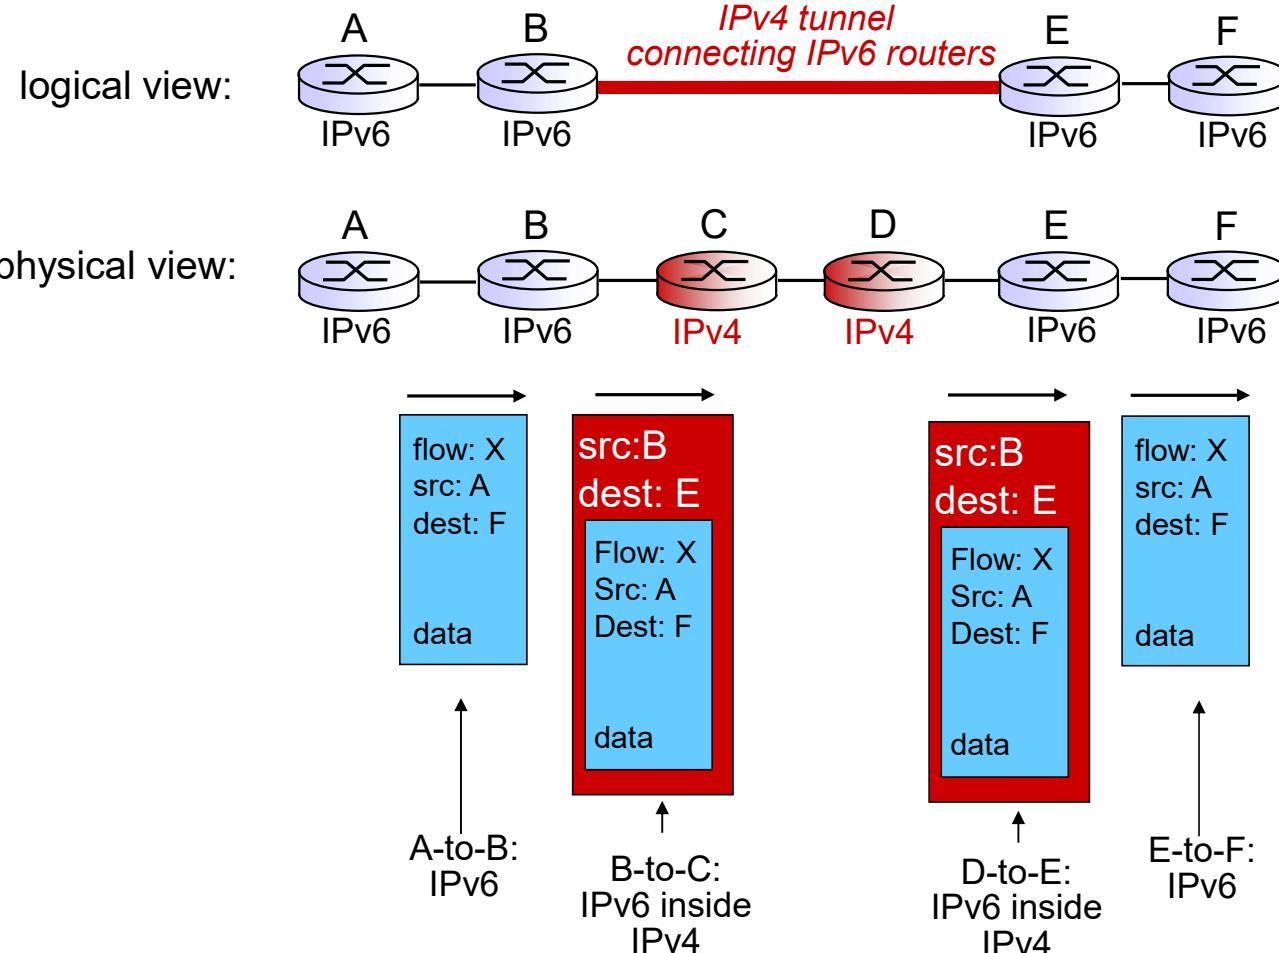
- additional message types, e.g. “Packet Too Big”
- multicast group management functions
- New DNS records for storing IPv6 addresses

## TRANSITION FROM IPv4 TO IPv6

- not all routers can be upgraded simultaneously
  - no “flag days”
  - how will network operate with mixed IPv4 and IPv6 routers?
- **tunneling:** IPv6 datagram carried as *payload* in IPv4 datagram among IPv4 routers



# TUNNELING



# ROUTERS – THE CORE NETWORKING DEVICES



## Cisco 8100 Routers

A fixed platform with up to 12.8 Tbps of capacity and optimized to reduce rack space and power costs.

- Chassis Form: Fixed
- Bandwidth: 6.4 to 12.8 Tbps
- Available Ports:  
32 QSFP28 100GbE;  
32 QSFP56-DD 400GbE;  
64 QSFP28 100GbE
- Height: 1 RU, 2 RU

## Cisco 8200 Routers

A fixed platform with 10.8 Tbps of capacity for deployment in space and power constrained facilities.

- Chassis Form: Fixed
- Bandwidth: 10.8 Tbps
- Available Ports:  
12/24 QSFP56-DD 400GbE
- Height: 1 RU, 2 RU



## Cisco 8800 Routers

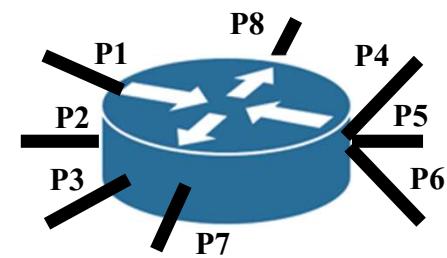
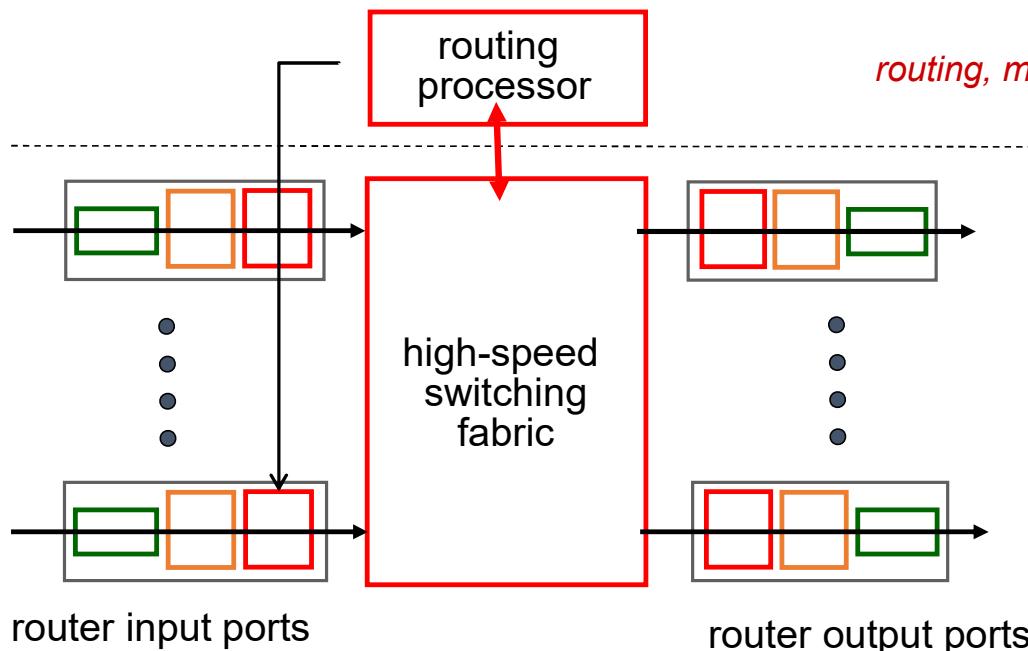
A high performance, high density modular platform with up to 259.2 Tbps of capacity that can consolidate the number of routers needed and reduce overall complexity.

- Chassis Form: Modular
- Bandwidth: Up to 259.2 Tbps
- Available Ports:  
36 QSFP56-DD 400GbE;  
48 QSFP28 100GbE with MACsec
- Height: 16 RU (8 slots), 21 RU (12 slots),  
33 RU (18 slots)

Source – Cisco Routers

# ROUTER ARCHITECTURE OVERVIEW

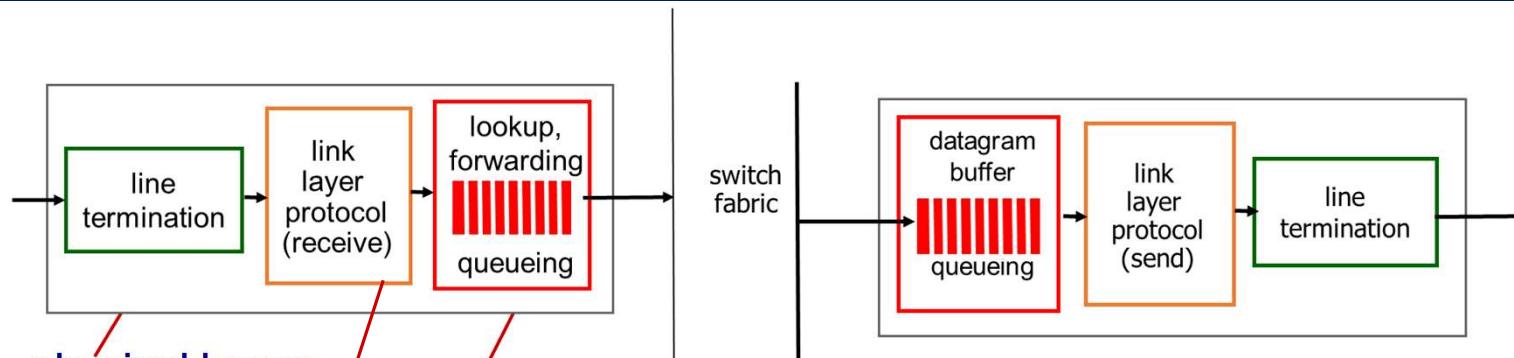
- high-level view of generic router architecture:



*routing, management: control plane (software)*  
operates in **second** time frame

*forwarding data plane (hardware)*  
operates in **nanosecond** timeframe

# INPUT PORT VS OUTPUT PORT FUNCTIONS



physical layer:  
bit-level reception

data link layer:  
Data Framing

Network layer:  
Packet switching

- **decentralized switching:**

Use datagram header fields to lookup output port in the forwarding table (**"match plus action"**)

decentralized switching modes:

- **destination-based forwarding:** forward based only on destination IP address (traditional)
- **generalized forwarding:** forward based on any set of header field values.

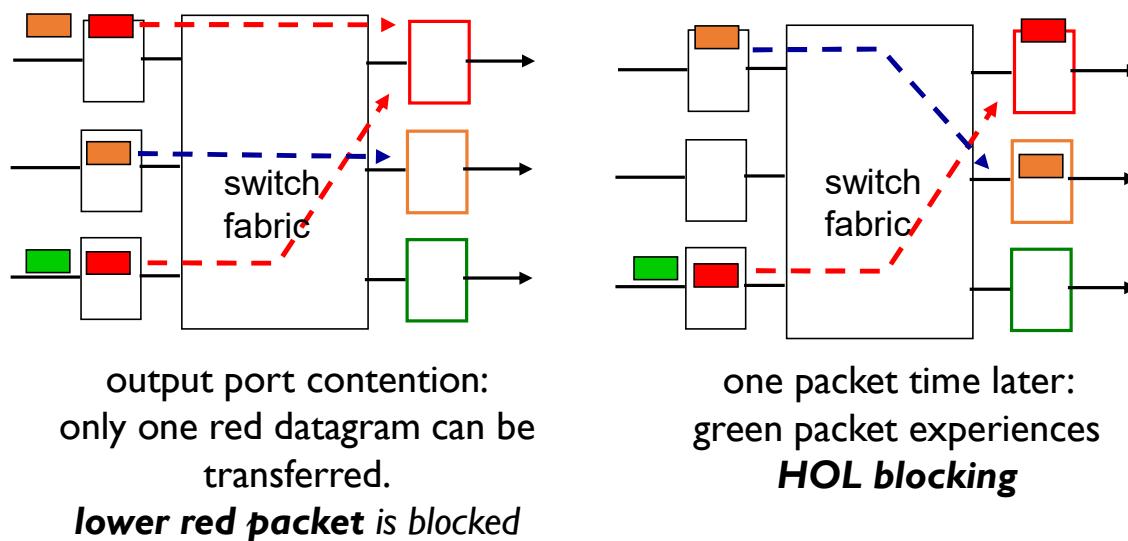
- **buffering** required when datagrams arrive from fabric at a rate slightly faster than the transmission rate

Datagram can be lost due to lack of buffers!

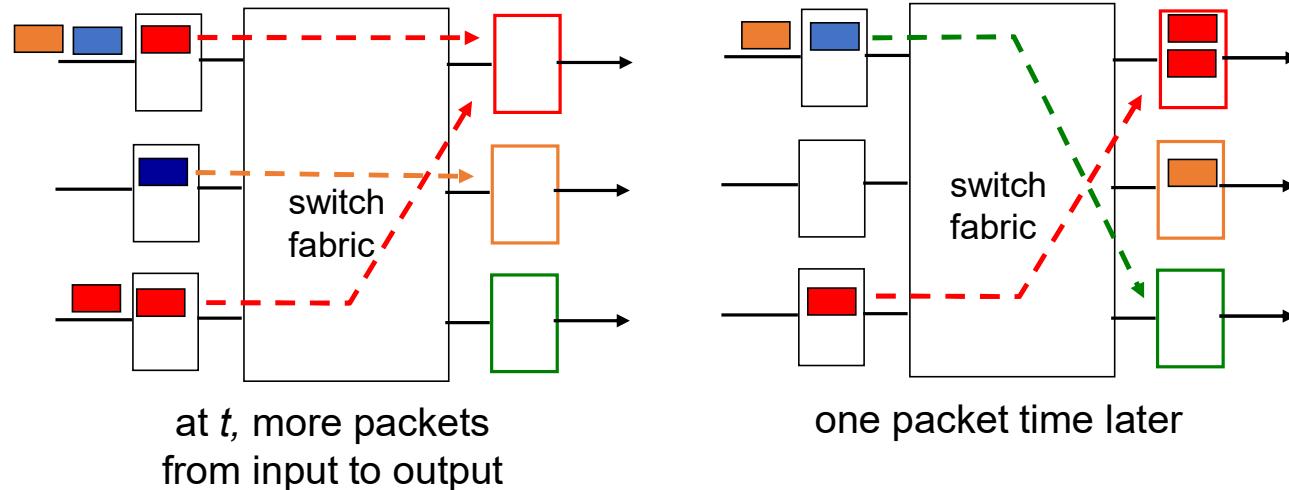
- **scheduling discipline** chooses among queued datagrams for transmission

## INPUT PORT QUEUING

- Queuing: when the datagrams arrive faster than forwarding rate of switch fabric.
- fabric slower than input ports combined -> queueing may occur at input queues
  - *queueing delay and loss due to input buffer overflow!*
- Head-of-the-Line (HOL) blocking: queued datagram at front of queue prevents others in queue from moving forward



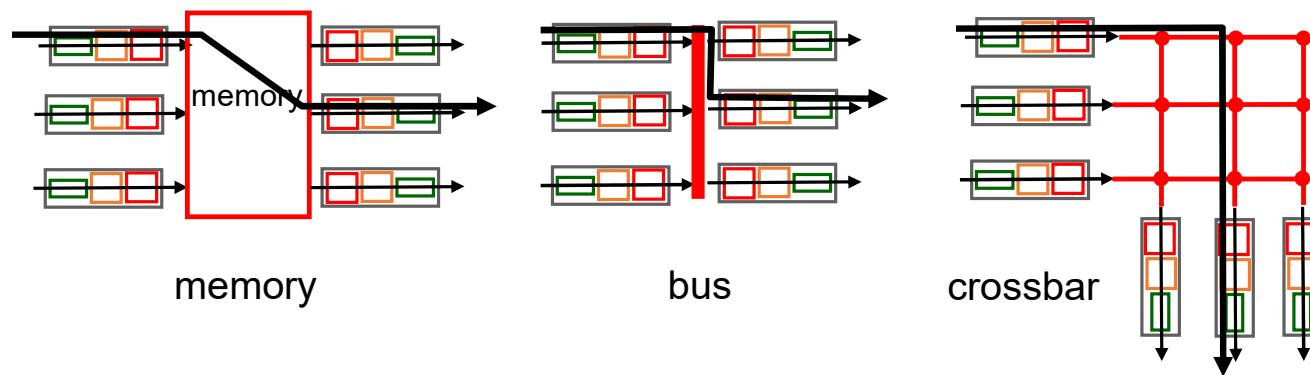
# OUTPUT PORT QUEUEING



- buffering when arrival rate via switch fabric exceeds output line speed
- *queueing (delay) and loss can also occur due to output port buffer overflow!*
- RFC 3439 rule of thumb: average buffering equal to “typical” RTT (say 250 msec) times link capacity C.
  - e.g., C = 10 Gpbs link: 2.5 Gbit buffer
- recommendation ([2004](#)): with N flows, buffering equal to =  $\frac{RTT \cdot C}{\sqrt{N}}$

# SWITCHING FABRICS

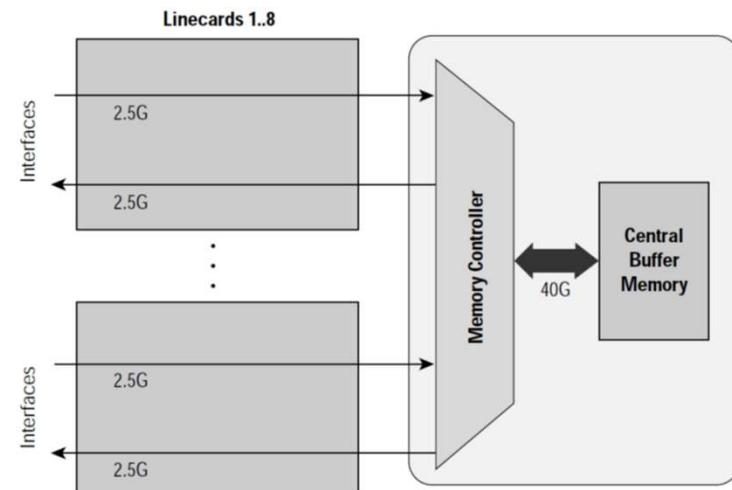
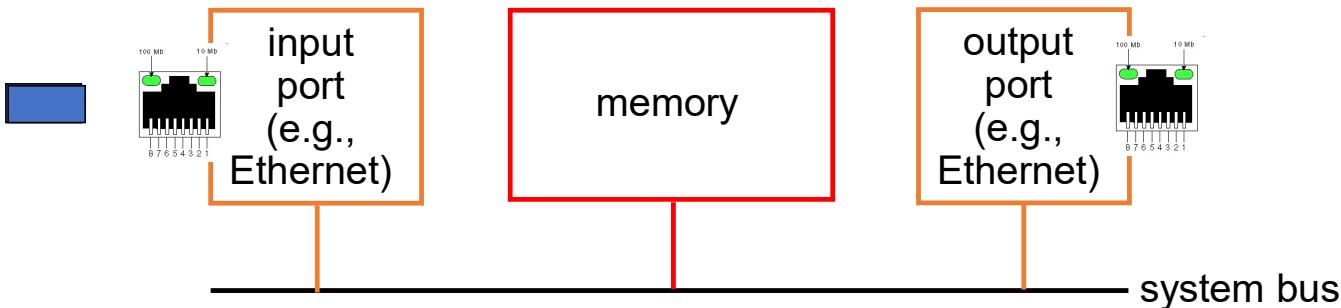
- transfer packet from input buffer to appropriate output buffer
- switching rate: rate at which packets can be transferred from inputs to outputs
  - often measured as multiple of input/output line rate
  - $N$  inputs: desirable switching rate  $N$  times line rate
- three types of switching fabrics



# SWITCHING VIA MEMORY

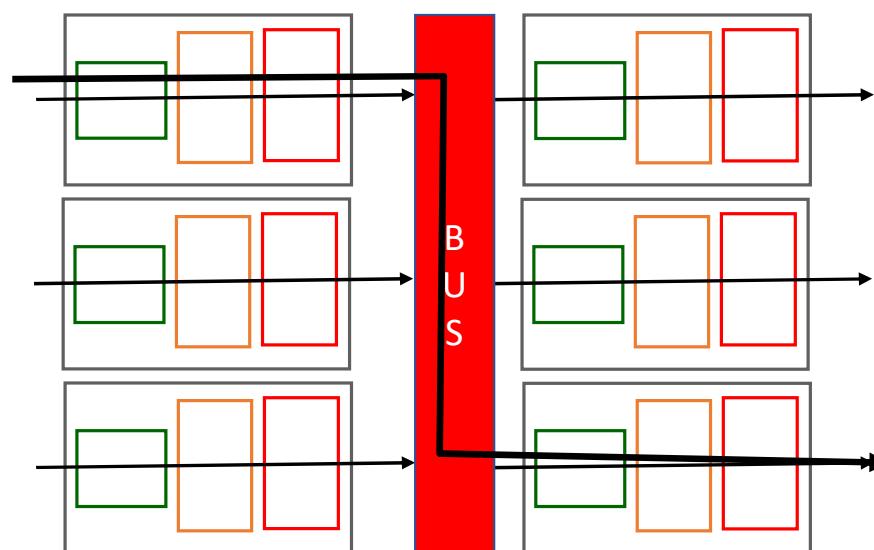
*first generation routers: Centralized shared Memory*

- Similar to traditional computers:  
with switching under direct control of CPU
- Incoming packet copied to system's memory
- Processing speed is limited by memory bandwidth  
(need 2 bus crossings per datagram)



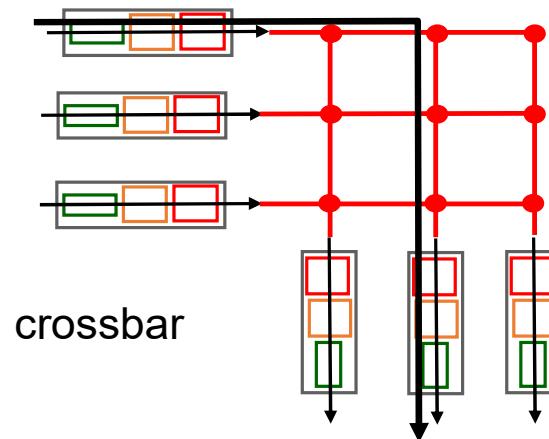
## SWITCHING VIA A BUS

- datagram from input port memory to output port memory via a shared bus
- *bus contention*: switching speed limited by bus bandwidth
  - 32 Gbps bus, Cisco 5600: sufficient speed for access and enterprise routers

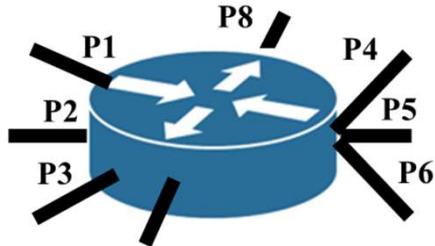


## SWITCHING VIA INTERCONNECTION NETWORK

- overcome bus bandwidth limitations
- banyan networks, crossbar, other interconnection nets initially developed to connect processors in multiprocessor
- advanced design: fragmenting datagram into fixed length cells, switch cells through the fabric.
- Cisco 12000: switches 60 Gbps through the interconnection network

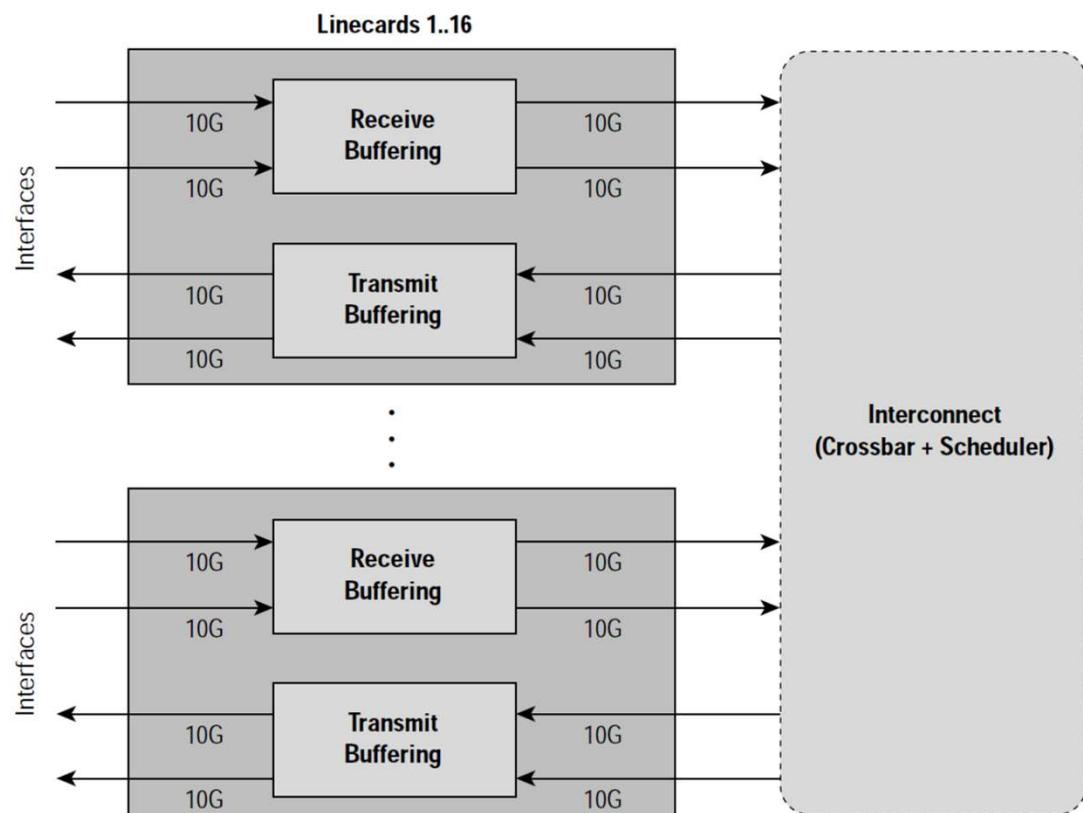


## TYPICAL ROUTER ARCHITECTURE



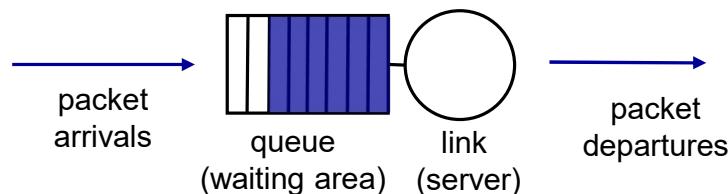
Router Jobs:

- Compute Best Path/Route → Role of a Router Processor (Control plane)  
Account for network policies & constraints
- Packet Forwarding → Role of a Forwarding Engine (Data Plane)
- Buffering and Scheduling → Role of a Scheduler (Data Plane)



# SCHEDULING MECHANISMS

- *scheduling*: choose next packet to send on link
- *discard policy*: if packet arrives to full queue: who to discard?
  - *tail drop*: drop arriving packet
  - *random*: drop/remove randomly
  - *priority*: drop/remove on priority basis

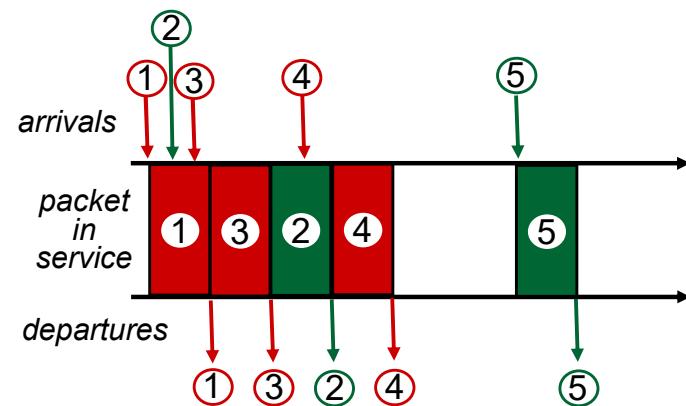
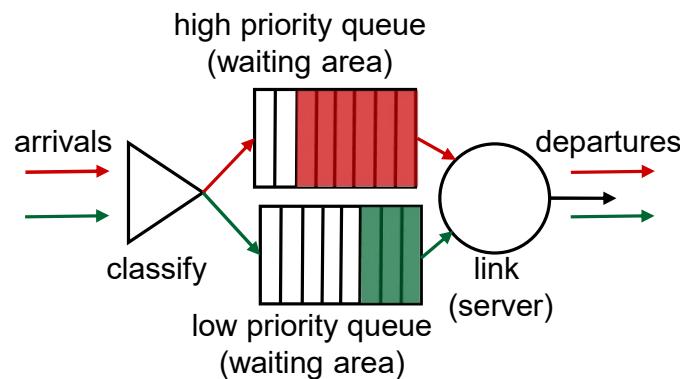


- *FIFO (first in first out) scheduling aka FCFS*: send in order of arrival to queue;

## SCHEDULING POLICIES: PRIORITY

**priority scheduling:** send highest priority queued packet

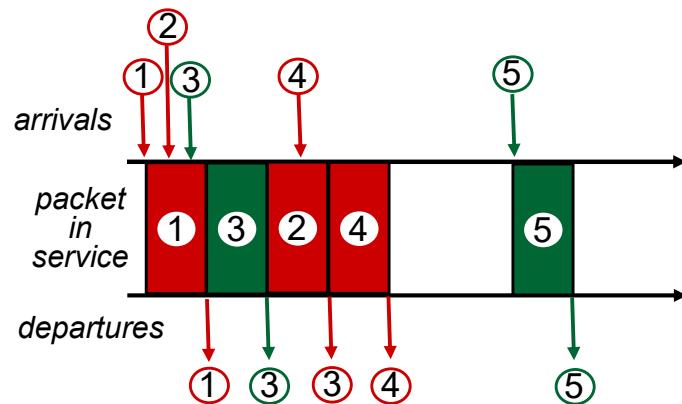
- multiple *classes*, with different priorities
  - class may depend on marking or other header info, e.g. IP source/dest, port numbers, etc.



## SCHEDULING POLICIES: STILL MORE

*Round Robin (RR) scheduling:*

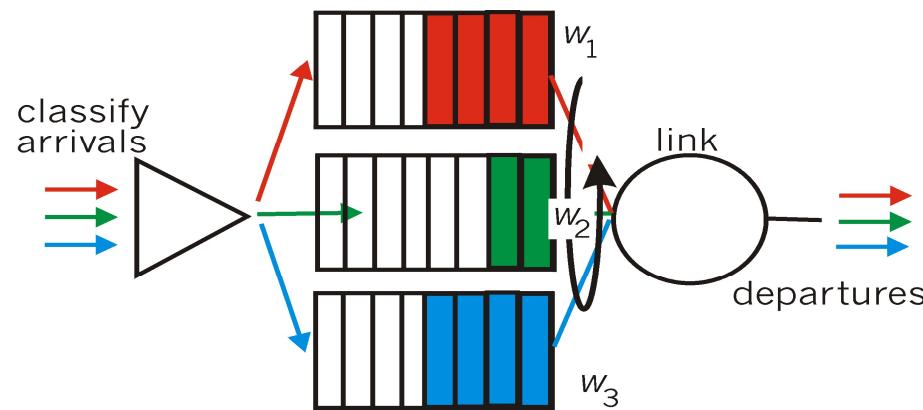
- multiple classes
- cyclically scan class queues, sending one complete packet from each class (if available)



## SCHEDULING POLICIES: STILL MORE

### *Weighted Fair Queuing (WFQ):*

- generalized Round Robin
- each class gets weighted amount of service in each cycle



# DESTINATION-BASED FORWARDING

lookup,  
forwarding  
  
queueing

forwarding table	
Destination Address Range	Link Interface
11001000 00010111 00010000 00000000 through (200.23.16.0 - 200.23.23.255) 11001000 00010111 00010111 11111111	0
11001000 00010111 00011000 00000000 through (200.23.24.0 - 200.23.24.255) 11001000 00010111 00011000 11111111	1
11001000 00010111 00011001 00000000 through (200.23.25.0 - 200.23.31.255) 11001000 00010111 00011111 11111111	2
otherwise	3



Q: but what happens if ranges don't divide up so nicely?

## LONGEST PREFIX MATCHING

lookup,  
forwarding  
  
queueing

### *longest prefix matching*

when looking for forwarding table entry for given destination address, use *longest* address prefix that matches destination address.

Destination Address Range	Link interface
11001000 00010111 00010*** *****	0
11001000 00010111 00011000 *****	1
11001000 00010111 00011*** *****	2
otherwise	3

examples:

DA: 11001000 00010111 00010110 10100001      which interface?

DA: 11001000 00010111 00011000 10101010      which interface?

- longest prefix matching: often performed using TCAMs
  - *content addressable*: retrieve address in one clock cycle, regardless of table size.
  - Cisco Catalyst: can up ~1M routing table entries in TCAM