


ROUTING PROTOCOLS

Goal: determine “good” paths (equivalently, routes), from sending hosts to receiving host, through network of routers

- **path:** sequence of routers packets will traverse in going from given initial source host to given final destination host
- **Metrics:** “good”: “least cost”, “fastest”, “least congested”




- routing: a “top-10” networking challenge!

GRAPH ABSTRACTION OF THE NETWORK

graph: $G = (N, E)$

$N = \text{set of routers} = \{ u, v, w, x, y, z \}$

$E = \text{set of links} = \{ (u,v), (u,x), (v,x), (v,w), (x,w), (x,y), (w,y), (w,z), (y,z) \}$



Metric cost of path: C

$c(x,x') = \text{cost of link } (x,x')$

e.g., $c(w,z) = 5$

cost could be directly related to latency,
or inversely related to length
or directly related to bandwidth,
or inversely related to congestion


cost of path $(x_1, x_2, x_3, \dots, x_p) = c(x_1, x_2) + c(x_2, x_3) + \dots + c(x_{p-1}, x_p)$

Note: Cost of non-existing link can be considered Infinity: $C(v,y) = \infty$

key question: what is the **least-cost** path between u and z ?

routing algorithm: algorithm that finds that least cost path

ROUTING (ROUTE BUILDING) ALGORITHM CLASSIFICATION



Q: global or decentralized information?

Global (Centralized):

- all routers have complete topology, link cost info
- “link state” algorithms

Decentralized:

- router knows physically-connected neighbors, link costs to neighbors
- iterative process of computation, exchange of info with neighbors
- “distance vector” algorithms

Q: static or dynamic?

static:

- routes change slowly over time

dynamic:

- routes change more quickly
 - periodic update
 - in response to link cost changes

LINK STATE ROUTING - CENTRALIZED

- Each node is assumed to know the ***state of links*** to its neighbors
- Step 1: Each node ***broadcasts*** its state to ***all*** other nodes
 - Reliable broadcast mechanism
 - flooding
 - *may have sequence number issues*
- Step 2: Each node ***locally computes*** shortest paths (***SSSP***) to all other nodes from its global state
 - Shortest path tree (SPT) algorithm
 - Dijkstra's SPT algorithm

A LINK-STATE ROUTING ALGORITHM

Dijkstra's algorithm (how to decide route)

- computes least cost paths from one node ('source') to all other nodes
 - gives *forwarding table* for that node
- iterative:** after k iterations, know least cost path to k destinations from a given source.

Given that we know: (preconditions)

- network topology, link costs of all the nodes.
 - accomplished via "**link state broadcast**"
 - all nodes have same *topology information*.

notation:

- $c(x,y)$: link cost from node x to y; set to ∞ if not direct neighbors
- $D(v)$: current value of cost of path from source to dest. v
- $p(v)$: predecessor node along path from source to v
- N' : set of nodes whose least cost path definitively known

DIJSKTRA'S ALGORITHM


```
1 Initialization:
2   N' = {u}
3   for all nodes v
4     if v adjacent to u
5       then D(v) = c(u,v)
6     else D(v) = ∞
7
8 Loop
9   find w not in N' such that D(w) is a minimum
10  add w to N'
11  update D(v) for all v adjacent to w and not in N' :
12    D(v) = min( D(v), D(w) + c(w,v) )
13  /* new cost to v is either old cost to v or known
14  shortest path cost to w plus cost from w to v */
15 until all nodes in N'
```

DIJKSTRA'S ALGORITHM (EXAMPLE)


Step	N'	$D(v)$ $p(v)$	$D(w)$ $p(w)$	$D(x)$ $p(x)$	$D(y)$ $p(y)$	$D(z)$ $p(z)$
0	u	7,u	3,u	5,u	∞	∞
1	uw	6,w		5,u	11,w	∞
2	uwx		6,w		11,w	14,x
3	uwxv				10,y	14,x
4	uwxvy					12,y
5	uwxvzy					

notes:

- ❖ construct shortest path tree by tracing predecessor nodes
- ❖ ties can exist (can be broken arbitrarily)




DIJKSTRA'S ALGORITHM: ANOTHER EXAMPLE



Step	N'	$D(v), p(v)$	$D(w), p(w)$	$D(x), p(x)$	$D(y), p(y)$	$D(z), p(z)$
0	u	2,u	5,u	1,u	∞	∞
1	ux	2,u	4,x		2,x	∞
2	uxy	2,u	3,y			4,y
3	uxyv		3,y			4,y
4	uxyvw					4,y
5	uxyvwz					

resulting shortest-path tree from u:



resulting forwarding table in u:

destination	link
v	(u,v)
x	(u,x)
y	(u,x)
w	(u,x)
z	(u,x)

* Check out the online interactive exercises for more examples: http://gaia.cs.umass.edu/kurose_ross/interactive/

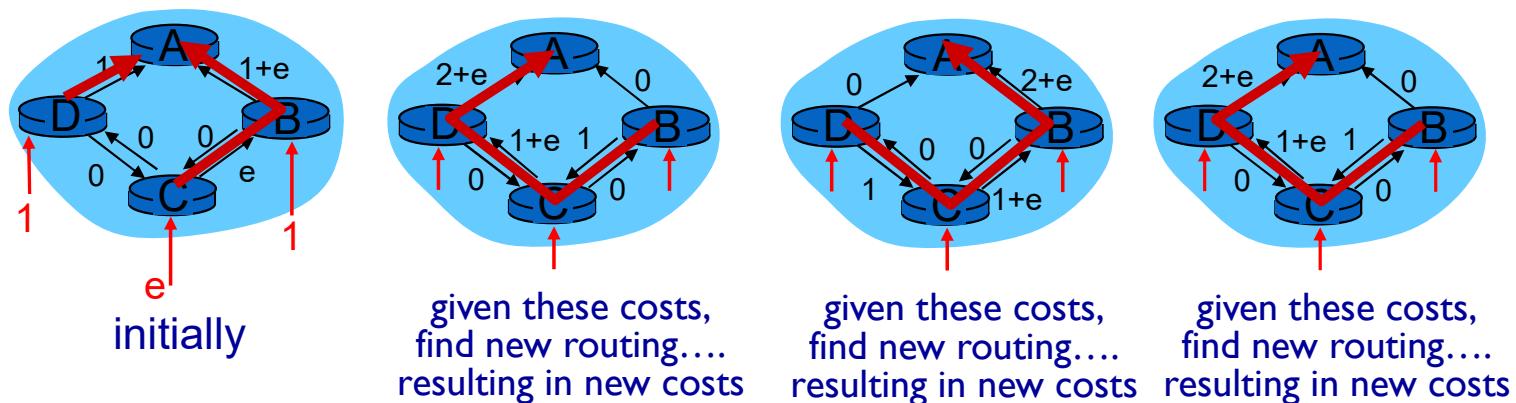
DIJKSTRA'S ALGORITHM, DISCUSSION

algorithm complexity: n nodes

- each iteration: need to check all nodes, w, not in N
- $n(n+1)/2$ comparisons: $O(n^2)$
- more efficient implementations possible: $O(n \log n)$

oscillations possible:



- e.g., support link cost equals amount of carried traffic:



LINK STATE ALGORITHM

Stage 1: Flooding:

- 1) Periodically distribute link-state advertisement (LSA) to neighbors
 - LSA contains cost (e.g., delay) to each neighbor
 - Link-state packet (LSP) – packet containing the LSA information.
- 2) Install received LSA in LS database
- 3) Re-distribute LSA to all neighbors
 - IF the LSP is the *most-recent* than the one stored (previously seen) THEN:
 - Flood LSP: Forward a copy on all links except on the link LSP was received.
 - Otherwise, discard LSP.



Stage 2: Path Computation

- 1) Use Dijkstra's shortest path algorithm to compute distances to all destinations
- 2) Install FIB <destination, nexthop> pair in the Forwarding Table

LINK STATE IN PRACTICE - OSPF

- OSPF (Open Shortest Path First Protocol) – *Intra AS Routing*
 - most important and most commonly used routing protocol on the Internet.
 - Also supports ***authentication, additional hierarchy, load balancing***
 - *OSPF Packet Structure:*

History and Evolution of OSPF

1989: RFC 1131 OSPFv1

1991: RFC1247 OSPFv2

CIDR based – classes routing


1994: RFC 1583 OSPFv2 (revised)

1997: RFC 2178 OSPFv2 (revised)

1998: **RFC 2328 OSPFv2** (current version for IPv4)


2008: RFC 5340 OSPF Version 3

support for IPv6



AuType	Description
0	Null authentication
1	Simple password
2	Cryptographic authentication
All others	Reserved for assignment by the IANA (iana@ISI.EDU)

OSPF MESSAGE STRUCTURE



OSPF packets are not carried as UDP or TCP payload!

OSPF has its own IP protocol number: **89**

TTL: set to 1 (in most cases)



Destination IP: neighbor's IP address or multicast address
for a broadcast environment:

224.0.0.5 (ALLSPFRouters) or 224.0.0.6 (AllIDRouters:
(designated and backup designated only)

32-bit LS sequence number field, does not wrap
On startup, router need not wait: can start with a lowest sequence number
LSP's compared on basis of sequence number
LSP's purged after about an hour
Synchronized expiration of LSPs
expired LSP reflooded with age zero

LS Age	Options	Type=1
Link state ID		
Advertising router		
LS sequence number		
LS checksum		Length

LINK STATE ALGORITHM –TYPICAL 2 STAGE WORKFLOW AT A ROUTER



Screenshot of NetworkMiner tool displaying an OSPF LS Update packet:

Frame 170: 90 bytes on wire (720 bits), 90 bytes captured (720 bits) on interface 0

Ethernet II, Src: c2:04:33:f4:00:00 (c2:04:33:f4:00:00), Dst: IPv4mcast_06 (01:00:5e:00:00:06)

Internet Protocol Version 4, Src: 192.168.10.25 (192.168.10.25), Dst: 224.0.0.6 (224.0.0.6)

Open Shortest Path First

LS Update Packet

Number of LSAs: 1

Summary-LSA (IP network)

- .000 0000 0000 0001 = LS Age (seconds): 1
- 0... = Do Not Age Flag: 0

Options: 0x22 (DC, E)

LS Type: Summary-LSA (IP network) (3)

Link State ID: 10.55.4.1 (10.55.4.1)

Advertising Router: 192.168.10.25 (192.168.10.25)

Sequence Number: 0x80000001

Csum: 0xbff9f

Length: 28

Netmask: 255.255.255.255 (255.255.255.255)

Metric: 11

Firewall.cx

PROBLEM: ROUTER FAILURE

- A failed router then comes up but does not remember the last sequence number it used before it crashed.
- New LSPs by this router may be ignored by others if they have sent with a lower sequence number

ONE SOLUTION: LSP AGING

- Nodes periodically decrement age (TTL) of stored LSPs
- LSPs expire when TTL reaches 0
 - LSP is re-flooded once TTL = 0
- Alternative for a rebooted router is to wait until all LSPs have expired.
- Trade-off between frequency of LSPs and router wait after reboot

HIERARCHICAL OSPF

- two-level hierarchy: local area, backbone.
 - link-state advertisements flooded only in area, or backbone
 - each node has detailed area topology; only knows direction to reach other destinations


area border routers:
“summarize” distances to
destinations in own area,
advertise in backbone

local routers:

- flood LS in area only
- compute routing within area
- forward packets to outside via area border router

boundary router:
connects to other ASes

backbone router:
runs OSPF limited to backbone



ANOTHER/~ARCHAIC APPROACH TO ROUTING

DISTANCE VECTOR –

I CAN KNOW THE DISTANCE(S) FROM MY NEIGHBOR
ONCE I KNOW THE DISTANCES, I CAN KNOW THE NETWORK


VS.

LINK STATE –

I CAN KNOW THE NETWORK TOPOLOGY FROM MY NEIGHBOR
ONCE I KNOW THE NETWORK, I CAN COMPUTE THE DISTANCES

DISTANCE VECTOR ALGORITHM:

each node:



iterative, asynchronous: each local iteration caused by:

- local link cost change
- DV update message from neighbor

distributed, self-stopping: each node notifies neighbors *only* when its DV changes

- neighbors then notify their neighbors – *only if necessary*
- no notification received, no actions taken!

DISTANCE VECTOR ALGORITHM

Bellman-Ford equation (dynamic programming)

Let $d_x(y) := \text{cost of least-cost path from } x \text{ to } y$


then

$$d_x(y) = \min_v \{ c(x,v) + d_v(y) \}$$

cost to neighbor v

min taken over all neighbors v of x

cost from neighbor v to destination y



clearly, $d_v(z) = 5$, $d_x(z) = 3$, $d_w(z) = 3$

- node x:

- x maintains distance vector $\mathbf{D}_x = [d_x(y): y \in N]$
- knows cost to each neighbor v: $c(x,v)$
- maintains its neighbors' distance vectors.
- For each neighbor v, x maintains $\mathbf{D}_v = [d_v(y): y \in N]$

B-F equation says:

$$\begin{aligned} d_u(z) &= \min \{ c(u,v) + d_v(z), \\ &\quad c(u,x) + d_x(z), \\ &\quad c(u,w) + d_w(z) \} \\ &= \min \{ 2 + 5, \\ &\quad \textcircled{1} + 3, \\ &\quad 5 + 3 \} = 4 \end{aligned}$$


DISTANCE VECTOR: EXAMPLE



$t=0$

- All nodes have distance estimates to nearest neighbors (only)
- All nodes send their local distance vector to their neighbors

DV in a:
$D_a(a)=0$
$D_a(b) = 8$
$D_a(c) = \infty$
$D_a(d) = 1$
$D_a(e) = \infty$
$D_a(f) = \infty$
$D_a(g) = \infty$
$D_a(h) = \infty$
$D_a(i) = \infty$



- A few asymmetries:
- missing link
 - larger cost


DISTANCE VECTOR EXAMPLE: ITERATION



t=1

All nodes:

- receive distance vectors from neighbors
- compute their new local distance vector
- send their new local distance vector to neighbors




DISTANCE VECTOR EXAMPLE: ITERATION



t=1

All nodes:

- receive distance vectors from neighbors
- compute their new local distance vector
- send their new local distance vector to neighbors




DISTANCE VECTOR EXAMPLE: ITERATION



t=1

All nodes:

- receive distance vectors from neighbors
- compute their new local distance vector
- send their new local distance vector to neighbors




DISTANCE VECTOR EXAMPLE: ITERATION



t=2

All nodes:

- receive distance vectors from neighbors
- compute their new local distance vector
- send their new local distance vector to neighbors




DISTANCE VECTOR EXAMPLE: ITERATION



t=2

All nodes:

- receive distance vectors from neighbors
- compute their new local distance vector
- send their new local distance vector to neighbors




DISTANCE VECTOR EXAMPLE: ITERATION



t=2

All nodes:

- receive distance vectors from neighbors
- compute their new local distance vector
- send their new local distance vector to neighbors



DISTANCE VECTOR EXAMPLE: ITERATION

.... and so on


Let's next take a look at the iterative *computations* at nodes

DISTANCE VECTOR EXAMPLE: COMPUTATION



t=1

- b receives DVs from a, c, e



DISTANCE VECTOR EXAMPLE: COMPUTATION




$t=1$

- b receives DVs from a, c, e, computes:

$$\begin{aligned}
 D_b(a) &= \min\{c_{b,a} + D_a(a), c_{b,c} + D_c(a), c_{b,e} + D_e(a)\} = \min\{8, \infty, \infty\} = 8 \\
 D_b(c) &= \min\{c_{b,a} + D_a(c), c_{b,c} + D_c(c), c_{b,e} + D_e(c)\} = \min\{\infty, 1, \infty\} = 1 \\
 D_b(d) &= \min\{c_{b,a} + D_a(d), c_{b,c} + D_c(d), c_{b,e} + D_e(d)\} = \min\{9, 2, \infty\} = 2 \\
 D_b(e) &= \min\{c_{b,a} + D_a(e), c_{b,c} + D_c(e), c_{b,e} + D_e(e)\} = \min\{\infty, \infty, 1\} = 1 \\
 D_b(f) &= \min\{c_{b,a} + D_a(f), c_{b,c} + D_c(f), c_{b,e} + D_e(f)\} = \min\{\infty, \infty, 2\} = 2 \\
 D_b(g) &= \min\{c_{b,a} + D_a(g), c_{b,c} + D_c(g), c_{b,e} + D_e(g)\} = \min\{\infty, \infty, \infty\} = \infty \\
 D_b(h) &= \min\{c_{b,a} + D_a(h), c_{b,c} + D_c(h), c_{b,e} + D_e(h)\} = \min\{\infty, \infty, 2\} = 2 \\
 D_b(i) &= \min\{c_{b,a} + D_a(i), c_{b,c} + D_c(i), c_{b,e} + D_e(i)\} = \min\{\infty, \infty, \infty\} = \infty
 \end{aligned}$$


DV in a:
$D_a(a)=0$
$D_a(b) = 8$
$D_a(c) = \infty$
$D_a(d) = 1$
$D_a(e) = \infty$
$D_a(f) = \infty$
$D_a(g) = \infty$
$D_a(h) = \infty$
$D_a(i) = \infty$




8

DV in b:
$D_b(a) = 8$
$D_b(f) = \infty$
$D_b(c) = 1$
$D_b(g) = \infty$
$D_b(d) = \infty$
$D_b(h) = \infty$
$D_b(e) = 1$
$D_b(i) = \infty$

DV in c:
$D_c(a) = \infty$
$D_c(b) = 1$
$D_c(c) = 0$
$D_c(d) = \infty$
$D_c(e) = \infty$
$D_c(f) = \infty$
$D_c(g) = \infty$
$D_c(h) = \infty$
$D_c(i) = \infty$



1



1

DV in b:
$D_b(a) = 8$
$D_b(f) = 2$
$D_b(c) = 1$
$D_b(g) = \infty$
$D_b(d) = 2$
$D_b(h) = 2$
$D_b(e) = 1$
$D_b(i) = \infty$


DV in e:
$D_e(a) = \infty$
$D_e(b) = 1$
$D_e(c) = \infty$
$D_e(d) = 1$
$D_e(e) = 0$
$D_e(f) = 1$
$D_e(g) = \infty$
$D_e(h) = 1$
$D_e(i) = \infty$

DISTANCE VECTOR EXAMPLE: COMPUTATION



t=1

- c receives DVs from b




DISTANCE VECTOR EXAMPLE: COMPUTATION



$t=1$

- c receives DVs from b computes:

$$\begin{aligned}
 D_c(a) &= \min\{c_{c,b} + D_b(a)\} = 1 + 8 = 9 \\
 D_c(b) &= \min\{c_{c,b} + D_b(b)\} = 1 + 0 = 1 \\
 D_c(d) &= \min\{c_{c,b} + D_b(d)\} = 1 + \infty = \infty \\
 D_c(e) &= \min\{c_{c,b} + D_b(e)\} = 1 + 1 = 2 \\
 D_c(f) &= \min\{c_{c,b} + D_b(f)\} = 1 + \infty = \infty \\
 D_c(g) &= \min\{c_{c,b} + D_b(g)\} = 1 + \infty = \infty \\
 D_c(h) &= \min\{c_{c,b} + D_b(h)\} = 1 + \infty = \infty \\
 D_c(i) &= \min\{c_{c,b} + D_b(i)\} = 1 + \infty = \infty
 \end{aligned}$$



DV in c:	
$D_c(a) = 9$	
$D_c(b) = 1$	
$D_c(c) = 0$	
$D_c(d) = 2$	
$D_c(e) = \infty$	
$D_c(f) = \infty$	
$D_c(g) = \infty$	
$D_c(h) = \infty$	
$D_c(i) = \infty$	


* Check out the online interactive exercises for more examples:
http://gaia.cs.umass.edu/kurose_ross/interactive/

DISTANCE VECTOR EXAMPLE: COMPUTATION



t=1


- e receives DVs from b, d, f, h




DISTANCE VECTOR: STATE INFORMATION DIFFUSION

Iterative communication, computation steps diffuses information through network:

-  t=0 c's state at t=0 is at c only
-  t=1 c's state at t=0 has propagated to b, and may influence distance vector computations up to **1** hop away, i.e., at b
-  t=2 c's state at t=0 may now influence distance vector computations up to **2** hops away, i.e., at b and now at a, e as well
-  t=3 c's state at t=0 may influence distance vector computations up to **3** hops away, i.e., at b,a,e and now at c,f,h as well
-  t=4 c's state at t=0 may influence distance vector computations up to **4** hops away, i.e., at b,a,e, c, f, h and now at g,i as well



DISTANCE VECTOR



$$D_x(y) = \min\{c(x,y) + D_y(y), c(x,z) + D_z(y)\}$$

$$= \min\{2+0, 7+1\} = 2$$

$$D_x(z) = \min\{c(x,y) + D_y(z), c(x,z) + D_z(z)\}$$

$$= \min\{2+1, 7+0\} = 3$$

node x table

from	x	y	z
x	0	2	7
y	∞	∞	∞
z	∞	∞	∞

node y table

from	x	y	z
x	∞	∞	∞
y	2	0	1
z	∞	∞	∞

node z table

from	x	y	z
x	∞	∞	∞
y	∞	∞	∞
z	7	1	0

from	x	y	z
x	0	2	3
y	2	0	1
z	7	1	0


from	x	y	z
x	0	2	7
y	2	0	1
z	7	1	0

from	x	y	z
x	0	2	7
y	2	0	1
z	3	1	0

from	x	y	z
x	0	2	3
y	2	0	1
z	3	1	0

from	x	y	z
x	0	2	3
y	2	0	1
z	3	1	0

from	x	y	z
x	0	2	3
y	2	0	1
z	3	1	0




$$D_z(x) = \min\{c(z,x) + D_x(x), c(z,y) + D_y(x)\}$$

$$= \min\{7+0, 1+2\} = 3$$

DISTANCE VECTOR: LINK COST CHANGES

link cost changes:

- node detects local link cost change
- updates routing info, recalculates local DV
- if DV changes, notify neighbors



“good news travels fast”

t_0 : y detects link-cost change, updates its DV, informs its neighbors.

t_1 : z receives update from y, updates its table, computes new least cost to x , sends its neighbors its DV.


t_2 : y receives z's update, updates its distance table. y's least costs do *not* change, so y does *not* send a message to z.

DISTANCE VECTOR: LINK COST CHANGES

link cost changes:

- node detects local link cost change and notify neighbors


“bad news travels slow” – count-to-infinity problem:



- y sees direct link to x has new cost 60, but z has said it has a path at cost of 5. So y computes “my new cost to x will be 6, via z); notifies z of new cost of 6 to x.
 - z learns that path to x via y has new cost 6, so z computes “my new cost to x will be 7 via y), notifies y of new cost of 7 to x.
 - y learns that path to x via z has new cost 7, so y computes “my new cost to x will be 8 via y), notifies z of new cost of 8 to x.
 - z learns that path to x via y has new cost 8, so z computes “my new cost to x will be 9 via y), notifies y of new cost of 9 to x.
- ...

DISTANCE VECTOR: HOW TO OVERCOME COUNT TO INFINITY!

Different Approaches have been proposed to avoid/prevent Routing Loops:



- ❖ **Split Horizon:** (Kind of Avoidance/Passive Prevention)
 - ❖ ‘Never advertise a route out of the interface through which you learned it’
 - ❖ Prevent a router from advertising the route back on the same interface from which it learns the route.
- ❖ **Route Poisoning:** (– almost the opposite of split horizon) Upon Failure detection, poison the route by advertising infinity metric to the neighbors.
 - ❖ Until the Route is activated – apply split horizon
 - ❖ When route fails, route poisoning overrules the split horizon.
- ❖ **Holddown Timer:** Upon ‘Route Poison’ updates disregard any other routing updates until the holddown time expires (typically 180 seconds)
- ❖ **Poison Reverse** (Kind of Active Prevention)
 - ❖ ‘Once you learn a route through an interface advertise it as unreachable back through the same interface’

Split Horizon with poisoned reverse:

Try to Read [RFC 1058](#)

Q: will this completely solve count to infinity problem?

COMPARISON OF LS AND DV ALGORITHMS

message complexity

- **LS:** with n nodes, E links, $O(nE)$ msgs sent
- **DV:** exchange between neighbors only
 - convergence time varies

Message Update interval

- **LS:** large duration (~30 mins) or on change
- **DV:** small time interval (~1 min) or on change.

speed of convergence

- **LS:** $O(n^2)$ algorithm requires $O(nE)$ msgs
 - may have oscillations
- **DV:** convergence time varies
 - may be routing loops
 - count-to-infinity problem

robustness: what happens if router malfunctions?

LS:

- node can advertise incorrect *link* cost
- each node computes only its own table

DV:

- DV node can advertise incorrect *path* cost
- each node's table used by others
 - error propagate through the network

CPU and Bandwidth tax:

LS:


- Flooding small packets
- each node computes only its own table

DV:

- DV node can advertise incorrect *path* cost
- each node's table used by others
 - error propagate through the network

PROBLEM

5. In the following network topology consider that it takes 1ms to transmit a message on each link. Consider running the Djkstra or Bellman-Ford algorithm to find the shortest path for the topology takes negligible time i.e. 0ms.(10 pts)



- What is the convergence time needed for LSR and DVR respectively? (2 pts)
- A total of how many packets are exchanged for route convergence in LSR and DVR respectively? Note: A total includes all the packets that are broadcast/flood in the network by different routers; (I.e. A packet is considered per hop) (4 pts)
- If the cost of link B-C changes from 1 to 5. How quickly can the LS and DVR converge again? (4 pts)


LSR

Node	Round T0 (# new msg x # links to send)	Round T1 (1 hop away)	Round T2 (2 hop away)
A	$1 \times 2 = 2$ [(B,C)]	$2 \times 1 = 2$ [B->C, C->B]	$2 \times 1 = 2$ (From D and E)
B	$1 \times 4 = 4$ (A,C,D,E)	$4 \times 3 = 12$ (From A,C,D,E to every other links)	---
C	$1 \times 3 = 3$ (A,B, E)	$3 \times 2 = 6$	----
D	$1 \times 2 = 2$ (B,E)	$2 \times 1 = 2$	$1 \times 1 = 1$ (From A)
E	$1 \times 3 = 3$ (C,D,E)	$3 \times 2 = 6$	$1 \times 2 = 2$ (From A)
Total Msgs	14	28	5

DVR

Node	Round T0 (# new msg x # links to send)	Round T1 (All nodes have updates due to T0)	Round T2 (Only A,D will have updates)
A	$1 \times 2 = 2$	$1 \times 2 = 2$	$1 \times 2 = 2$
B	$1 \times 4 = 4$	$1 \times 4 = 4$	----
C	$1 \times 3 = 3$	$1 \times 3 = 3$	----
D	$1 \times 2 = 2$	$1 \times 2 = 2$	$1 \times 2 = 2$
E	$1 \times 3 = 3$	$1 \times 3 = 3$	----
Total Msgs	14	14	4

MAKING ROUTING SCALABLE – MEET INTERNET DEMANDS



scale: with billions of destinations:


- can't store all destinations in routing tables!
- routing table exchange would swamp links!

administrative autonomy

- **internet = network of networks**
- each network admin may want to control routing in its own network.

INTERNET APPROACH CONNECTED AS'ES:: INTERCONNECTED AS'ES

aggregate routers into regions known as “autonomous systems” (AS) (a.k.a. “domains”)



- forwarding table configured by both intra- and inter-AS routing algorithm
 - intra-AS routing determine entries for destinations within AS
 - inter-AS & intra-AS determine entries for external destinations

intra-AS routing


- routing among hosts, routers within the same AS (“network”)
- all routers in AS must run *same* intra-domain routing protocol
- routers in *different* AS can run *different* intra-domain routing protocol
- Gateway/Edge router in AS: has link(s) to router(s) in other AS'es

inter-AS routing

- routing among the AS'es
- gateways perform inter-domain routing (as well as intra-domain routing)

INTER-AS TASKS

- How should any router in AS2 forward datagram destined outside of AS2?



Inter-AS routing:

1. learn about destinations reachable through other As'es (AS1 & AS3)
2. propagate this information to all routers within AS2.

- *Intra-AS routing a.k.a interior gateway protocols (IGP):* most common protocols:

- RIP: Routing Information Protocol (now mostly RIPv2) – *DVR*
- OSPF: Open Shortest Path First -- *LSR*
- IS-IS: Intermediate System routing protocol essentially same as OSPF -- *LSR*
- IGRP: Interior Gateway Routing Protocol (Cisco proprietary *until 2016*) – *DVR*


- *Inter-AS routing a.k.a exterior gateway protocols:* most common protocols:

- BGP: Border Gateway Protocol – *PVR*
- EGP: Exterior Gateway Protocol (way back in 1982, now obsolete) – *??*

INTER-AS ROUTING: BORDER GATEWAY PROTOCOL (BGP)


BGP: *the de facto inter-domain routing protocol*

- “glue that holds the ***Internet*** together”
 - allows subnet to advertise itself, and the destinations it can reach, to rest of Internet: “*I am here, here is who I can reach, and how*”.
- BGP provides each AS a means to:
- **eBGP:** obtain subnet reachability information from neighboring AS'es
 - **iBGP:** propagate reachability information to all AS-internal routers.
 - determine “good” routes to other networks based on reachability information and ***policy***



BGP OUTLINE OF KEY OPERATIONS

- **BGP session:** BGP routers (“peers”) exchange BGP messages over a TCP connection: **port 179**
 - advertise **paths** to different destination network prefixes (BGP is a “path vector” protocol)
- Using **eBGP**, AS3 gateway router 3a advertises path **AS3,X** to AS2 gateway router 2c:
 - AS3 **promises** to AS2 it will forward datagrams towards X
- Using **iBGP**, AS2 router 2c advertises path **AS3,X** to all of the AS2 routers.




- BGP advertised route: **prefix + attributes**
 - **prefix**: destination being advertised
 - two important **attributes**:
 - **AS-PATH**: list of ASes through which prefix advertisement has passed
 - **NEXT-HOP**: indicates specific internal-AS router to the next-hop AS

Policy-based routing:

gateway receiving route advertisement uses **import policy** to accept/decline path (e.g., never route through AS2). AS policy also determines whether to **advertise** path to other neighboring ASes or not!


BGP, OSPF, FORWARDING TABLE ENTRIES

Q: how does router set forwarding table entry to distant prefix?



- recall: 1a, 1b, 1d learn via iBGP from 1c: “path to X goes through 1c”
- at 1d: OSPF intra-domain routing: to get to 1c, use interface 1
- at 1d: to get to X, use interface 1
- at 1a: OSPF intra-domain routing: to get to 1c, use interface 2
- at 1a: to get to X, use interface 2


HOT POTATO ROUTING



- 2d learns (via iBGP) it can route to X via 2a or 2c
- *hot potato routing*: choose local gateway that has least intra-domain cost
e.g., 2d chooses 2a, even though more AS hops to X): don't worry about inter-domain cost!

BGP: ACHIEVING POLICY VIA ADVERTISEMENTS

Q: Suppose an ISP does not want to carry transit traffic between other ISPs?



- A,B,C are *provider networks*
- X,W,Y are customer (of provider networks)
- X is *dual-homed*: attached to two networks
- *policy to enforce*: X does not want to route from B to C via X
 - Solution: X will not advertise to B a route to C
- *policy to enforce*: B does not want to route from C to A via B
 - A advertises path Aw to B and to C
- **Solution: B *chooses not to advertise* BAw to C:**
 - C does not learn about CBAw path
 - C will route CAw (not using B) to get to w

SUMMARY: WHY DIFFERENT INTRA-AS AND INTER-AS ROUTING?

policy:

- inter-AS: admin wants control over how its traffic is routed, who routes through its network, etc.
- intra-AS: single admin, so no policy decisions needed

scale:

- hierarchical routing saves table size, reduced update traffic

performance:

- intra-AS: can focus on performance.
- inter-AS: policy may dominate over performance.

BGP MESSAGES

- BGP messages exchanged between peers over TCP connection; using TCP port 179.
- BGP messages:
 - **OPEN**: opens TCP connection to remote BGP peer and authenticates sending BGP peer
 - **UPDATE**: advertises new path (or withdraws old)
 - **KEEPALIVE**: keeps connection alive in absence of UPDATES; also ACKs OPEN request
 - **NOTIFICATION**: reports errors in previous msg; also used to close connection

BGP ROUTE SELECTION

- router may learn about more than one route to destination AS, selects route based on:
 1. local preference value attribute: policy decision
 2. shortest AS-PATH
 3. closest NEXT-HOP router: hot potato routing
 4. additional criteria