

CS 331: Computer Networks

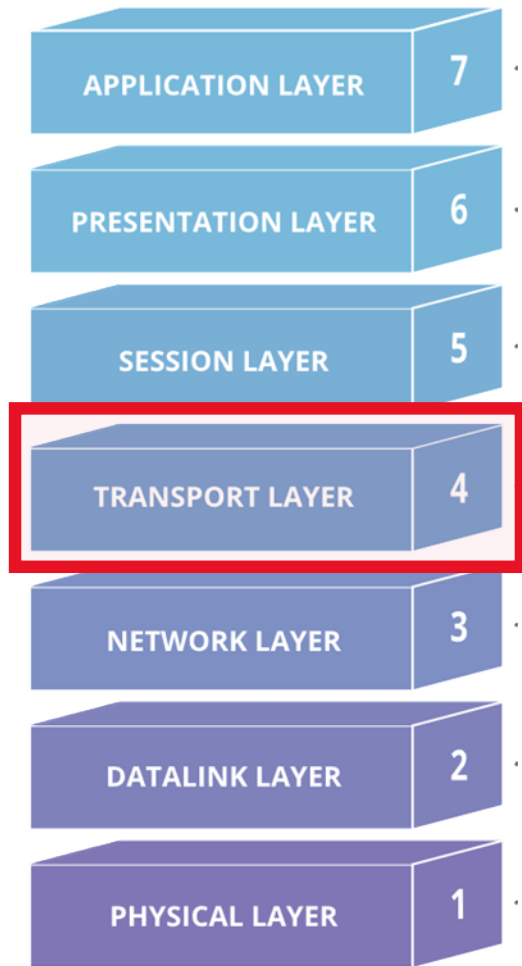
Sameer G Kulkarni

Assistant Professor,

*Department of Computer Science and Engineering,
Indian Institute of Technology, Gandhinagar*



TRANSPORT LAYER



A note on the use of these ppt slides:

We're making these slides freely available to all (faculty, students, readers). They're in PowerPoint form so you see the animations; and can add, modify, and delete slides (including this one) and slide content to suit your needs. They obviously represent a *lot* of work on our part. In return for use, we only ask the following:


- ❖ If you use these slides (e.g., in a class) that you mention their source (after all, we'd like people to use our book!)
- ❖ If you post any slides on a www site, that you note that they are adapted from (or perhaps identical to) our slides, and note our copyright of this material.

Thanks and enjoy! JFK/KWR ©

All material copyright 1996-2019
J.F Kurose and K.W. Ross, All Rights Reserved

Few Slides are borrowed from Prof. Rajat Chakraborty and Prof. Sandip Chakraborty, Computer Science Department, IIT Kharagpur

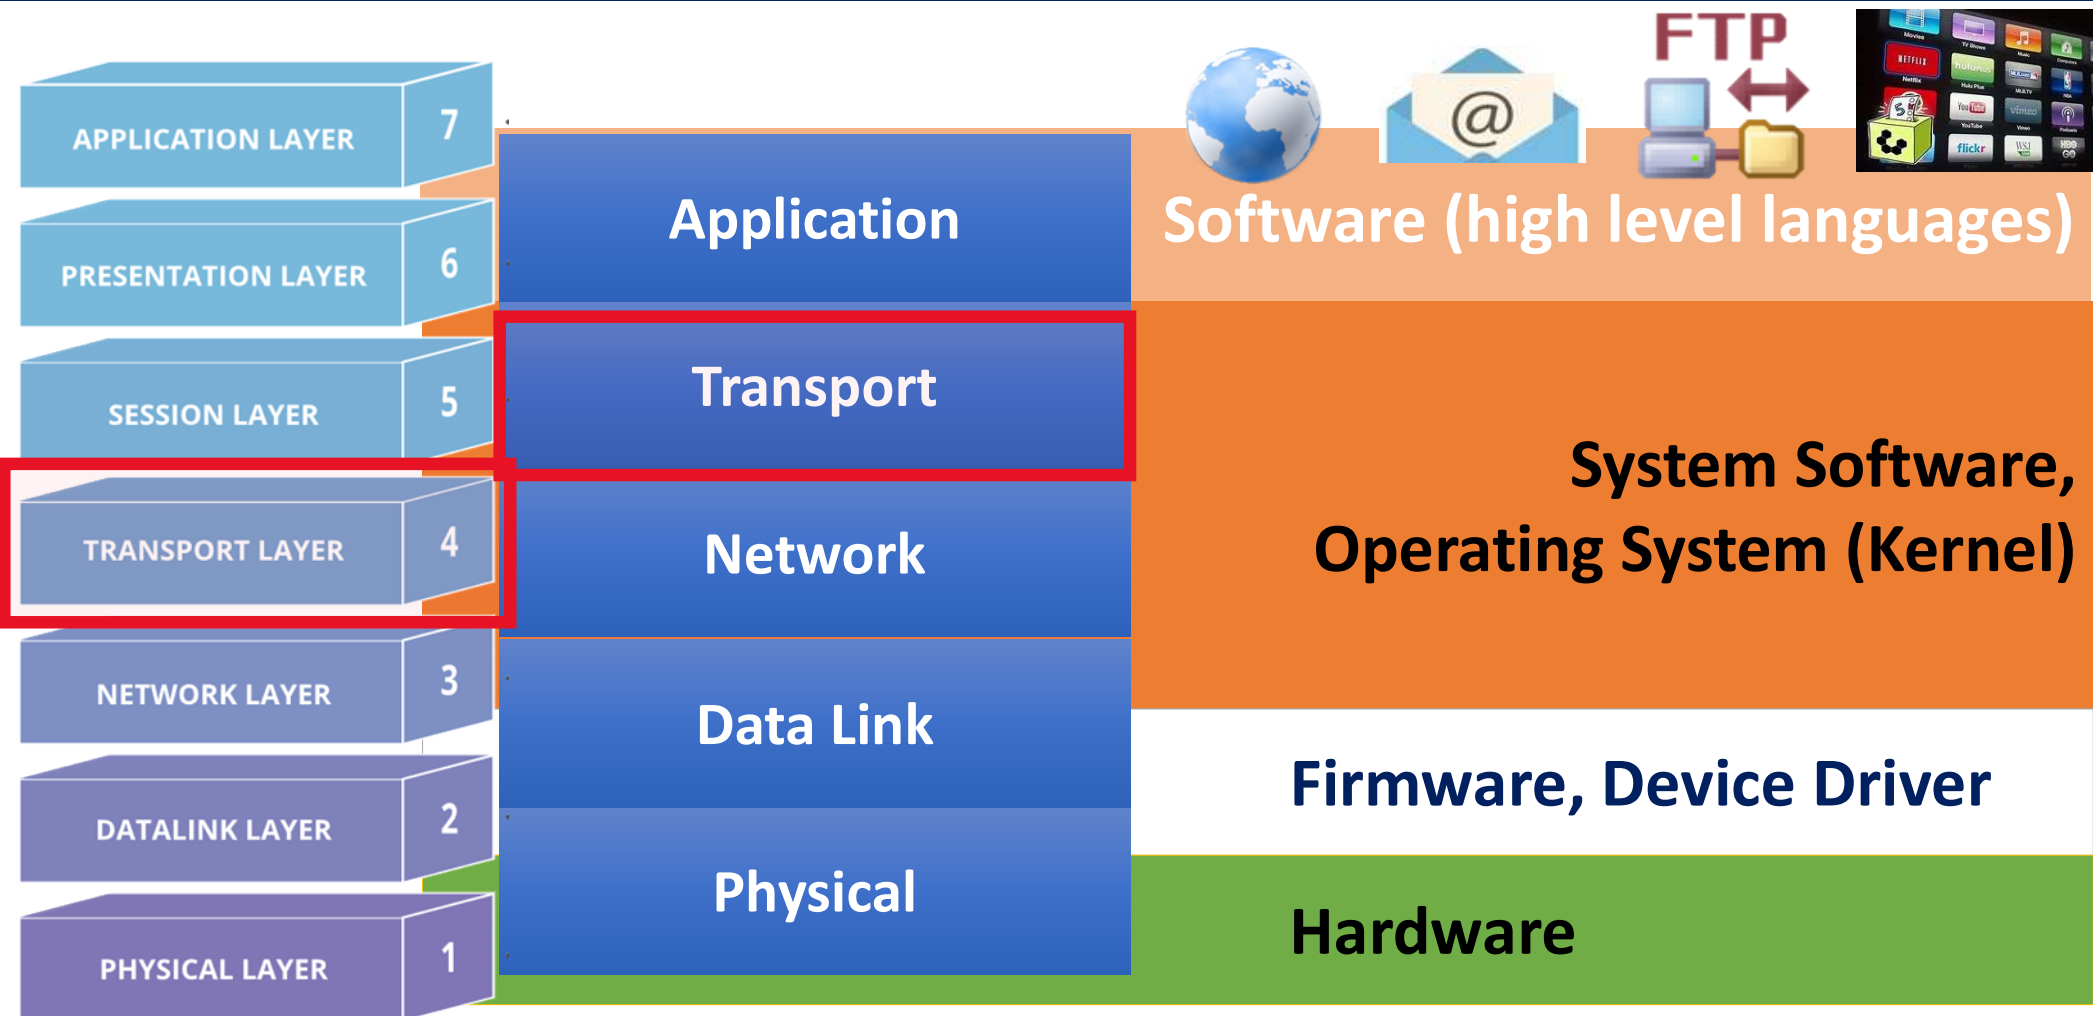
Transport Layer: Study of Protocols: TCP, UDP. Reliable Data transfer, Flow control, congestion control. TCP variants.



Chapter 3

*Computer
Networking: A Top
Down Approach*
8th edition
Jim Kurose, Keith Ross
Addison-Wesley

PROTOCOL STACK IMPLEMENTATION IN A HOST



TRANSPORT LAYER SERVICES

End to end
packet delivery

UDP

Reliable and
Ordered Data
Delivery

Connection
Establishment

Flow Control

Congestion
Control

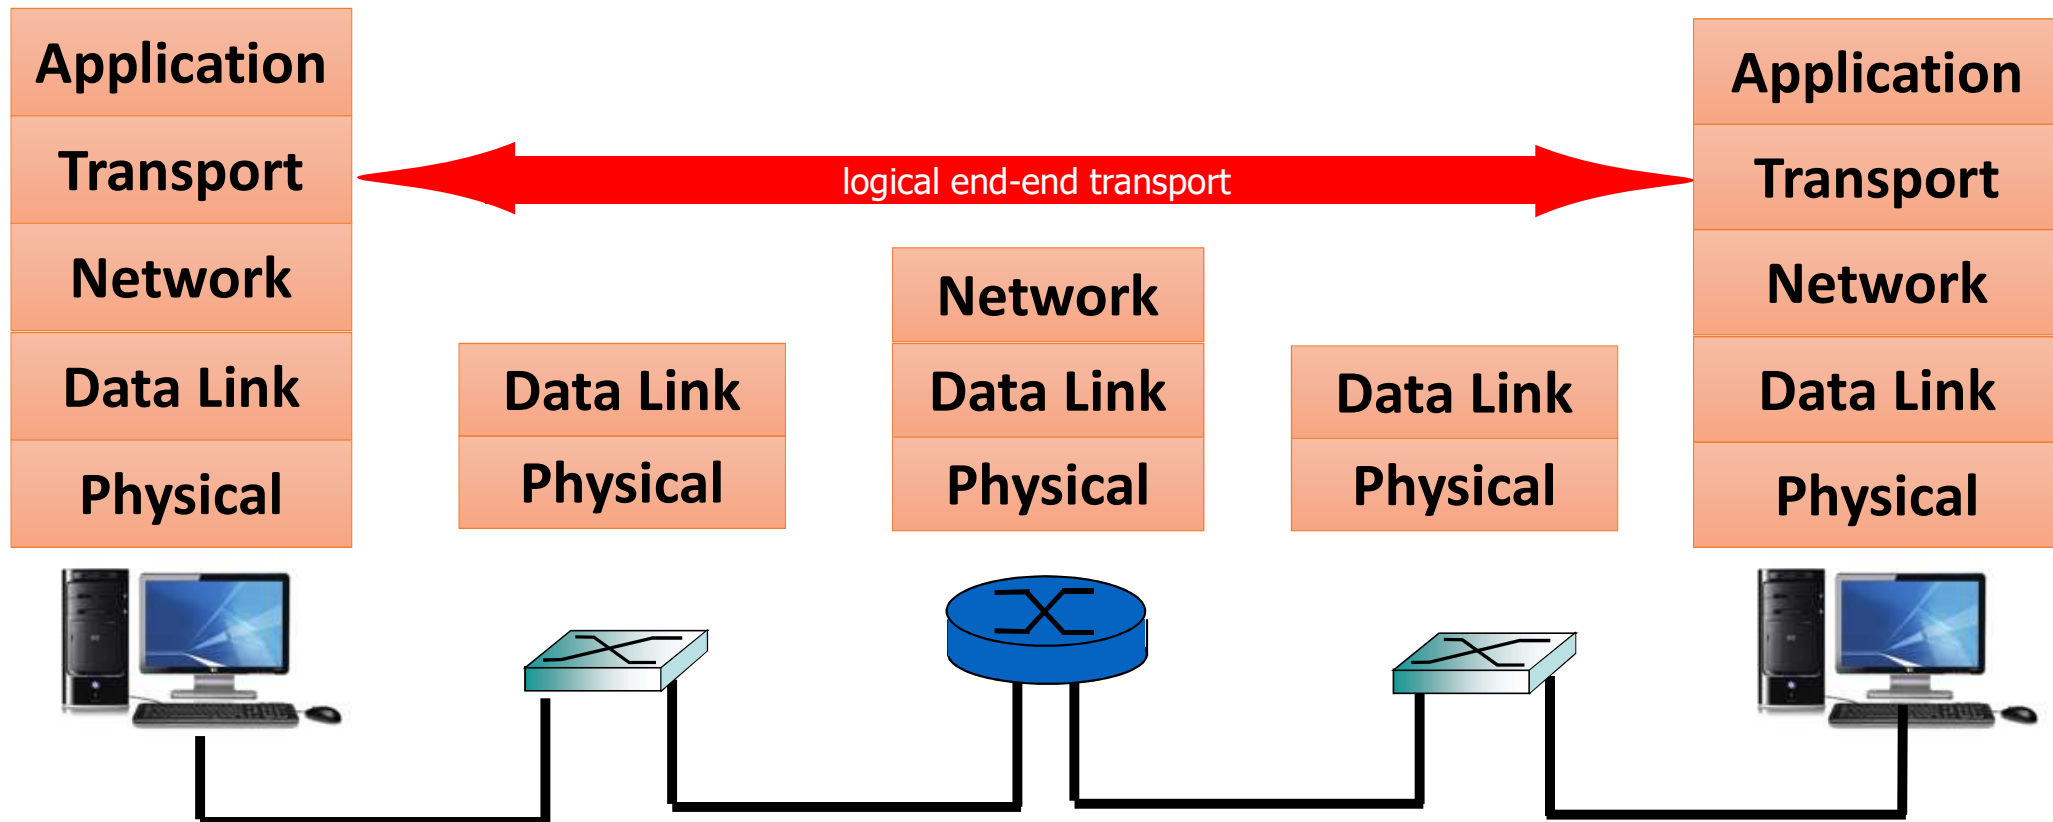
TCP

Transport Layer

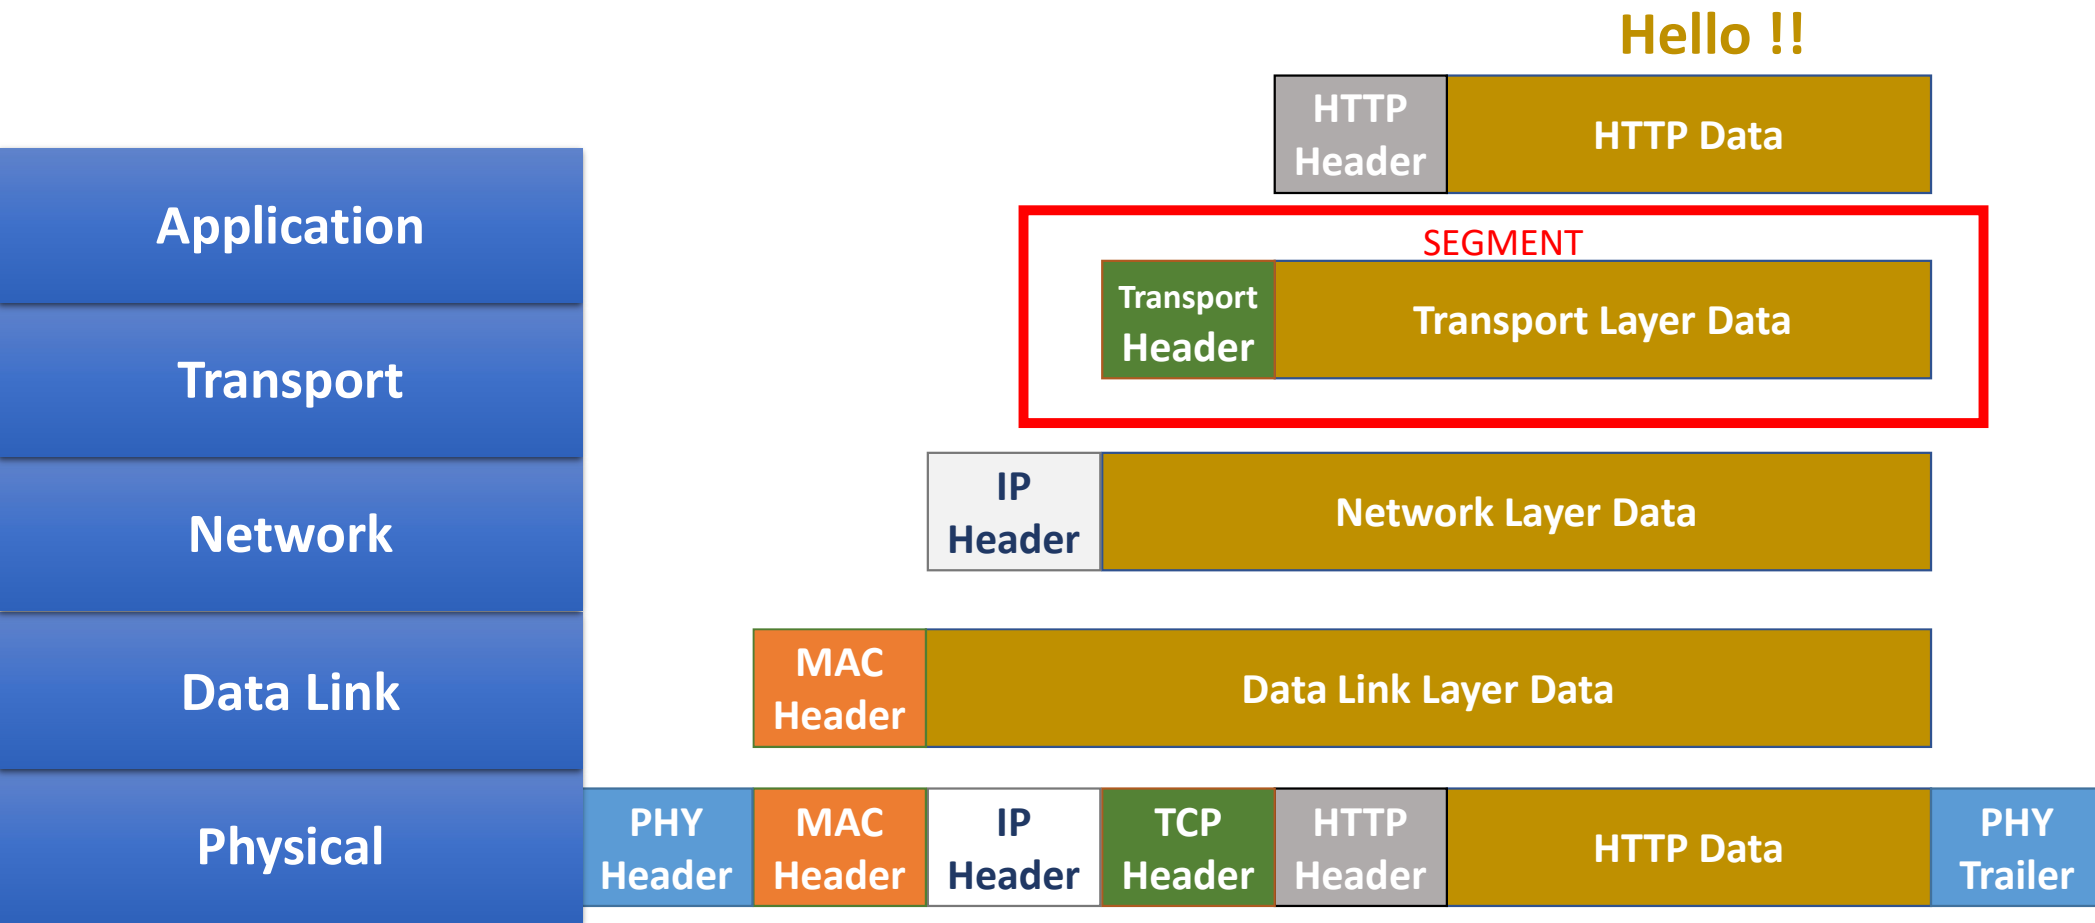
IP Layer

Packet delivery (unreliable)

DATA TRANSFER BETWEEN TWO REMOTE MACHINES

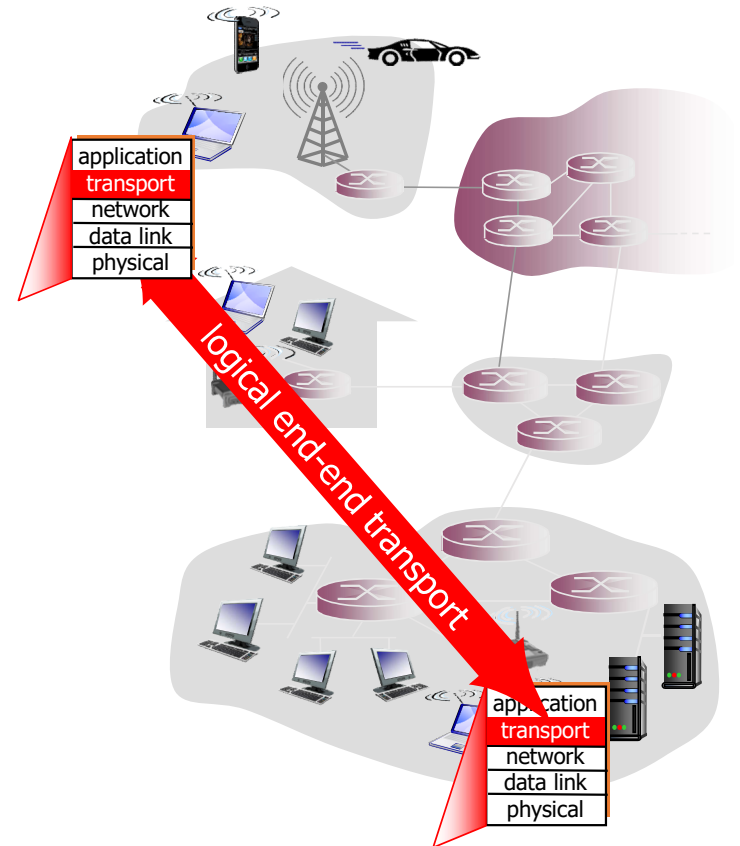


HOW APPLICATION DATA PASSES THROUGH DIFFERENT LAYERS



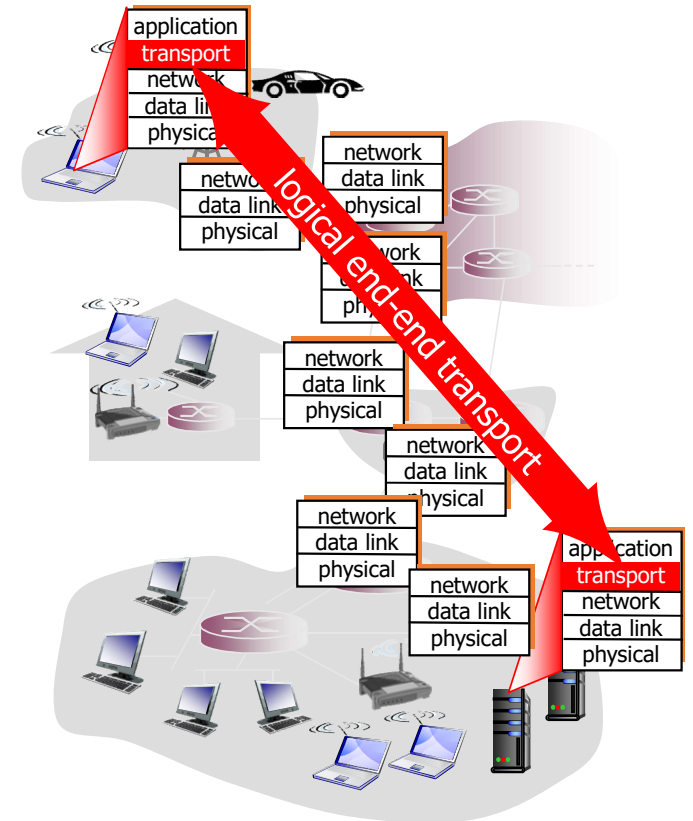
TRANSPORT SERVICES AND PROTOCOLS

- ❖ provide *logical communication* between app processes running on different hosts
- ❖ transport protocols run in end systems
 - send side: breaks app messages into *segments*, passes to network layer
 - rcv side: reassembles segments into messages, passes to app layer
- ❖ more than one transport protocol available to apps
 - Internet: TCP and UDP

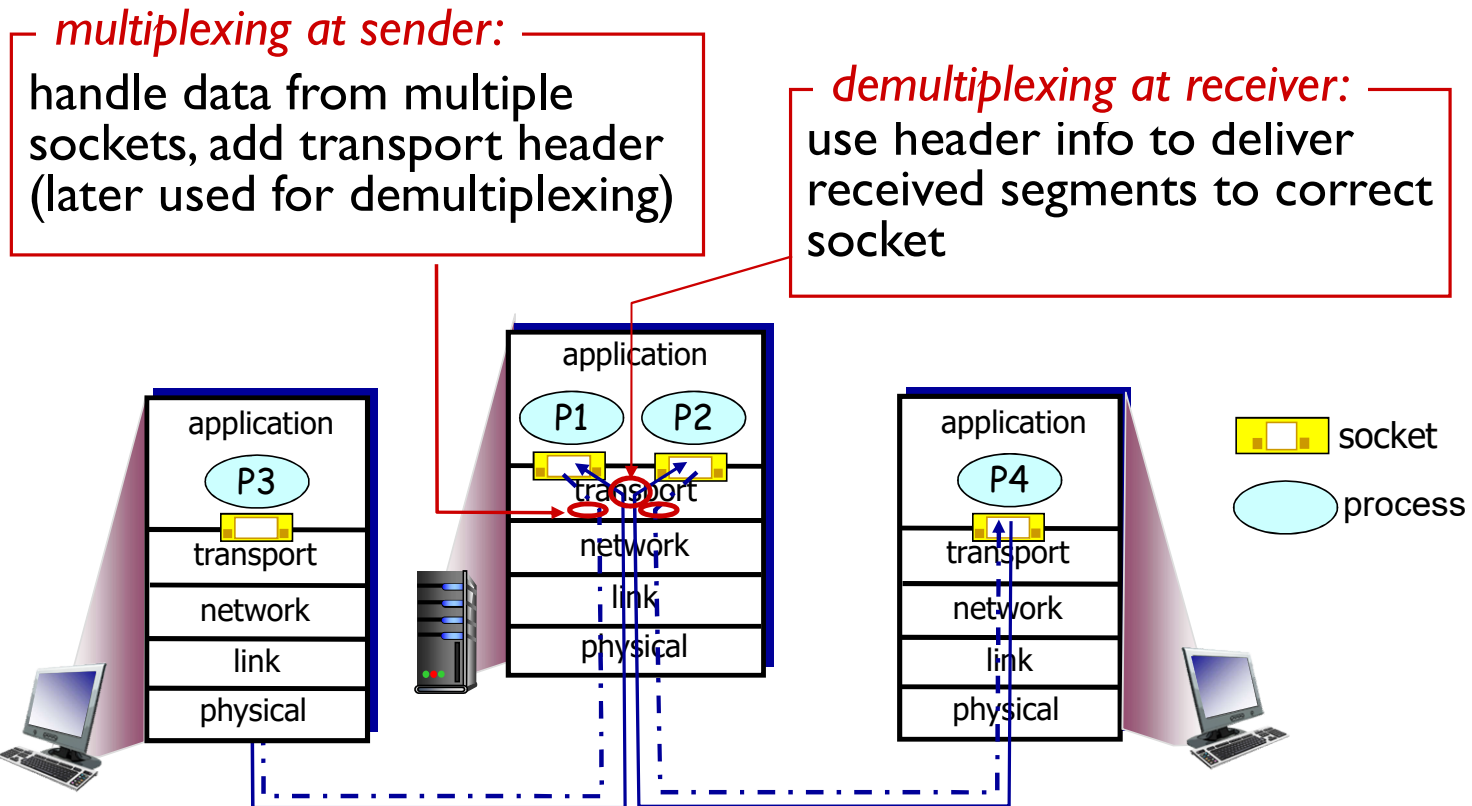


TRANSPORT-LAYER PROTOCOLS

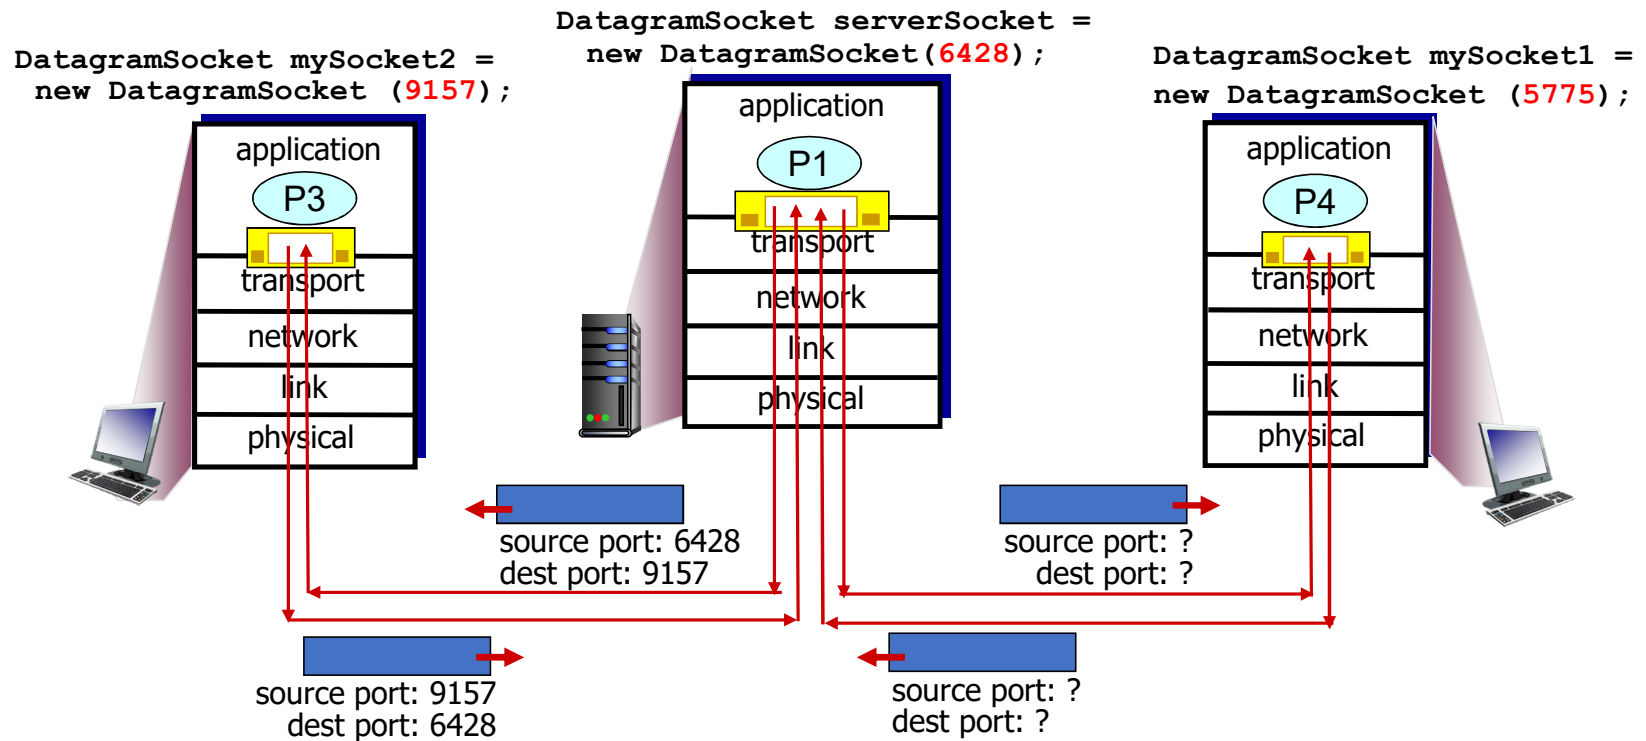
- unreliable, unordered delivery: UDP
 - no-frills extension of “best-effort” IP
- reliable, in-order delivery (TCP)
 - connection setup
 - congestion control
 - flow control
- services not available:
 - delay guarantees
 - bandwidth guarantees



MULTIPLEXING/DEMULTIPLEXING



CONNECTIONLESS DEMUX: EXAMPLE

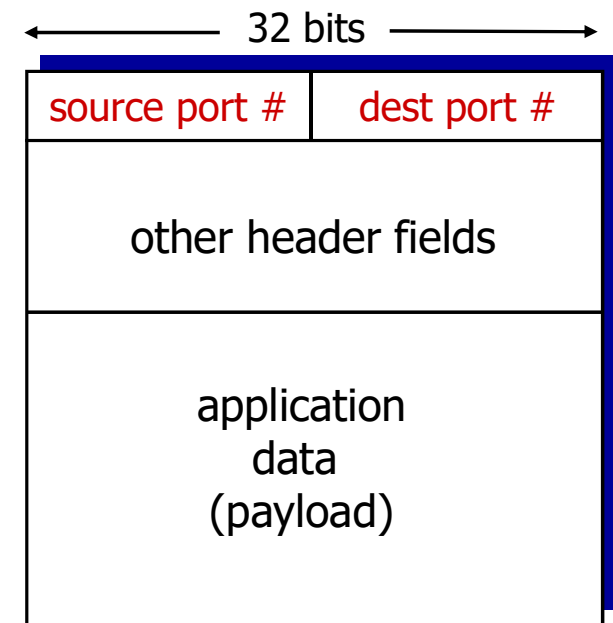


HOW DEMULTIPLEXING WORKS

❖ host receives IP datagrams

- each datagram has source IP address, destination IP address
- each datagram carries one transport-layer segment
- each segment has source, destination port number

❖ host uses *IP addresses & port numbers* to direct segment to appropriate socket

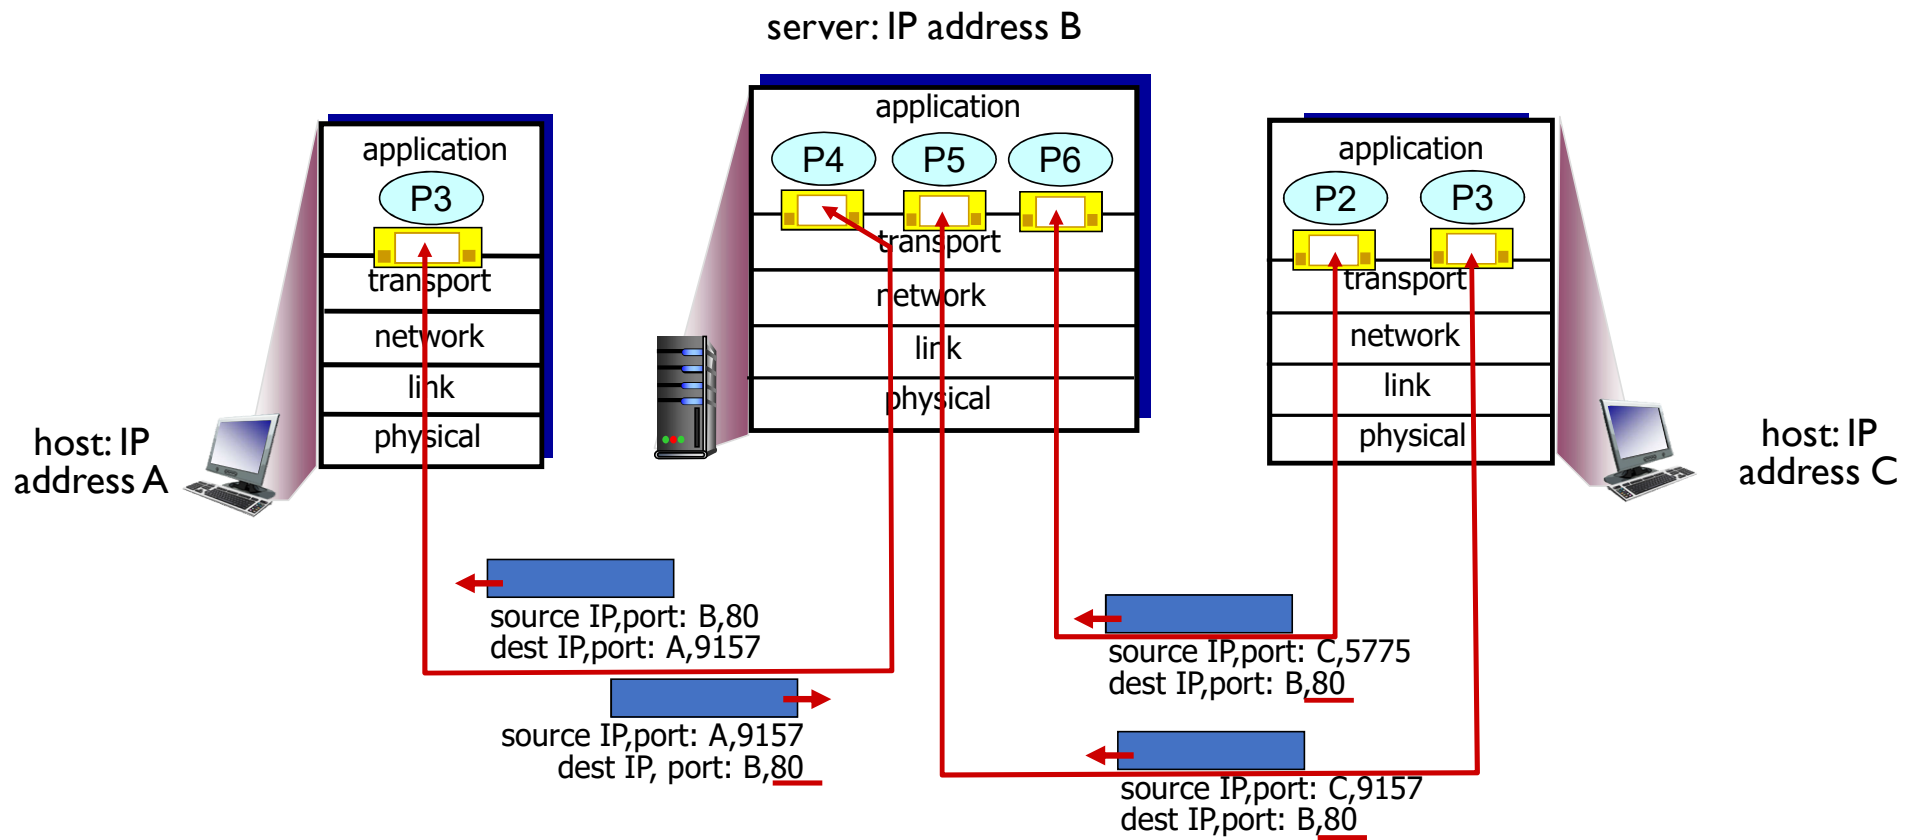


TCP/UDP segment format

CONNECTION-ORIENTED DEMUX

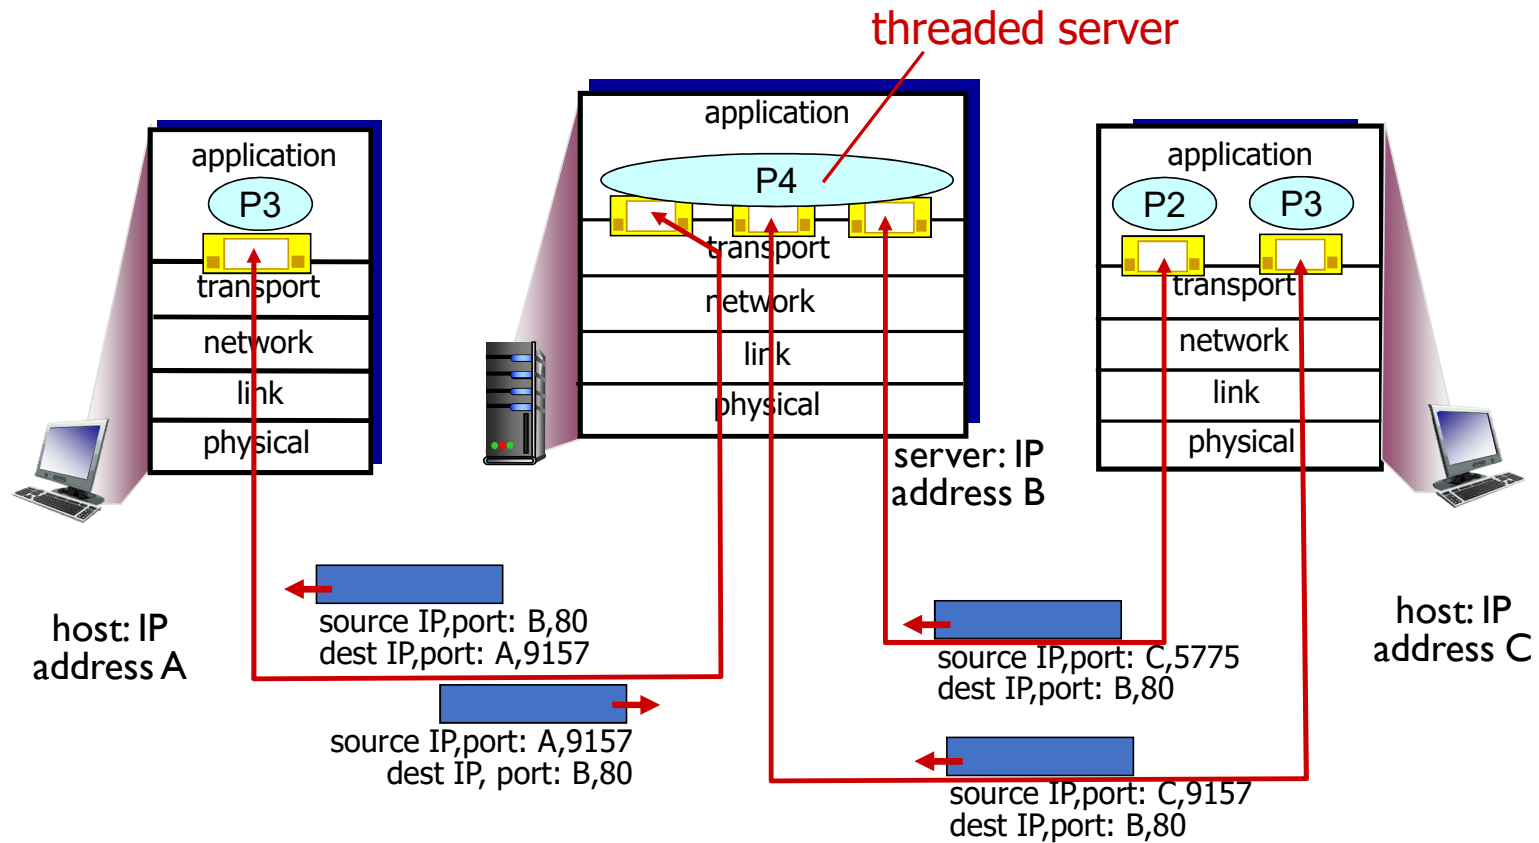
- ❖ Transport layer connection is identified by 4-tuple:
 - source IP address
 - source port number
 - dest IP address
 - dest port number
- ❖ demux: receiver uses all four values to direct segment to appropriate socket
- ❖ server host may support many simultaneous sockets:
 - each socket identified by its own 4-tuple
- ❖ web servers have different sockets for each connecting client.
 - non-persistent HTTP will have different socket for each request

CONNECTION-ORIENTED DEMUX: EXAMPLE



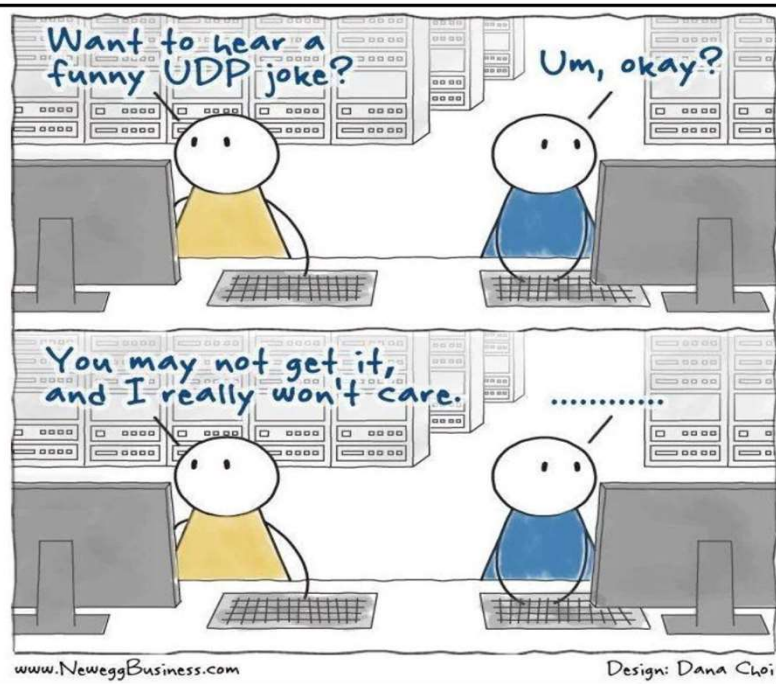
three segments, all destined to IP address: B,
dest port: 80 are demultiplexed to *different* sockets

CONNECTION-ORIENTED DEMUX: EXAMPLE



TCP AND UDP – A CLASSIC NETWORKING PUN

UDP



TCP

"Hi, I'd like to hear a TCP joke."
"Hello, would you like to hear a TCP joke?"
"Yes, I'd like to hear a TCP joke."
"OK, I'll tell you a TCP joke."
"Ok, I will hear a TCP joke."
"Are you ready to hear a TCP joke?"
"Yes, I am ready to hear a TCP joke."
"Ok, I am about to send the TCP joke. It will last 10 seconds, it has two characters, it does not have a setting, it ends with a punchline."
"Ok, I am ready to get your TCP joke that will last 10 seconds, has two characters, does not have an explicit setting, and ends with a punchline."
"I'm sorry, your connection has timed out. ...
Hello, would you like to hear a TCP joke?"

TRANSPORT LAYER SERVICES

End to end
packet delivery

UDP

Connection
Establishment

Reliable and
Ordered Data
Delivery

Flow Control

Congestion
Control

TCP

Transport Layer

IP Layer

Packet delivery (unreliable)

UDP: USER DATAGRAM PROTOCOL [[RFC 768](#)]

INTERNET STANDARD

RFC 768

J. Postel
ISI
28 August 1980


User Datagram Protocol

Introduction

This User Datagram Protocol (UDP) is defined to make available a datagram mode of packet-switched computer communication in the environment of an interconnected set of computer networks. This protocol assumes that the Internet Protocol (IP) [[1](#)] is used as the underlying protocol.

This protocol provides a procedure for application programs to send messages to other programs with a minimum of protocol mechanism. The protocol is transaction oriented, and delivery and duplicate protection are not guaranteed. Applications requiring ordered reliable delivery of streams of data should use the Transmission Control Protocol (TCP) [[2](#)].

Format

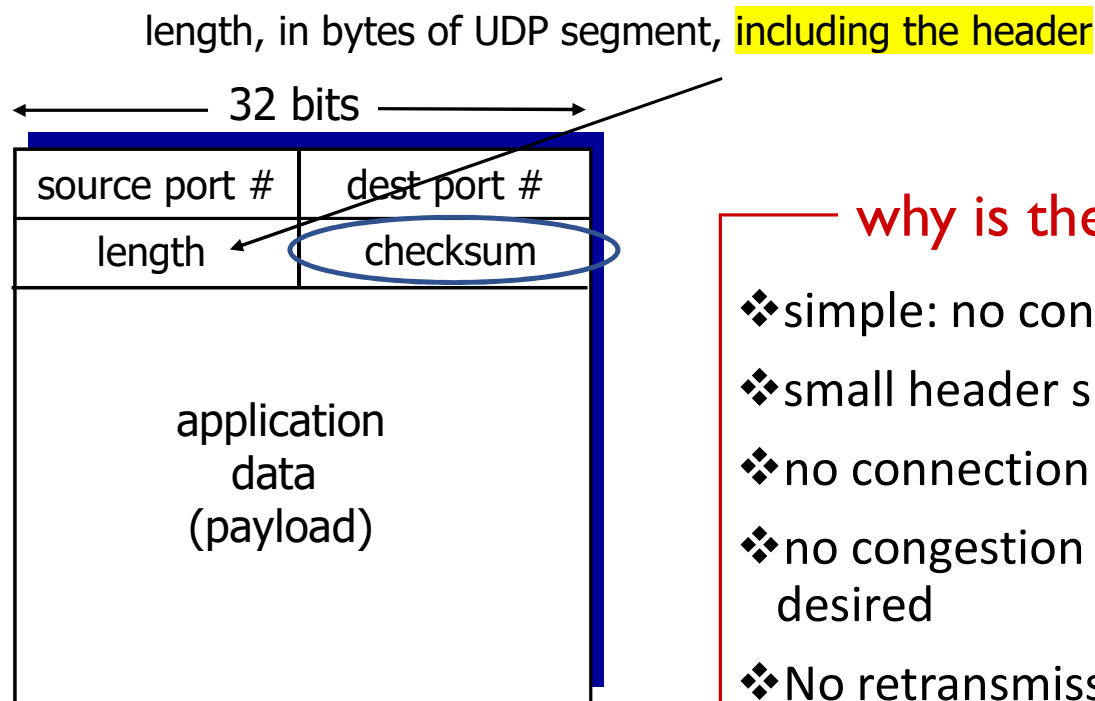


UDP: USER DATAGRAM PROTOCOL [RFC 768]

- *“best effort” service*
 - No guarantees on delivery,
 - delivered out-of-order
- *connectionless:*
 - *no handshaking* between UDP sender, receiver
- *Message/Datagram:*
 - each UDP segment handled independently of others
- ❖ UDP usage:
 - streaming multimedia apps (loss tolerant, rate sensitive)
 - DNS
 - DHCP
 - SNMP
- ❖ reliable transfer over UDP:
 - add reliability at application layer
 - application-specific error recovery!

Explore more on UDP-Lite, RUDP, DCCP and SCTP

UDP: SEGMENT HEADER



UDP segment format

why is there a UDP?

- ❖ simple: no connection state at sender, receiver
- ❖ small header size
- ❖ no connection establishment (which can add delay)
- ❖ no congestion control: UDP can blast away as fast as desired
- ❖ No retransmissions or in-order delivery: perceived - better control on latency
- ❖ ***Multicast***


UDP ERROR DETECTION

Goal: detect errors (*i.e.*, flipped bits) in transmitted segment

	1 st number	2 nd number	sum
Transmitted:	5	6	11

Received:

4	6	11
receiver-computed checksum		sender-computed checksum (as received)
≠		



```
int16_t checksum = 0;
int16_t comp_checksum (uint32_t count, int8_t* addr)
{
    while(1 < count) {
        sum += *(uint16_t)addr++;
        count -=2;
    }
    if (count)
        sum += *(uint16_t)addr++;

    while(sum >> 16)
        sum = (sum & 0xffff) + (sum >>16);

    return (checksum = ~sum);
}
```

UDP/INTERNET CHECKSUM

Goal: detect errors (*i.e.*, flipped bits) in transmitted segment

sender:

- treat contents of UDP segment (including UDP header and *IP header fields*) as sequence of 16-bit integers.
- **checksum:** addition (one's complement sum) of segment content.
- Update UDP checksum field with the computed checksum value.

receiver:

- compute checksum of received segment
- check if computed checksum equals checksum field value:
 - Not equal - error detected
 - Equal - no error detected. *But maybe errors nonetheless?* More later

<https://datatracker.ietf.org/doc/html/rfc1071>

INTERNET CHECKSUM: AN EXAMPLE

example: add two 16-bit integers

Note: When adding numbers, a carry from the msb needs to be added back to the result.

		1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
		1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1	1
sum		1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum		0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1
		1	1	1	0	0	1	1	0	0	1	1	0	0	1	1	0
		1	1	0	1	0	1	0	1	0	1	0	1	0	1	0	1
wraparound	1	1	0	1	1	1	0	1	1	1	0	1	1	1	0	1	1
sum		1	0	1	1	1	0	1	1	1	0	1	1	1	1	0	0
checksum		0	1	0	0	0	1	0	0	0	1	0	0	0	0	1	1

Binary sum

Add carry

1s complement

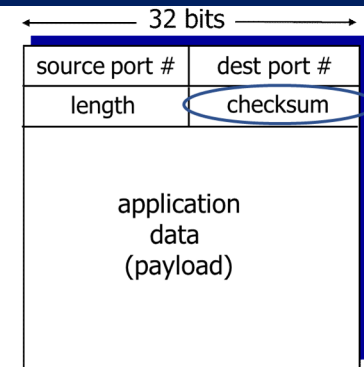
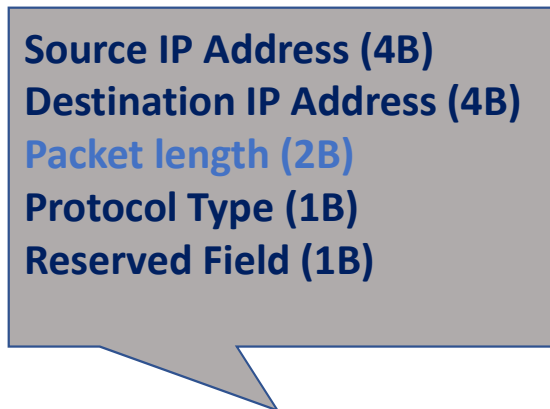
0 1

1 0

What about two or multi-bit errors?

Even though numbers have changed (bit flips), **no** change in checksum!

HOW IS UDP CHECKSUM COMPUTED?



Divide and slice the content into 16 bit pieces (2 Bytes).



At the Sender:

- Transport layer computes the Checksum and adds it to the Checksum field.
- Note: During the calculation of checksum, the checksum field is set with 0's. After the calculation, the computed checksum value will be updated.

At the Receiver:

- Transport layer computes the Checksum and *compares with sent value*.

- **point-to-point:**
 - one sender, one receiver
- **reliable, in-order *byte stream*:**
 - no “message boundaries”
- **full duplex data:**
 - bi-directional data flow in same connection
 - MSS: maximum segment size
- **pipelined:**
 - TCP congestion and flow control set window size
- **connection-oriented:**
 - handshaking (exchange of control msgs) inits sender and receiver state before data exchange
- **Flow controlled:**
 - sender will not overwhelm receiver
- **Congestion controlled:**
 - sender will not overwhelm the network/ adapt to the network capacity.

Today's Focus:
Understanding TCP Protocol
Connection Setup and Teardown
Sequence and Acknowledgment Numbers
Round Trip time and RTT Estimation

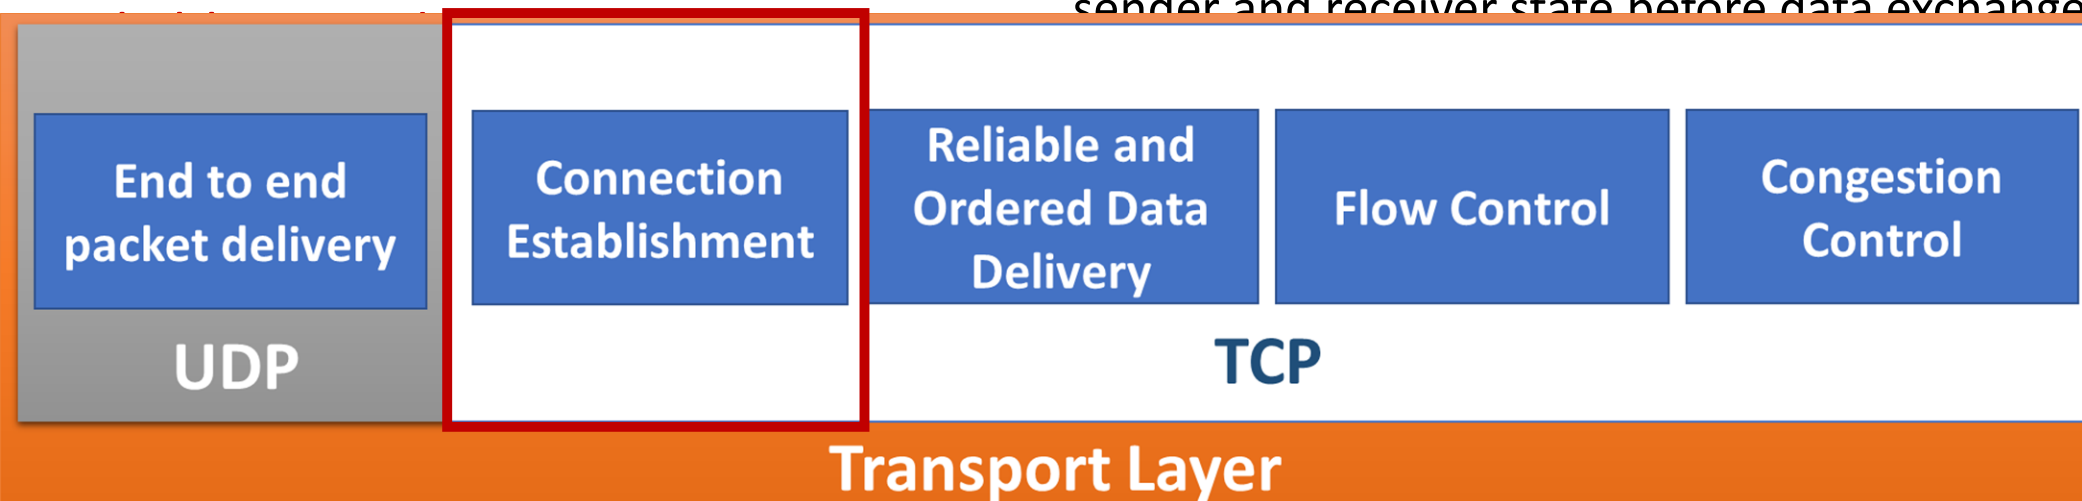
TCP: OVERVIEW RFCs: [793](#), [1011](#), [2873](#), [5961](#), [6093](#), [6429](#), [6528](#), [6691](#), [9293](#)

- point-to-point:

- one sender, one receiver

- connection-oriented:

- handshaking (exchange of control msgs) initiates sender and receiver state before data exchange



- TCP congestion and flow control set window size

to the network capacity.

Today's Focus:

Understanding TCP Protocol

Connection Setup and Teardown

Sequence and Acknowledgment Numbers

Round Trip time and RTT Estimation

TCP SEGMENT STRUCTURE

