

CS330 - Operating Systems

Deadlocks

13-10-2025

Deadlocks



Deadlocks

- Deadlock happens when
 - processes compete for access to limited resources
 - and are not synchronized properly

Thread 1:

```
pthread_mutex_lock(L1);  
pthread_mutex_lock(L2);
```

Thread 2:

```
pthread_mutex_lock(L2);  
pthread_mutex_lock(L1);
```

- Conditions for deadlock
 - Mutual exclusion
 - Hold-and-wait
 - No preemption
 - Circular wait


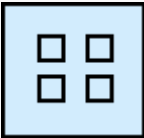
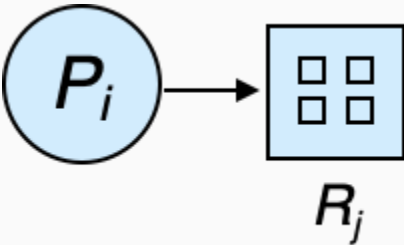
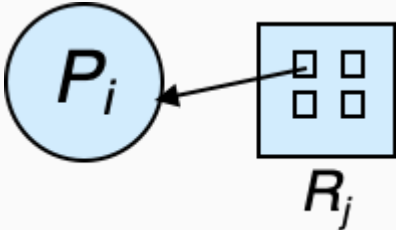
Handling Deadlocks

- Prevention
- Avoidance
- Recovery
- Ignore!


Preventing Deadlocks

- Make it impossible to have deadlocks
- Eliminate one of the four conditions
- Resource Allocation Graphs
 - System is a graph - A set of vertices V and a set of edges E
 - Processes and resources are vertices -
 $P = \{P_1, P_2, \dots, P_n\}$, the set consisting of all the processes
 $R = \{R_1, R_2, \dots, R_m\}$, the set consisting of all resource types
 - request edge : directed edge $P_i \rightarrow R_j$
 - assignment edge : directed edge $R_j \rightarrow P_i$


Resource Allocation Graphs

- Process 
- Resource Type with 4 instances 
- P_i requests instance of R_j 
- P_i is holding an instance of R_j 


Resource Allocation Graphs



Resource Allocation Graphs



Resource Allocation Graphs



Resource Allocation Graphs

- If graph contains no cycles, then no deadlock
- If graph contains a cycle
 - if only one instance per resource type, then definitely deadlock
 - if several instances per resource type, possibility of deadlock

Preventing Deadlocks

- Mutual exclusion
 - Must hold only for non-sharable resources
- Hold and wait
 - must guarantee that whenever a process requests a resource, it does not hold any other resources
 - All resources allocated upfront
 - Can request resource when it does not hold any
 - Low resource utilization; may lead to starvation

Preventing Deadlocks

- No Preemption
 - Release all resources currently being held
 - Preempted resources are added to the list of resources for which the process is waiting
 - Process restarts only when it can regain its old resources, as well as the new ones that it is requesting
- Circular wait
 - impose a total ordering of all resource types, and require that each process requests resources in an increasing order

Avoiding Deadlocks

- Control the allocation of resources such that it does not lead to deadlock
 - Requires that the system has some prior information available
 - Simplest model requires that each process **declares** the maximum number of resources of each type that it may need
 - Dynamically examine the resource-allocation state to ensure that there can never be a circular-wait condition
 - Resource-allocation state is defined by the number of available and allocated resources, and the maximum demands of the processes

Avoiding Deadlocks - Safe state

- System is in **safe state** if there exists a sequence $\langle P_1, P_2, \dots, P_n \rangle$ of ALL the processes in the system such that for each P_i , the resources that P_i can still request can be satisfied by currently available resources + resources held by all the P_j , with $j < i$
- If P_i resource needs are not immediately available, then P_i can wait until all P_j have finished
- When P_j is finished, P_i can obtain needed resources, execute, return allocated resources, and terminate
- When P_i terminates, P_{i+1} can obtain its needed resources, and so on

Avoiding Deadlocks - Safe state

- If a system is in safe state, then no deadlocks
- If a system is in unsafe state, then there is a possibility of deadlock
- Avoidance ensures that a system will never enter an **unsafe state**.


Avoiding Deadlocks

- Single instance of a resource type
 - Use a resource-allocation graph
- Multiple instances of a resource type
 - Use the banker's algorithm


Avoiding Deadlocks

- Resources must be claimed a priori in the system
- **Claim edge** $P_i \twoheadrightarrow R_j$ indicates that process P_i **may** request resource R_j
- Claim edge converts to request edge when a process requests a resource
- Request edge converted to an assignment edge when the resource is allocated to the process
- When a resource is released by a process, assignment edge reconverts to a claim edge

Avoiding Deadlocks



Avoiding Deadlocks



Avoiding Deadlocks

- Suppose that process P_i requests a resource R_j
 - The request can be granted only if converting the request edge to an assignment edge does not result in the formation of a cycle in the resource allocation graph

Avoiding Deadlocks

- Banker's Algorithm
 - Multiple instances
 - Each process must claim maximum use at the beginning
 - When a process requests a resource it may have to wait
 - When a process gets all its resources it must return them in a finite amount of time
 - Use different data structures to enforce this
 - Max matrix
 - Allocation matrix
 - Need matrix
 - Available list

Banker's Algorithm

- 5 processes P0 through P4;
- 3 resource types:
 - A (10 instances), B (5 instances), and C (7 instances)
- Snapshot at time T:

	Allocation			Max			Available		
	A	B	C	A	B	C	A	B	C
P0	0	1	0	7	5	3	3	3	2
P1	2	0	0	3	2	2			
P2	3	0	2	9	0	2			
P3	2	1	1	2	2	2			
P4	0	0	2	4	3	3			

Banker's Algorithm

- Available:

A B C

3 3 2

- Snapshot at time T:

	Allocation				Max				Need		
	A	B	C		A	B	C		A	B	C
P0	0	1	0		7	5	3		7	4	3
P1	2	0	0		3	2	2		1	2	2
P2	3	0	2		9	0	2		6	0	0
P3	2	1	1		2	2	2		0	1	1
P4	0	0	2		4	3	3		4	3	1

Banker's Algorithm

- Available:

A	B	C
3	3	2

- Snapshot at time T:

	Allocation		
	A	B	C
P0	0	1	0
P1	2	0	0
P2	3	0	2
P3	2	1	1
P4	0	0	2

	Max		
	A	B	C
P0	7	5	3
P1	3	2	2
P2	9	0	2
P3	2	2	2
P4	4	3	3

	Need		
	A	B	C
P0	7	4	3
P1	1	2	2
P2	6	0	0
P3	0	1	1
P4	4	3	1

Is the system in safe state?

Banker's Algorithm

- Available:

A	B	C
3	3	2

- Snapshot at time T:

	Allocation		
	A	B	C
P0	0	1	0
P1	2	0	0
P2	3	0	2
P3	2	1	1
P4	0	0	2

	Max		
	A	B	C
P0	7	5	3
P1	3	2	2
P2	9	0	2
P3	2	2	2
P4	4	3	3

	Need		
	A	B	C
P0	7	4	3
P1	1	2	2
P2	6	0	0
P3	0	1	1
P4	4	3	1

Yes, there exists a sequence
<P1, P3, P4, P2, P0>

Deadlock Detection


- Allow system to enter deadlock state
- Detection algorithm
- Recovery scheme

Deadlock Detection with Single Instances

- Maintain **wait-for** graph
 - Nodes are processes
 - $P_i \rightarrow P_j$ if P_i is waiting for P_j
- Periodically, invoke an algorithm that searches for cycles in the graph. If there is a cycle, there exists a deadlock!
- What is the complexity of the algorithm?


Deadlock Detection with Single Instances

Resource-Allocation Graph



(a)

Wait-for Graph



(b)

Several Instances?

Deadlock Recovery

- Abort all deadlocked processes (or)
- Abort 1 process at a time until the cycle is eliminated
 - In which order should we choose to abort?
 - Priority of the process
 - How long process has computed, and how much longer to completion
 - Resources the process has used
 - Resources process needs to complete
 - How many processes will need to be terminated
 - Is process interactive or batch
- Select a victim while minimizing cost
- **Rollback** to some safe state, restart process for that state
- Same process may always be picked as victim (starvation)