

# CS330 - Operating Systems

## Locks and CVs

15-09-2025

# Implementing Synchronization

- Interrupt-controls
- Peterson's solution
  - Classic solution (works on early architectures) for two threads
  - Uses load and store (which were atomic in early arch.)

```
boolean flag[2];
int turn;

void init() {
    flag[0] = flag[1] = false;
}

Can lead to starvation on one or
both threads
```

```
void lock() {
    flag[i] = true;

    while (flag[j])
        // wait until it's i's turn
}

void unlock() {
    flag[i] = false;
}
```

# Implementing Synchronization

- Interrupt-controls
- Peterson's solution
  - Classic solution (works on early architectures) for two threads
  - Uses load and store (which were atomic in early arch.)

```
boolean flag[2];
int turn;

void init() {
    flag[0] = flag[1] = false;
    // whose turn is it? (thread 0 or 1)
    turn = 0;
}
```

```
void lock() {
    flag[i] = true;
    turn = j; // make it other thread's turn
    while (flag[j] && (turn == j));
        // wait until it's i's turn
}

void unlock() {
    flag[i] = false;
}
```

# Implementing Synchronization

- Interrupt-controls
- Peterson's solution
- Test and Set

```
int TestAndSet(int *old_ptr, int new)
    // fetch old value at old_ptr
    int old = *old_ptr;
    // store 'new' into old_ptr
    *old_ptr = new;
    // return the old value
    return old;
}
```

```
typedef struct __lock_t {
    int flag;
} lock_t;

void init(lock_t *lock) {
    // 0: lock is available, 1: lock is held
    lock->flag = 0;
}

void lock(lock_t *lock) {
    while (TestAndSet(&lock->flag, 1) == 1)
        ; // spin-wait (do nothing)
}

void unlock(lock_t *lock) {
    lock->flag = 0;
}
```

# Implementing Synchronization

- Interrupt-controls
- Peterson's solution
- Test and Set
- Compare and Swap

```
int CompareAndSwap(int *ptr, int expected, int new) {  
    int original = *ptr;  
    if (original == expected)  
        *ptr = new;  
    return original;  
}  
  
void lock(lock_t *lock) {  
    while (CompareAndSwap(&lock->flag, 0, 1) == 1)  
        ; // spin  
}
```

# Implementing Synchronization

- Cons of spin-wait locks
  - Performance
    - Interrupt thread inside critical section
  - Priority Inversion

<u>T1</u>	<u>T2</u>
Ready	Running
Ready	Blocked
Running (Lock)	Ready
Ready (Lock)	Running (Wait) (Waits forever)

# Too Much Spinning

- Context switch in critical section
  - Performance issues
  - Hung execution (Not fair)



"What Really Happened on Mars?" by Glenn E. Reeves.  
[research.microsoft.com/en-us/um/people/mbj/Mars\\_Pathfinder/Authoritative\\_Account.html](http://research.microsoft.com/en-us/um/people/mbj/Mars_Pathfinder/Authoritative_Account.html)

# Too Much Spinning

- Yield
  - Assume `yield` function
  - Instead of spinning, yield the CPU

```
while (TestAndSet(&flag, 1) == 1) { yield(); }
```
  - What happens with multiple threads?
- Keep track of threads that are waiting in a queue

# Example: Bounded buffer

- Producer(s) put things into a shared buffer
  - Consumer(s) take them out
  - Need synchronization to coordinate producers/consumers
- 
- Multi-threaded web server
  - Pipes
- ```
grep foo file.txt | wc -l
```

# Example: Bounded buffer

- Producer(s) put things into a shared buffer
- Consumer(s) take them out
- Need synchronization to coordinate producers/consumers

```
int buffer;  
int count = 0; // initially, empty
```

```
void put (int value) {  
    count = 1;  
    buffer = value;  
}
```

```
int get() {  
    count = 0;  
    return buffer;  
}
```

# Example: Bounded buffer

- Producer(s) put things into a shared buffer
- Consumer(s) take them out
- Need synchronization to coordinate producers/consumers

```
int buffer;  
int count = 0; // initially, empty
```

```
void put (int value) {  
    count = 1;  
    buffer = value;  
}
```

```
int get() {  
    count = 0;  
    return buffer;  
}
```

```
void *producer(int arg) {  
    int i;  
    int loops = arg;  
    for (i = 0; i < loops; i++) {  
        put(i);  
    }  
}
```

```
void *consumer() {  
    while (1) {  
        int tmp = get();  
        printf("%d\n", tmp);  
    }  
}
```

# Synchronization between Parent and Child

```
1 void *child(void *arg) {  
2     printf("child\n");  
3     // XXX how to indicate we are done?  
4     return NULL;  
5 }  
6  
7 int main(int argc, char *argv[]) {  
8     printf("parent: begin\n");  
9     pthread_t c;  
10    Pthread_create(&c, NULL, child, NULL); // child  
11    // XXX how to implement the wait for child?  
12    printf("parent: end\n");  
13    return 0;  
14 }
```

```
1 void done() {  
2     cond_signal(&c);  
3 }  
4  
5 }  
6  
7 void join() {  
8     cond_wait(&c);  
9 }  
10 }  
11 }
```

# Synchronization between Parent and Child

```
1 void *child(void *arg) {  
2     printf("child\n");  
3     // XXX how to indicate we are done?  
4     return NULL;  
5 }  
6  
7 int main(int argc, char *argv[]) {  
8     printf("parent: begin\n");  
9     pthread_t c;  
10    Pthread_create(&c, NULL, child, NULL); // child  
11    // XXX how to implement the wait for child?  
12    printf("parent: end\n");  
13    return 0;  
14 }
```

```
1 void done() {  
2     lock(&m);  
3     cond_signal(&c);  
4     unlock(&m);  
5 }  
6  
7 void join() {  
8     lock(&m);  
9     cond_wait(&c);  
10    unlock(&m);  
11 }
```

# Synchronization between Parent and Child

```
1 void *child(void *arg) {  
2     printf("child\n");  
3     // XXX how to indicate we are done?  
4     return NULL;  
5 }  
6  
7 int main(int argc, char *argv[]) {  
8     printf("parent: begin\n");  
9     pthread_t c;  
10    Pthread_create(&c, NULL, child, NULL); // child  
11    // XXX how to implement the wait for child?  
12    printf("parent: end\n");  
13    return 0;  
14 }
```

```
1 void done() {  
2     lock(&m);  
3     done = 1;  
4     cond_signal(&c);  
5     unlock(&m);  
6 }  
7  
8 void join() {  
9     lock(&m);  
10    if (done == 0)  
11        cond_wait(&c, &m);  
12    unlock(&m);  
13 }
```

# Condition Variables

- Waiting inside a critical section
  - Called only when holding a lock
- Wait: atomically release lock and relinquish processor
  - Reacquire the lock when wakened
- Signal: wake up a waiter, if any
- Broadcast: wake up all waiters, if any

```
pthread_cond_wait(pthread_cond_t *c, pthread_mutex_t *m);  
pthread_cond_signal(pthread_cond_t *c);
```

# Producer-consumer with CV

- Single producer and consumer

```
int loops; // must initialize somewhere...
cond_t cond;
mutex_t mutex;

void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);
        if (count == 1)
            Pthread_cond_wait(&cond, &mutex);
        put(i);
        Pthread_cond_signal(&cond);
        Pthread_mutex_unlock(&mutex);
    }
}
```

```
void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);
        if (count == 0)
            Pthread_cond_wait(&cond, &mutex);
        int tmp = get();
        Pthread_cond_signal(&cond);
        Pthread_mutex_unlock(&mutex);
        printf("%d\n", tmp);
    }
}
```

# Multiple Producers and Consumers

- Mesa Semantics
  - Assume we have two consumers  $Tc1$  and  $Tc2$  and producer  $Tp$
  - Consumer  $Tc1$  runs and *waits* because buffer is empty
  - Producer runs and fills the buffer
    - Signals the buffer is filled
    - Moves  $Tc1$  from waiting/sleeping state to ready state
  - $Tp$  sleeps as buffer is full
  - $Tc2$  starts running and eats up the buffer; goes to sleep
  - $Tc1$  is already in the ready queue
    - Runs next and finds that the buffer is empty!!!

# Producer-consumer with CV

- Modified code

(All code is available at: <https://github.com/remzi-arpacidusseau/ostep-code/>)

```
int loops;
cond_t cond;
mutex_t mutex;

void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);
        while (count == 1)
            Pthread_cond_wait(&cond, &mutex);
        put(i);
        Pthread_cond_signal(&cond);
        Pthread_mutex_unlock(&mutex);
    }
}

void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);
        while (count == 0)
            Pthread_cond_wait(&cond, &mutex);
        int tmp = get();
        Pthread_cond_signal(&cond);
        Pthread_mutex_unlock(&mutex);
        printf("%d\n", tmp);
    }
}
```

# Producer-consumer Problem

- Use two condition variables

```
cond_t empty, fill;
mutex_t mutex;

void *producer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);
        while (count == 1)
            Pthread_cond_wait(&empty, &mutex);
        put(i);
        Pthread_cond_signal(&fill);
        Pthread_mutex_unlock(&mutex);
    }
}
```


```
void *consumer(void *arg) {
    int i;
    for (i = 0; i < loops; i++) {
        Pthread_mutex_lock(&mutex);
        while (count == 0)
            Pthread_cond_wait(&fill, &mutex);
        int tmp = get();
        Pthread_cond_signal(&empty);
        Pthread_mutex_unlock(&mutex);
        printf("%d\n", tmp);
    }
}
```

# Condition Variables

- Hold lock when calling `wait()` and `signal()`
  - For shared state; always hold lock when accessing shared state
- Condition variables are memoryless
  - If signal happens when no one is waiting, NO-OP
  - If not wait before signal wakes up the waiter
- Wait atomically releases lock
- Woken thread may not run immediately
  - signal puts thread on ready list; anyone can acquire it after release
- Wait must happen in “loop”

# Semaphores

Semaphore flag signals



© 2013 Encyclopædia Britannica, Inc.

# Semaphores

- Single primitive for all synchronization (Dijkstra)
- Positive Integer
  - Initialize its value to any positive integer, but after that the only operations you are allowed to perform are increment and decrement
  - Cannot read the current value
  - When a thread decrements the semaphore, if the result is negative, the thread blocks itself and cannot continue until another thread increments the semaphore
  - When a thread increments the semaphore, if there are other threads waiting, one of the waiting threads gets unblocked

# Semaphores

- Signal (V)
  - increment\_and\_wake\_a\_waiting\_process\_if\_any
  - sema\_up
- Wait (P)
  - decrement\_and\_block\_if\_the\_result\_is\_negative
  - sema\_down

# Semaphores

- Signal (V)
  - increment\_and\_wake\_a\_waiting\_process\_if\_any
- Wait (P)
  - decrement\_and\_block\_if\_the\_result\_is\_negative
- Impose deliberate constraints to avoid errors.
- Often clean and organized
- Implemented efficiently on many systems

# Rendezvous

Given two threads T1, T2:

T1 : a1 a2

T2 : b1 b2

Guarantee that a1 happens before b2 and b1 before a2

# Semaphore Usage

- Binary semaphore (or locks (or mutexes))
  - Enforce mutual exclusion

## Thread

```
mutex.wait()  
// critical section  
count = count + 1  
mutex.signal()
```

- Multiplex
  - Generalize the previous solution so that it allows multiple threads to run in the critical section at the same time but no more than  $n$  threads can run in the critical section at the same time

# Semaphores Usage

- Counting Semaphores
  - Similar in use to condition variables
- Finite buffer producer-consumer problem

```
mutex = Semaphore(1)
```

```
items = Semaphore(0)
```

```
spaces = Semaphore(buffer.size())
```

```
items.wait()
```

```
mutex.wait()
```

```
    buffer.get()
```

```
mutex.signal()
```

```
spaces.signal()
```

```
spaces.wait()
```

```
mutex.wait()
```

```
    buffer.put(i)
```

```
mutex.signal()
```

```
items.signal()
```

# Producer-consumer Problem with Semaphores

- Finite buffer producer-consumer problem

```
int buffer[SIZE]
items = Semaphore(0)
spaces = Semaphore(buffer.size())

items.wait()
    buffer.get()
spaces.signal()

spaces.wait()
    buffer.put(i)
items.signal()
```

# Producer-consumer Problem with Semaphores

```
int buffer[MAX];
int fill = 0;
int use=0;

void put(int value) {
    buffer[fill] = value;
    fill = (fill+1)%MAX;
}

items = Semaphore(0)
spaces = Semaphore(buffer.size())

items.wait()
    buffer.get()
spaces.signal()

int get() {
    int tmp = buffer[use];
    use = (use+1)%MAX;
    return tmp;
}
```

# Producer-consumer Problem with Semaphores

```
mutex = Semaphore(1)
items = Semaphore(0)
spaces = Semaphore(buffer.size())
```

```
mutex.wait()
items.wait()
    buffer.get()
spaces.signal()
mutex.signal()
```


Consumer

```
mutex.wait()
spaces.wait()
    buffer.put(i)
items.signal()
mutex.signal()
```

Producer


# Dining Philosophers Problem

- Five philosophers
  - Each pair shares a fork
  - Each need two to eat
  - Think when do not eat
- ```
while (true) {  
    think ();  
    get_forks (p);  
    eat ();  
    put_forks (p);  
}
```




# Dining Philosophers Problem

```
int left(int p) {  
    return p; }  
  
int right(int p) {  
    return (p + 1) % 5; }  
  
void put_forks(int p) {  
}  
  
void get_forks(int p) {  
}
```



# Dining Philosophers Problem

```
int left(int p) {  
    return p; }  
  
int right(int p) {  
    return (p + 1) % 5; }  
  
void put_forks(int p) {  
    sema_up(&forks[left(p)]);  
    sema_up(&forks[right(p)]);}  
  
void get_forks(int p) {  
    sema_down(&forks[left(p)]);  
    sema_down(&forks[right(p)]);}
```



# Deadlocks



# Deadlocks

## Thread 1:

```
pthread_mutex_lock(L1);  
pthread_mutex_lock(L2);
```

## Thread 2:

```
pthread_mutex_lock(L2);  
pthread_mutex_lock(L1);
```

- Conditions for deadlock

- Mutual exclusion
  - exclusive access of locks
- Hold-and-wait
  - hold locks when waiting
- No preemption
  - locks cannot be removed from threads
- Circular wait
  - circular chain of threads holding each others' locks

# Real-world Concurrency Threats - Example

- Race in the Therac-25 radiation therapy machine caused massive overdose
  - resulted in patient deaths and serious injuries
  - Data Entry and Keyboard Handler routines share a variable, which recorded whether the technician had completed entering commands.
  - A race condition bug of this shared variable cause the UI to display the wrong mode to operators