

CS 330 - Operating Systems

Paging


01-09-2025

Page Tables

- Stores virtual-to-physical address translations
- One page table per process
- Where are page tables stored?

Page Tables

- Stores virtual-to-physical address translations
- One page table per process
- Where are page tables stored?
 - Can get large
- Page table entry



Slowness of Paging

```
movl 21, %eax
```

```
VPN      = (VirtualAddress & VPN_MASK) >> SHIFT  
PTEAddr = PageTableBaseRegister + (VPN * sizeof(PTE))
```

```
offset   = VirtualAddress & OFFSET_MASK  
PhysAddr = (PFN << SHIFT) | offset
```

Summary of Paging

- Manage memory in fixed size units, or pages
- Each process has its own page table
 - Stored in physical memory
 - Hardware registers
 - pointer to page table start
 - page table length
- Cons
 - **Too slow**
 - Too much memory

Translation Lookaside Buffer (TLB)

- In memory management unit (MMU)
- Cache of certain VA to PA translations
 - Address translation cache

TLB Basic Algorithm

1. Extract VPN from VA
2. Check if translation for VPN is in TLB
3. If **TLB hit**
 - a. Retrieve the translation entry
 - b. Extract PFN from TLB entry
 - c. Concatenate offset from VA to obtain PA
 - d. (access checks should not fail)
4. If **TLB miss**
 - a. Access the page table to retrieve translation (if valid)
 - b. Update TLB with translation
 - c. *Retry instruction (step 2)*

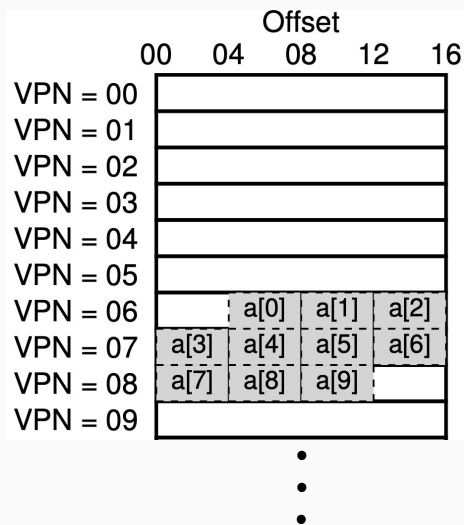
Example

```
int sum = 0;  
for (i = 0; i < 10; i++) {  
    sum += a[i]; }  
}
```


Example

```
int sum = 0;
for (i = 0; i < 10; i++) {
    sum += a[i]; }

```



Example

```
int sum = 0;
for (i = 0; i < 10; i++) {
    sum += a[i]; }

```

Diagram illustrating a memory layout with 16 words (VPNs) and an access time of 100 ns. The words are indexed by Offset (00, 04, 08, 12, 16) and VPN (00 to 09). The layout shows a sequence of words, with the last three words (VPN 06, 07, 08) containing specific data elements (a[0] through a[9]).

VPN	Offset	Content
00	00	
01	04	
02	08	
03	12	
04	16	
05		
06		a[0], a[1], a[2]
07		a[3], a[4], a[5], a[6]
08		a[7], a[8], a[9]
09		

•
•
•

a[0]	a[1]	a[2]	a[3]	a[4]	a[5]	a[6]	a[7]	a[8]	a[9]
M	H	H	M	H	H	H	M	H	H

TLB Hit Rate = No. of hits/No. of accesses

Caching with TLB

- Spatial locality
 - Access to a means program may access $a+1$
 - Normally the next instruction is accessed
- Temporal locality
 - Re-access accessed data or instruction (loops)
- Why not big caches?

Memory Cycle Rate =

$\text{Hit rate} * \text{Hit cycles} + \text{Miss rate} * (\text{Miss cycles} + \text{Hit cycles})$

TLB Contents

- Fully associative

- Translation may be anywhere
- TLB entry VPN | PFN | other bits
- Other bits
 - Valid bit
 - Protection bit
 - Address-space identifier
 - Dirty bit
 - ...

Sharing TLB

Process 1

VPN	PFN
1	21
3	39

Sharing TLB

Process 1 and Process 2

VPN	PFN
1	21
3	39
1	49

Flushing TLB

Process 1 and Process 2

VPN	PFN	Valid bit
1	21	0
3	39	0
1	49	1

Address space identifiers

Process 1 and Process 2

VPN	PFN	Valid bit	ASID
1	21	1	1
3	39	1	1
1	49	1	2

TLB Entry Replacement

- Installing a new entry
 - Replace when no more space in TLB
- Cache replacement policies

Techniques So Far

- Base & Bounds
 - Pros: Very quick, 2 registers
 - Cons: Contiguous block of memory -> fragmentation .
- Segmentation
 - Pros: Still relatively simple, 3 registers, lesser fragmentation
 - Cons: Still contiguous block of memory for segment
- Paging
 - Pros: Very low chances of segmentation
 - Cons: Slow, lots of memory accesses; overhead/process is huge!
- Paging + TLB
 - Pros: Improves address translation (spatial & temporal locality)
 - Cons: Limited in size, overhead/process still huge

Paging Memory Overheads

Assuming a

32-bit address space (2^{32} bytes),

4KB (2^{12} bytes) pages,

and 4-byte PTE

what is the size of the page table?

Paging Memory Overheads

Assuming a

32-bit address space (2^{32} bytes),

4KB (2^{12} bytes) pages,

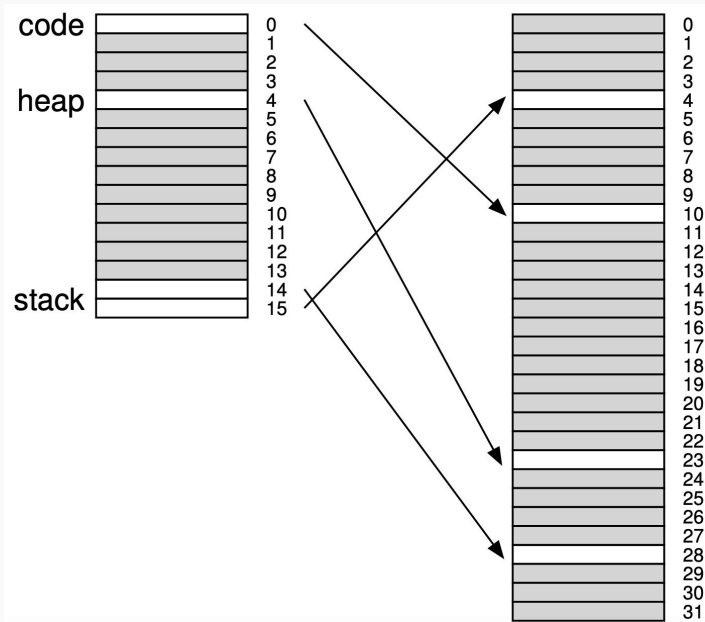
and 4-byte PTE

what is the size of the page table?

With 100 processes running simultaneously?

Combine Pages and Segments


16 KB address spaces having 1 KB pages



PFN	Valid	...
10	1	..
-	0	
-	0	
-	0	
23	1	
-	0	
-	0	
-	0	
-	0	
-	0	
-	0	
-	0	
-	0	
-	0	
28	1	
4	1	

Combine Pages and Segments

- Per segment PT
- Different size
- Separate base & bounds per segment
- Base register
 - Stores PA of start of PT
- Bounds
 - Stores PA of end of PT (or number of valid entries in PT)



Paging with Segmentation

- Suppose we have the following virtual address design:

Segment No.		Page No.		Offset	

2 bits		8 bits		8 bits	

Virtual addresses are translated into 16-bit physical addresses of the following form:

Physical Frame No.		Offset	

8 bits		8 bits	

What is the page size?

What is the maximum size of the physical memory?

Paging with Segmentation

- Suppose we have the following virtual address design:

Segment No.	Page No.	Offset
-----	-----	-----
2 bits	6 bits	8 bits

0x4225?

Virtual addresses are translated into 16-bit physical addresses:

Physical Frame No.	Offset
-----	-----
8 bits	8 bits

Segment Table:

Segment	Base Addr	Limit
0x0	0x2100	0x1000
0x1	0x3500	0x2000
0x2	0x1500	0x0200

Page tables at:

0x2100	0x3500	0x1500
0x10	0x35	0x75
0x15	0x45	0x80
0x20	0x55	0x85
0x25	0x65	0x90

Combine Pages and Segments


- Pros:
 - Leads to memory saving (Large gaps between segments)
- Cons:
 - Uses segmentation
 - Assumes certain usage pattern of address space
 - Sparsely used segments (sub-segment internal fragmentation) have same space waste issue
 - Variable size page tables can lead to fragmentation

Combine Pages and Pages

- Avoid segmentation with multi-level paging
- Linear page table → tree (hierarchical structure)
- Multi-level paging
 - Break PT into pages
 - Allocate page of PT if at least one entry in PT is valid
 - **Page directory**
 - Indicates where a page of the PT is (or) if invalid

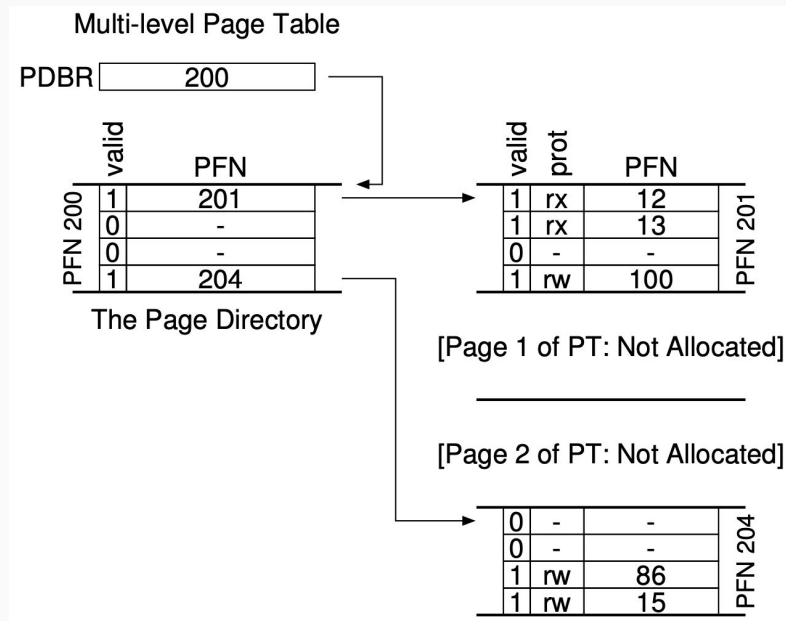

Multi-level Paging

- Suppose a 30-bit VA space and 512 byte page-size
 - VA has 21-bit VPN and 9-bit offset
 - With 4-bytes PTE, 128 PTEs forms one page (7 bits for PN)
 - Remaining 14 bits are for page directory
 - Spans 128 pages!



- TLB with multi-level paging
 - Multiple additional memory accesses to look up translation

Multi-level Paging



Multi-level Paging (Example)

Address space - 16 KB

What is the size of the virtual address in bits?

Page size - 64 bytes

PTE - 4 bytes

No. of bits for VPN and offset?

No. of entries in linear page table?

Size of the linear page table?

For 64-byte pages, with 1 KB table: 16 64-byte pages, i.e., each second level PT will hold 16 PTE

Multi-level Paging (Example)

Address space - 16 KB
(14-bit VA - 8 for VPN and 6 for offset)


Page size - 64 bytes

PTE - 4 bytes

Code - pages 0 and 1

Heap - pages 4 and 5

Stack - pages 254 and 255



Multi-level Paging (Example)

- Assume that the page directory has two entries for two page tables stored at 100 and 101:

Page Directory		Page of PT (@PFN:100)			Page of PT (@PFN:101)		
PFN	valid?	PFN	valid	prot	PFN	valid	prot
100	1	10	1	r-x	—	0	—
—	0	23	1	r-x	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	80	1	rw-	—	0	—
—	0	59	1	rw-	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	—	0	—
—	0	—	0	—	55	1	rw-
101	1	—	0	—	45	1	rw-

Code - pages 0 and 1
Heap - pages 4 and 5
Stack - pages 254 and 255

0x3F80?

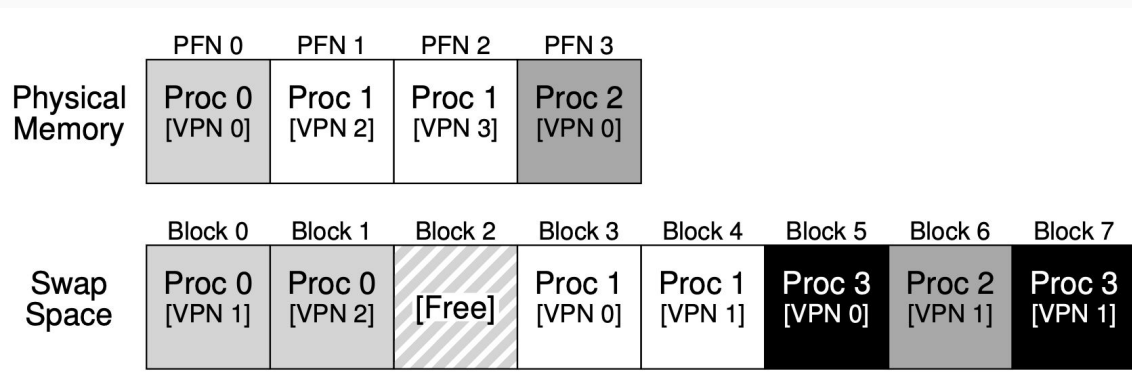
0x0030?

Larger Virtual Memory

- Virtual memory may be larger than physical memory
- Stash address spaces from physical memory that are not in use
 - Normally, in a location that has more capacity
 - Hard disk
- Larger devices to provide illusion of larger virtual address space
 - Helps with multiprogramming
- Memory overlays in older systems

Demand Paging

- Disk space to move pages back and forth
 - Swap space
 - Remember disk address of a page



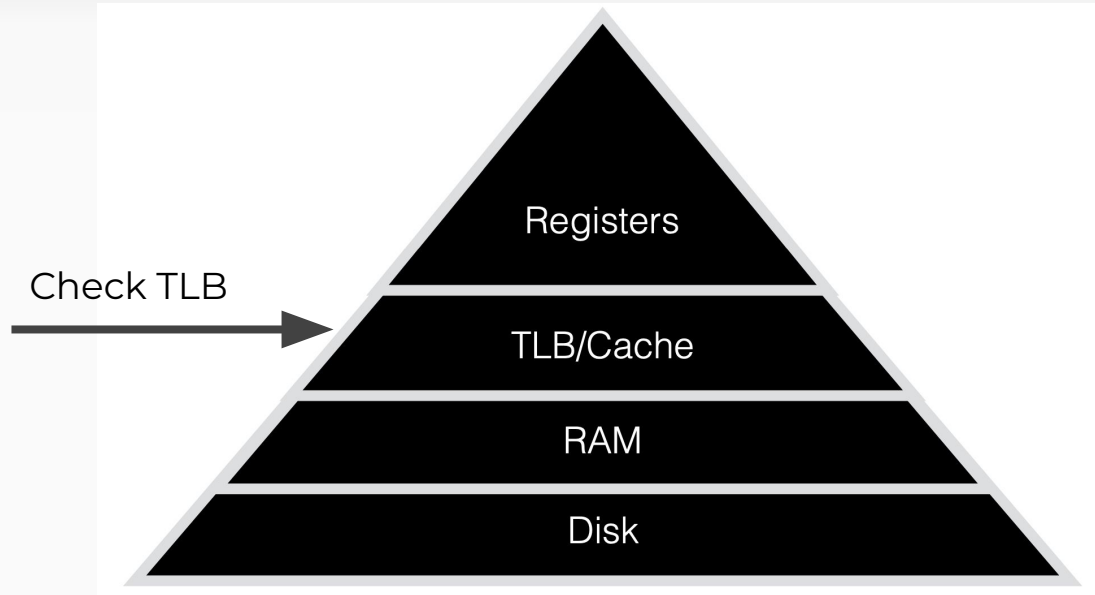
Cache Management

- Main memory caches some/all pages of a process
- Minimize number of misses when replacing pages
- Given the number of hits and misses, the **average memory access time** for a program is:

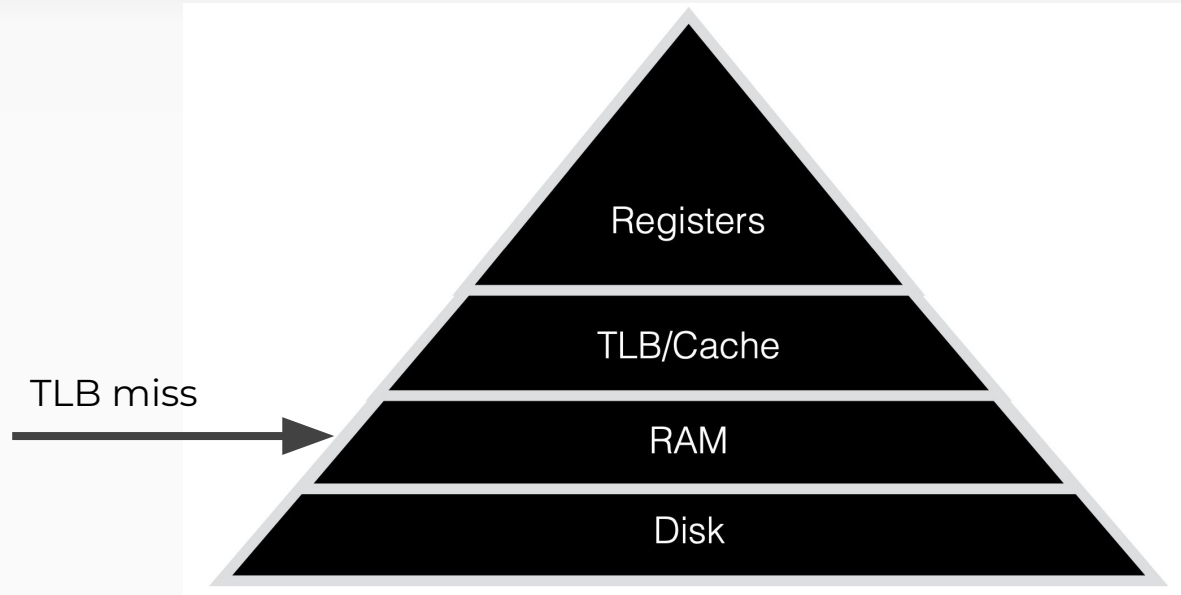
$$AMAT = TM + (PMiss * TD)$$

where TM represents the cost of accessing memory,
TD the cost of accessing disk, and
PMiss the probability of a miss

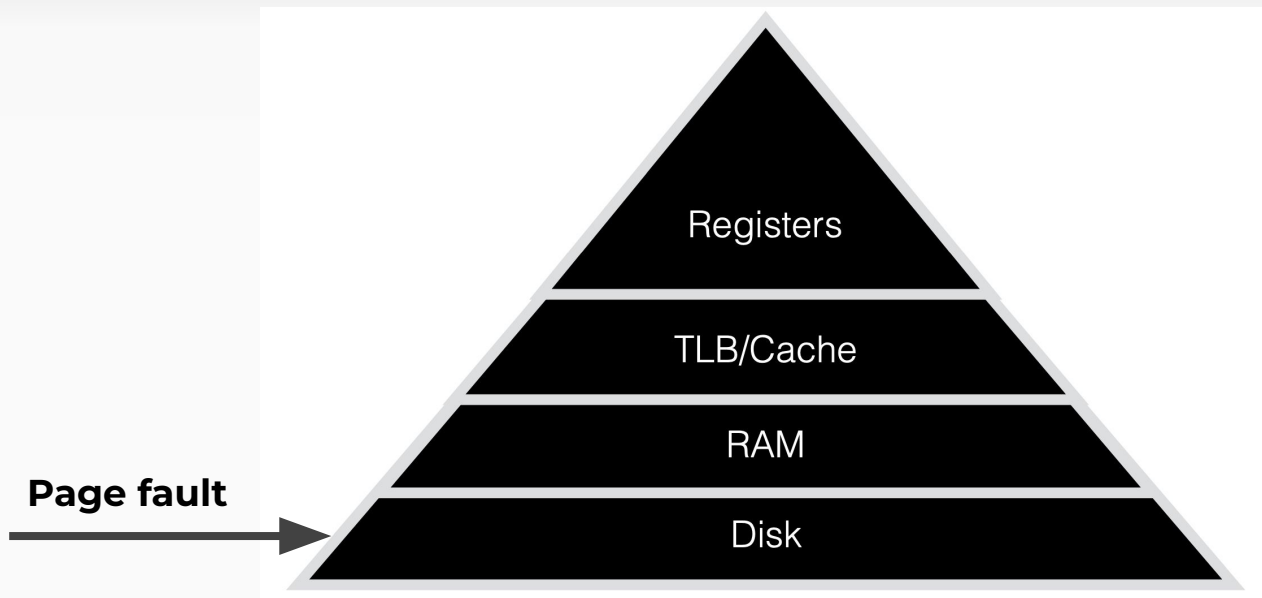
Page faults



Page faults



Page faults



$$\text{Avg. Mem. Access Time} = \text{Time-Mem} + (\text{PMiss} * \text{Time-Disk})$$

Present Bit

- Program generates virtual memory address
- Translated into physical address
 - Extracts VPN from VA
 - Checks TLB (TLB hit or not)
 - TLB miss
 - Locates PT in memory
 - Looks up PTE and gets the page (if in memory)
 - Page not present in physical memory
 - Identify using Present bit
 - Present bit set to one means the page is in physical memory
 - If not, it's on the disk (**page fault**)

Page faults

- When page is not found, page-fault handler in the OS is triggered
 - Looks up PTE bits for obtaining this data
 - Fetch page from disk address in memory
- Swap in (Disk → Memory)
 - Mark page as present in PT
- Swap out (Memory → Disk)
- Expensive to handle page faults
 - Helps with multiprogramming

Memory Full?

- Memory may be full
 - *High watermark* and *low watermark*
 - Start evicting when fewer than LW pages
 - Stop when there are HW free pages available in memory
- Swap out (page out) one or more pages
- Pick a page to replace in memory
 - Page replacement policy
 - Need good page replacement policies
- Can swap a group of pages for efficiency

Optimal Page Replacement Strategy

- Example
 - Assume the following page accesses:
1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- If we have 3 physical pages, how many page faults?
- With 4?

Optimal Page Replacement Strategy

- Optimal strategy requires knowing the future
 - Replace page that will not be used for longest period of time

FIFO Page Replacement Strategy

- Evict oldest fetched page in system
- Example
 - Assume the following page accesses:
1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- If we have 3 physical pages, how many page faults?


FIFO Page Replacement Strategy

- Evict oldest fetched page in system
- Example
 - Assume the following page accesses:
1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- If we have 3 physical pages,

9 page faults
- With 4 physical pages?

Belady's Anomaly

- More memory does not always mean lesser page faults



LRU Page Replacement Strategy

- Least recently used
 - Because past often predicts the future
 - Replace page that has least recently been used
- Example
 - Assume the following page accesses:
1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- If we have 3 physical pages, how many page faults?
- With 4 pages?

Clock Algorithm

- Use accessed bit supported by most hardware
 - E.g., x86 will write 1 to A bit in PTE on first access
 - Software managed TLBs like MIPS can do the same
- Do FIFO but skip accessed pages
- Keep pages in circular FIFO list
- Scan:
 - page's A bit = 1, set to 0 & skip
 - else if A = 0, evict
- Second-chance replacement

Clock Algorithm

- Example
 - Assume the following page accesses:
1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5
- If we have 3 physical pages, how many page faults?
- With 4 pages?

Thrashing

- Too much paging