

CS 330 - Operating Systems

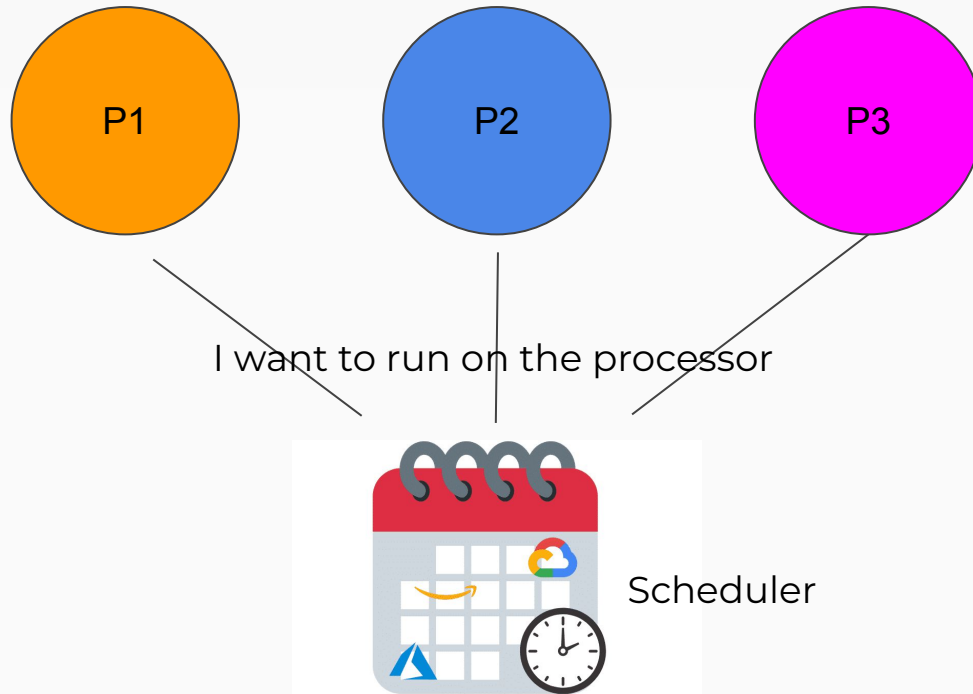
Scheduling

14-08-2025

Scheduling



What is a Scheduler?



Scheduling Policies

- How should we develop a basic framework for thinking about scheduling policies?
- What are the key assumptions about the **workload**?
- What **metrics** are important?
- What basic approaches have been used in the earliest of computer systems?

Scheduling Parameters

- When the process is created?
- Time-taken for the job to complete
- Interrupting a process
- Does the process use CPU and also do I/O operations?

Designing a Scheduling Policy


- Parameters
- Scheduling metrics

Designing a Scheduling Policy

- Parameters
- Scheduling metrics
 - Performance
 - Fairness

Scheduling Metrics


Turnaround time



$$\text{Turnaround-time}(p) = \text{Completion-time}(p) - \text{Arrival-time}(p)$$

Scheduling Metrics


Turnaround time



$$\text{Turnaround-time}(P1) = t2 - t1$$


$$\text{Turnaround-time}(P2) = t4 - t3$$

Scheduling Disciplines - FIFO



Scheduling Disciplines - FIFO

First-in First-out or FCFS



Arrival-time(P1) = 0

Completion-time(P1) = 20

Arrival-time(P2) = 10


Completion-time(P2) = 40

Arrival-time(P3) = 30

Completion-time(P3) = 40

Scheduling Disciplines - FIFO

First-in First-out or FCFS



Arrival-time(P1) = 10

Completion-time(P1) = 50

Arrival-time(P2) = 0


Completion-time(P2) = 30

Arrival-time(P3) = 20

Completion-time(P3) = 60

Scheduling Disciplines - FIFO

First-in First-out or FCFS



Arrival-time(P1) = 10

Completion-time(P1) = 50


Arrival-time(P2) = 0

Completion-time(P2) = 30

Arrival-time(P3) = 20


Completion-time(P3) = 60

Scheduling Disciplines - SJF



Scheduling Disciplines - SJF

Shortest Job First



Arrival-time(P1) = 0

Completion-time(P1) = 30

Arrival-time(P2) = 0


Completion-time(P2) = 60

Arrival-time(P3) = 0

Completion-time(P3) = 10

Scheduling Disciplines - SJF

Shortest Job First



Arrival-time(P1) = 10

Completion-time(P1) = 30

Arrival-time(P2) = 0


Completion-time(P2) = 30

Arrival-time(P3) = 10

Completion-time(P3) = 40

Scheduling Disciplines - STCF / SRTF

Shortest **Time-to-Completion (Remaining Time)** First



Arrival-time(P1) = 10

Completion-time(P1) = 40

Arrival-time(P2) = 0

Completion-time(P2) = 60

Arrival-time(P3) = 20


Completion-time(P3) = 30

Scheduling Disciplines - Priority



Scheduling Disciplines - Priority

Priority



Arrival-time(P1) = 0

Completion-time(P1) = 60

Arrival-time(P2) = 10

Completion-time(P2) = 50

Arrival-time(P3) = 20


Completion-time(P3) = 30

Scheduling Discipline - Round robin




Round-robin Scheduling

- Run job for a **time slice** and then switch
 - Scheduling quantum
- Repeats until the jobs complete
- Length of time slice is multiple of timer interrupt
- Pessimal with turnaround time



Other Scheduling Metrics

Response time



$\text{Response-time}(P) = \text{First-run-time}(P) - \text{Arrival-time}(P)$

$\text{Response-time}(P1) = t_1 - t_1$


$\text{Response-time}(P2) = t_3 - t_3$

RR scheduling - Amortization


- Length of time slice critical
 - Shorter : better performance under response time
 - Too small time slice : Context-switching costs
- Amortization is used when fixed cost of operations
 - Incur that cost less often
 - Time slice at 10 sec and context switch 1 sec (10% overhead)
 - Time slice at 100 sec and context switch 1 sec (1% overhead)

Scheduling with I/O

- Programs perform I/O
- Processes block when performing I/O



Bad usage



Good usage

Job Lengths

- How does the OS know each job's running time in SJF/STCF?
 - Normally, it doesn't

Multi-Level Feedback Queue (MLFQ)

- Described by Corbato et al. (1962)
 - Compatible Time-Sharing System (CTSS)
 - Turing Award



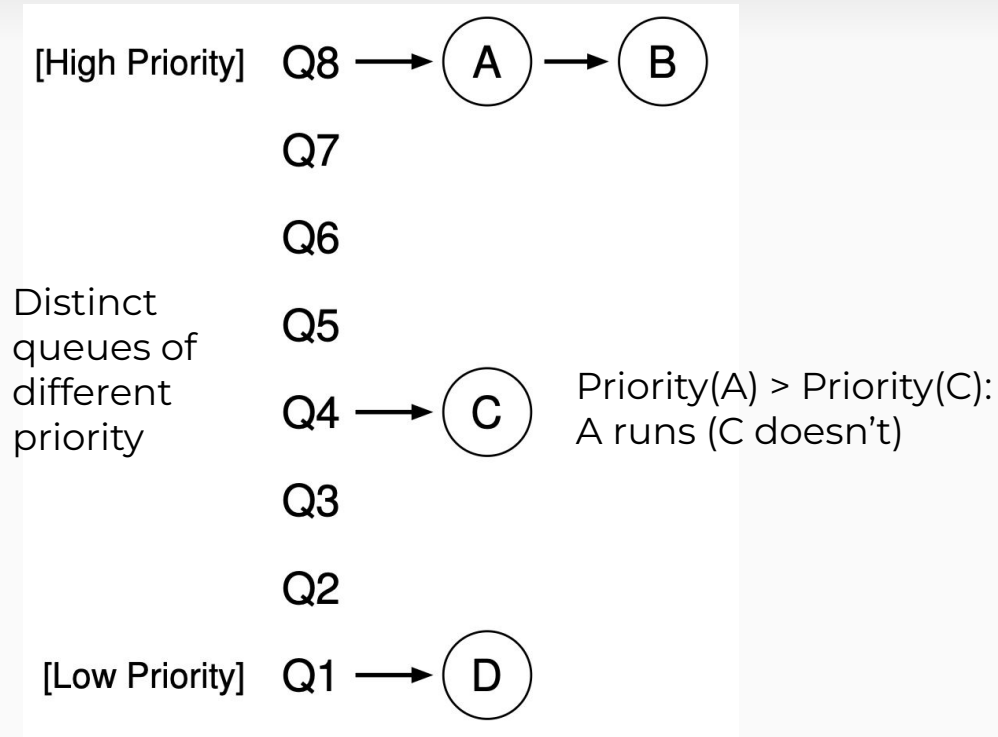
Multi-Level Feedback Queue (MLFQ)

- Described by Corbato et al. (1962)
 - Compatible Time-Sharing System (CTSS)
 - Turing Award
- Optimize turnaround time
 - Do not know running time
- Optimize response time
 - Round-robin scheduling (but with better turnaround time?)

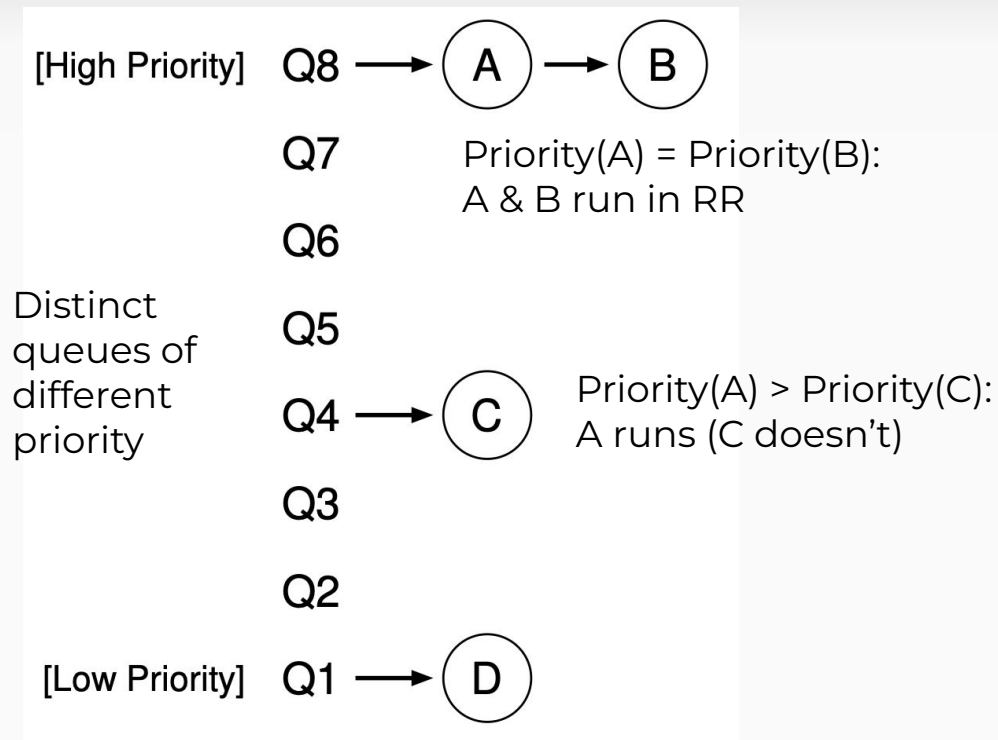
MLFQ

- Number of distinct queues
 - each assigned a different priority level (decreasing order)
- A job that is ready to run is on a single queue
 - priorities to decide which job should run at a given time
- More than one job may be on a given queue
 - have the same priority
 - round-robin scheduling among those jobs.

MLFQ : Rules



MLFQ : Rules



MLFQ: Priorities


- How to set the priorities
 - Vary priority based on behavior
 - Interactive job (repeatedly switches b/w CPU and I/O)
 - High priority
 - CPU Intensive
 - Reduce priority

MLFQ: How to change the priorities

- Mix of interactive and CPU-intensive jobs
- Algorithm


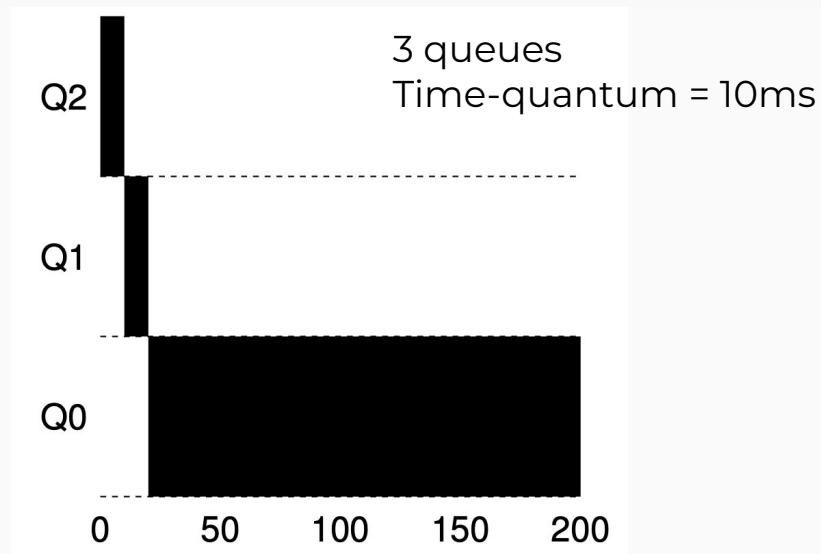
MLFQ: How to change the priorities

- Mix of interactive and CPU-intensive jobs
- Algorithm



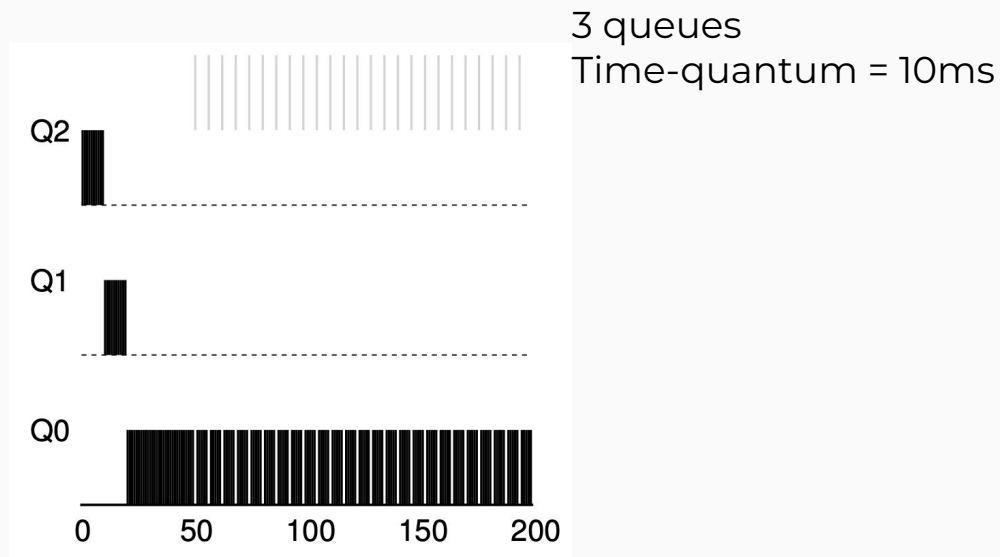
MLFQ: How to change the priorities

- Mix of interactive and CPU-intensive jobs
- Algorithm



MLFQ: How to change the priorities

- Mix of interactive and CPU-intensive jobs
- Algorithm




MLFQ: How to change the priorities

Problems with current algorithm:

- Starvation
- Trick the scheduler
- Benign program's behavior change


MLFQ: Priority Boost

- After some time period S , move all the jobs in the system to the topmost queue



MLFQ: Accounting of CPU time

- Once a process has used its allotment, it is demoted to the next priority queue



MLFQ

- If $\text{Priority}(A) > \text{Priority}(B)$, A runs (B doesn't)
- If $\text{Priority}(A) = \text{Priority}(B)$, A & B run in RR
- When a job enters the system, it is placed at the highest priority queue
- Once a job uses up its time allotment at a given level its priority is reduced (i.e., it moves down one queue)
- After some time period S , move all the jobs in the system to the topmost queue

Uniprocessor Scheduling

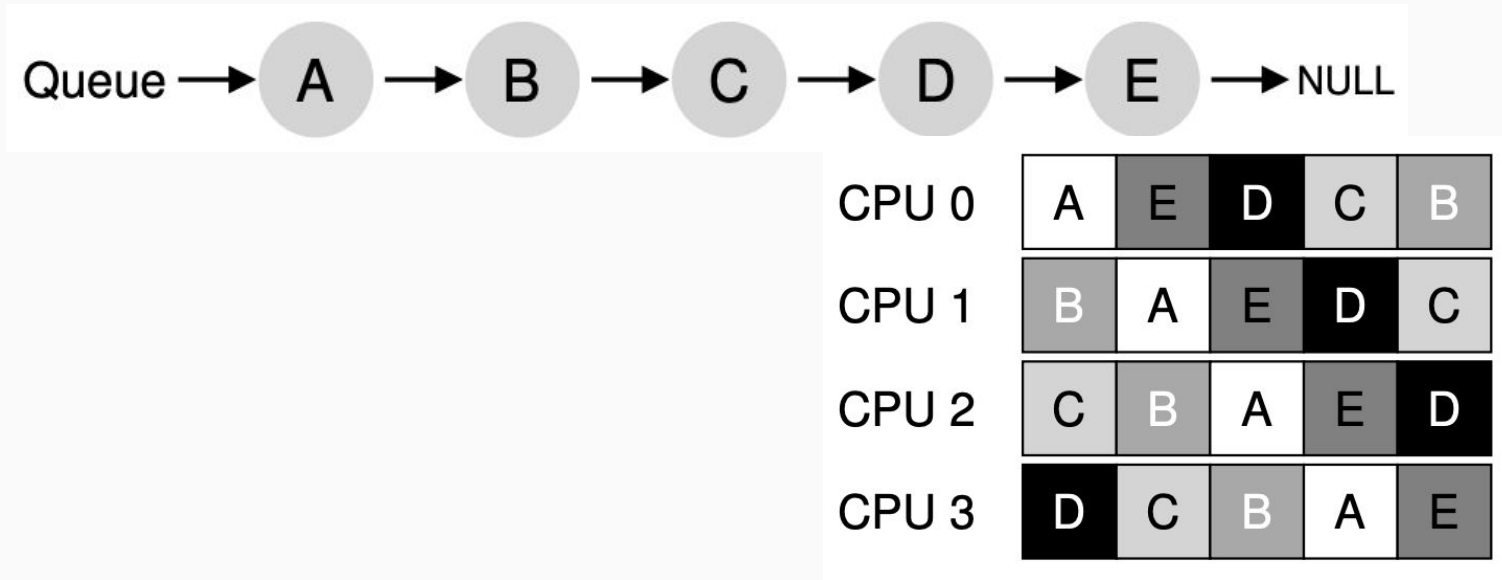
- FIFO is simple and minimizes overhead.
 - If tasks are variable in size, then FIFO can have very poor average turnaround time.
 - If tasks are equal in size, FIFO is optimal in terms of average turnaround time.
- SJF is optimal in terms of avg turnaround time.
 - SJF is pessimal in terms of variance in turnaround time.
 - If tasks are variable in size, Round Robin approximates SJF.
- For equal tasks, RR will have poor avg turnaround time.
 - Tasks that intermix processor and I/O benefit from SJF and can do poorly under Round Robin.
 - RR avoids starvation

Multiprocessor Scheduling

- Different from single-processor
 - Contains multiple CPUs
 - Multiple caches
 - Data sharing across multiple processors
- Issues due to caches
 - Data cached in CPU 1 may be required in CPU 2 due to scheduling pattern
 - Cache coherence
 - Locality - temporal and spatial - affected
 - Cache affinity


Multiprocessor Scheduling

- Single queue
 - Simple



Multiprocessor Scheduling

- Single queue
 - Simple
 - Shortcomings?
 - Scalability - multiple processors cannot access the queue at the same time
 - Cache affinity may be affected



CPU 0	A	E	D	C	B
CPU 1	B	A	E	D	C
CPU 2	C	B	A	E	D
CPU 3	D	C	B	A	E



Multiprocessor Scheduling

- Single queue
 - Simple
 - Shortcomings
 - Scalability
 - Cache affinity
- Multiple queues
 - Multiple scheduling queues follow their own scheduling algo.
 - OS decides which CPU to schedule on
 - Scalable with cache affinity




Multiprocessor Scheduling

- Multiple queues
 - Multiple scheduling queues follow their own scheduling algo.
 - OS decides which CPU to schedule on
 - Scalable with cache affinity



Load Imbalance and Migration

- Suppose job C finishes



Load Imbalance and Migration

- Or both job A & C finish




CPU 0




Load Imbalance and Migration

- Switch jobs occasionally



Load Imbalance and Migration

- Switch jobs occasionally



- How to migrate?
 - Work stealing