

CS330 - Operating Systems

Synchronization

11-09-2025

Interprocess Communication (IPC)

- Cooperating processes
 - Want to communicate with each other
 - Why?

Interprocess Communication (IPC)

- Cooperating processes
 - Want to communicate with each other
 - Allow information sharing, better performance, modularity
- But what about the restriction from the OS?
 - Prevent one process from violating other's space
 - Need some other mechanism

Producer-Consumer Problem

- Two distinct processes
 - Producer (writer) produces some data that the consumer (reader) has to consume
- How to allow this communication?

Producer-Consumer Problem

- Two distinct processes
 - Producer (writer) produces some data that the consumer (reader) has to consume
- How to allow this communication?
 - Shared memory
 - Message-passing

Producer-Consumer Problem

- Two distinct processes
 - Producer (writer) produces some data that the consumer (reader) has to consume
- Can we use a file?
 - Parent and child share file-descriptors
 - Use persistent storages like files to exchange data
 - Expensive if only transient communication is required
- Use memory-buffers
 - Finite (cannot assume unbounded buffers)
 - If producer tries to write when buffer full, it blocks
 - If consumer tries to read when buffer empty, it blocks

Threads with Shared Data

- May not give desired result
- May not give consistent result

```
static volatile int counter = 0;

// mythread()
//
// Simply adds 1 to counter repeatedly, in a loop
// No, this is not how you would add 10,000,000 to
// a counter, but it shows the problem nicely.
//
void *mythread(void *arg) {
    printf("%s: begin\n", (char *) arg);
    int i;
    for (i = 0; i < 1e7; i++) {
        counter = counter + 1;
    }
    printf("%s: done\n", (char *) arg);
    return NULL;
}
```

Uncontrolled Scheduling

- `counter = counter + 1`
 - May be translated as:
`mov 0xc %eax`
`add $0x1 %eax`
`mov %eax 0xc`

Uncontrolled Scheduling

OS	Thread 1	Thread 2	(after instruction)		
			PC	eax	counter
			100	0	50
	mov 0xc, %eax		105	50	50
	add \$0x1, %eax		108	51	50
interrupt					
<i>save T1</i>					
<i>restore T2</i>					
			100	0	50
		mov 0xc, %eax	105	50	50
		add \$0x1, %eax	108	51	50
		mov %eax, 0xc	113	51	51
interrupt					
<i>save T2</i>					
<i>restore T1</i>					
			108	51	51
	mov %eax, 0xc		113	51	51

Race Condition

Critical Section

Question: can this *panic*?

Thread 1

```
p = someComputation();  
pInitialized = true;
```

Thread 2

```
while (!pInitialized) ;  
  
q = someFunction(p);  
if (q != someFunction(p))  
    panic
```

Why *Synchronize*?

- When threads concurrently read/write shared memory, program behavior is undefined
 - Two threads write to the same variable; which one should win?
- Thread schedule is non-deterministic
 - Behavior changes when re-run program
- Compiler/hardware instruction reordering
- Multi-word operations are not ***atomic***

Atomicity

- All or none execution of an instruction
 - Execute as a unit
 - Guarantee that no other instruction will be run in between

```
mov 0xc, %eax  
add $0x1, %eax  
mov %eax, 0xc
```

- Synchronization primitives
 - Introduce some order

Definitions

- **Race condition:** output of a concurrent program depends on the order of operations between threads
- **Critical section:** piece of code that only one thread can execute at once
- **Mutual exclusion:** only one thread does a particular thing at a time
- **Synchronization:** using atomic operations to ensure cooperation between threads

Example : Buy milk

- **Ideal:**

Someone buys, if needed **and** at most one person buys

Progress

Safety

Example : Buy milk

- **Ideal:**

Someone buys, if needed **and** at most one person buys

Progress

Safety

Time

Person A

Person B

T1

Out of milk

T2

Leave for store

T3

Arrive at store; buy milk

T4

Arrive home with milk

T5

Out of milk

Leave for store

Arrive at store; buy milk

T6

Arrive home with milk

(Too much milk)

Example : Buy milk

- Leave a note

If no note

 If no milk

 Leave note

 Buy milk (go to store; buy; arrive at home)

 Remove note

Example : Buy milk

- Multiple threads

Leave noteA

If no noteB

 If no milk

 Buy milk

Remove noteA

Leave noteB

If no noteA

 If no milk

 Buy milk

Remove noteB

Example : Buy milk

- Multiple threads

Leave noteA

while noteB

do nothing

If no milk

Buy milk

Remove noteA

Leave noteB

if no noteA

If no milk

Buy milk

Remove noteB

At **while** and **if**, can guarantee that:
Safe for me to buy **or** other will buy

Locks



Image from codegym.cc

Locks

- Add locks (mutex) to code
 - Around critical sections
 - Guarantee atomic execution of critical section
- Acquire
 - wait until lock is free, then take it
- Release
 - release lock, waking up anyone waiting for it
- Correctness of Locks
 - At most one lock holder at a time (safety)
 - If no one holding, acquire gets lock (progress)
 - If all lock holders finish and no higher priority waiters, waiter eventually gets lock (progress)

Example : Buy milk

- Using locks

```
pthread_mutex_t lock = PTHREAD_MUTEX_INITIALIZER;  
// lock.acquire()
```

```
Pthread_mutex_lock(&lock); // wrapper; exits on failure
```

If no milk

 Buy milk

```
// lock.release()
```

```
Pthread_mutex_unlock(&lock);
```

Allow concurrent code to be much simpler

Locks for allocating/freeing memory

```
char *malloc (n) {  
    heaplock.acquire();  
    p = allocate n memory  
    heaplock.release();  
    return p;  
}
```

```
void free (char *p) {  
    heaplock.acquire();  
    put p back on free list  
    heaplock.release();  
}
```

Rules for using Locks

- Lock is initially free
- Always acquire before accessing shared data structure
 - Beginning of procedure!
- Always release after finishing with shared data
 - End of procedure!
 - Only the lock holder can release
 - DO NOT throw lock for someone else to release
- Never access shared data without lock
 - Danger!

Goals when Building Locks

- Mutual exclusion
 - Prevent multiple threads from entering a critical section
- Fairness
 - Each thread has fair shot of acquiring lock
- Performance
 - Time overheads

Implementing Synchronization

- Interrupt-controls
 - Disable interrupts when in critical section

Implementing Synchronization

- Interrupt-controls
 - Disable interrupts when in critical section
 - Doesn't work on multiprocessors, requires privileged op., inefficient

Implementing Synchronization

- Interrupt-controls
- Peterson's solution
 - Classic solution (works on early architectures) for two threads
 - Uses `load` and `store` (which were atomic in early arch.)

```
boolean flag[2];
int turn;

void init() {
    flag[0] = flag[1] = false;
    // whose turn is it? (thread 0 or 1)
    turn = 0;
}
```

```
void lock() {
    flag[i] = true;
    turn = j; // make it other thread's turn
    while (flag[j] && (turn == j));
    // wait until it's i's turn
}

void unlock() {
    flag[i] = false;
}
```